



# CUDA MATH API

v8.0 | February 2016

## API Reference Manual



# TABLE OF CONTENTS

<b>Chapter 1. Modules.....</b>	<b>1</b>
1.1. Mathematical Functions.....	1
1.2. Single Precision Mathematical Functions.....	2
acosf.....	2
acoshf.....	2
asinf.....	3
asinhf.....	3
atan2f.....	3
atanf.....	4
atanhf.....	4
cbrtf.....	5
ceilf.....	5
copysignf.....	5
cosf.....	6
coshf.....	6
cospif.....	6
cyl_bessel_i0f.....	7
cyl_bessel_i1f.....	7
erfcf.....	8
erfcinvf.....	8
erfcxf.....	8
erff.....	9
erfinvf.....	9
exp10f.....	10
exp2f.....	10
expf.....	10
expm1f.....	11
fabsf.....	11
fdimf.....	12
fdividef.....	12
floorf.....	12
fmaf.....	13
fmaxf.....	13
fminf.....	14
fmodf.....	14
frexpf.....	15
hypotf.....	15
ilogbf.....	16
isfinite.....	16
isinf.....	17

isnan.....	17
j0f.....	17
j1f.....	18
jnf.....	18
ldexpf.....	19
lgammaf.....	19
llrintf.....	20
llroundf.....	20
log10f.....	20
log1pf.....	21
log2f.....	21
logbf.....	21
logf.....	22
lrintf.....	22
lroundf.....	22
modff.....	23
nanf.....	23
nearbyintf.....	24
nextafterf.....	24
norm3df.....	24
norm4df.....	25
normcdf.....	25
normcdfinvf.....	26
normf.....	26
powf.....	26
rcbrtf.....	27
remainderf.....	28
remquof.....	28
rhypotf.....	29
rintf.....	29
rnorm3df.....	29
rnorm4df.....	30
rnormf.....	30
roundf.....	31
rsqrtf.....	31
scalblnf.....	31
scalbnf.....	32
signbit.....	32
sincosf.....	32
sincospif.....	33
sinf.....	33
sinhf.....	34
sinpif.....	34

sqrtf.....	35
tanf.....	35
tanhf.....	36
tgammaf.....	36
truncf.....	36
y0f.....	37
y1f.....	37
ynf.....	38
1.3. Double Precision Mathematical Functions.....	38
acos.....	38
acosh.....	39
asin.....	39
asinh.....	40
atan.....	40
atan2.....	40
atanh.....	41
cbrt.....	41
ceil.....	42
copysign.....	42
cos.....	42
cosh.....	43
cospi.....	43
cyl_bessel_i0.....	43
cyl_bessel_i1.....	44
erf.....	44
erfc.....	45
erfcinv.....	45
erfcx.....	45
erfinv.....	46
exp.....	46
exp10.....	47
exp2.....	47
expm1.....	47
fabs.....	48
fdim.....	48
floor.....	48
fma.....	49
fmax.....	49
fmin.....	50
fmod.....	50
frexp.....	51
hypot.....	51
ilogb.....	52

isfinite.....	52
isinf.....	53
isnan.....	53
j0.....	53
j1.....	54
jn.....	54
ldexp.....	55
lgamma.....	55
llrint.....	56
llround.....	56
log.....	56
log10.....	57
log1p.....	57
log2.....	58
logb.....	58
lrint.....	58
lround.....	59
modf.....	59
nan.....	60
nearbyint.....	60
nextafter.....	60
norm.....	61
norm3d.....	61
norm4d.....	62
normcdf.....	62
normcdfinv.....	62
pow.....	63
rcbrt.....	64
remainder.....	64
remquo.....	64
rhypot.....	65
rint.....	65
rnorm.....	66
rnorm3d.....	66
rnorm4d.....	67
round.....	67
rsqrt.....	67
scalbln.....	68
scalbn.....	68
signbit.....	69
sin.....	69
sincos.....	69
sincospi.....	70

sinh.....	70
sinpi.....	71
sqrt.....	71
tan.....	72
tanh.....	72
tgamma.....	72
trunc.....	73
y0.....	73
y1.....	74
yn.....	74
1.4. Single Precision Intrinsics.....	75
__cosf.....	75
__exp10f.....	75
__expf.....	76
__fadd_rd.....	76
__fadd_rn.....	76
__fadd_ru.....	77
__fadd_rz.....	77
__fdiv_rd.....	77
__fdiv_rn.....	78
__fdiv_ru.....	78
__fdiv_rz.....	78
__fdividef.....	79
__fmaf_rd.....	79
__fmaf_rn.....	80
__fmaf_ru.....	80
__fmaf_rz.....	81
__fmul_rd.....	81
__fmul_rn.....	82
__fmul_ru.....	82
__fmul_rz.....	82
__frcp_rd.....	83
__frcp_rn.....	83
__frcp_ru.....	83
__frcp_rz.....	84
__frsqrt_rn.....	84
__fsqrt_rd.....	85
__fsqrt_rn.....	85
__fsqrt_ru.....	85
__fsqrt_rz.....	86
__fsub_rd.....	86
__fsub_rn.....	86
__fsub_ru.....	87

__fsub_rz.....	87
__log10f.....	88
__log2f.....	88
__logf.....	88
__powf.....	89
__saturatef.....	89
__sincosf.....	90
__sinf.....	90
__tanf.....	90
1.5. Double Precision Intrinsics.....	91
__dadd_rd.....	91
__dadd_rn.....	91
__dadd_ru.....	92
__dadd_rz.....	92
__ddiv_rd.....	92
__ddiv_rn.....	93
__ddiv_ru.....	93
__ddiv_rz.....	94
__dmul_rd.....	94
__dmul_rn.....	94
__dmul_ru.....	95
__dmul_rz.....	95
__drcp_rd.....	95
__drcp_rn.....	96
__drcp_ru.....	96
__drcp_rz.....	97
__dsqrt_rd.....	97
__dsqrt_rn.....	97
__dsqrt_ru.....	98
__dsqrt_rz.....	98
__dsub_rd.....	99
__dsub_rn.....	99
__dsub_ru.....	99
__dsub_rz.....	100
__fma_rd.....	100
__fma_rn.....	101
__fma_ru.....	101
__fma_rz.....	102
1.6. Integer Intrinsics.....	102
__brev.....	102
__brevll.....	103
__byte_perm.....	103
__clz.....	103

__clzll.....	104
__ffs.....	104
__ffsll.....	104
__hadd.....	105
__mul24.....	105
__mul64hi.....	105
__mulhi.....	106
__popc.....	106
__popcll.....	106
__rhadd.....	106
__sad.....	107
__uhadd.....	107
__umul24.....	107
__umul64hi.....	108
__umulhi.....	108
__urhadd.....	108
__usad.....	109
1.7. Type Casting Intrinsics.....	109
__double2float_rd.....	109
__double2float_rn.....	109
__double2float_ru.....	110
__double2float_rz.....	110
__double2hiint.....	110
__double2int_rd.....	110
__double2int_rn.....	111
__double2int_ru.....	111
__double2int_rz.....	111
__double2ll_rd.....	112
__double2ll_rn.....	112
__double2ll_ru.....	112
__double2ll_rz.....	112
__double2loint.....	113
__double2uint_rd.....	113
__double2uint_rn.....	113
__double2uint_ru.....	114
__double2uint_rz.....	114
__double2ull_rd.....	114
__double2ull_rn.....	114
__double2ull_ru.....	115
__double2ull_rz.....	115
__double_as_longlong.....	115
__float2half_rn.....	116
__float2int_rd.....	116



__float2int_rn.....	116
__float2int_ru.....	116
__float2int_rz.....	117
__float2ll_rd.....	117
__float2ll_rn.....	117
__float2ll_ru.....	118
__float2ll_rz.....	118
__float2uint_rd.....	118
__float2uint_rn.....	118
__float2uint_ru.....	119
__float2uint_rz.....	119
__float2ull_rd.....	119
__float2ull_rn.....	120
__float2ull_ru.....	120
__float2ull_rz.....	120
__float_as_int.....	121
__float_as_uint.....	121
__half2float.....	121
__hiloint2double.....	121
__int2double_rn.....	122
__int2float_rd.....	122
__int2float_rn.....	122
__int2float_ru.....	123
__int2float_rz.....	123
__int_as_float.....	123
__ll2double_rd.....	123
__ll2double_rn.....	124
__ll2double_ru.....	124
__ll2double_rz.....	124
__ll2float_rd.....	125
__ll2float_rn.....	125
__ll2float_ru.....	125
__ll2float_rz.....	125
__longlong_as_double.....	126
__uint2double_rn.....	126
__uint2float_rd.....	126
__uint2float_rn.....	127
__uint2float_ru.....	127
__uint2float_rz.....	127
__uint_as_float.....	127
__ull2double_rd.....	128
__ull2double_rn.....	128
__ull2double_ru.....	128

__ull2double_rz.....	129
__ull2float_rd.....	129
__ull2float_rn.....	129
__ull2float_ru.....	130
__ull2float_rz.....	130
1.8. SIMD Intrinsics.....	130
__vabs2.....	130
__vabs4.....	131
__vabsdiffs2.....	131
__vabsdiffs4.....	131
__vabsdiffu2.....	132
__vabsdiffu4.....	132
__vabsss2.....	132
__vabsss4.....	133
__vadd2.....	133
__vadd4.....	133
__vaddss2.....	134
__vaddss4.....	134
__vaddus2.....	134
__vaddus4.....	135
__vavg2.....	135
__vavg4.....	135
__vavg2.....	136
__vavg4.....	136
__vcmpeq2.....	136
__vcmpeq4.....	137
__vcmpges2.....	137
__vcmpges4.....	137
__vcmpgeu2.....	138
__vcmpgeu4.....	138
__vcmpgts2.....	138
__vcmpgts4.....	139
__vcmpgtu2.....	139
__vcmpgtu4.....	139
__vcmples2.....	140
__vcmples4.....	140
__vcmpleu2.....	140
__vcmpleu4.....	141
__vcmplt2.....	141
__vcmplt4.....	141
__vcmpltu2.....	142
__vcmpltu4.....	142
__vcmpne2.....	142

__vcmpne4.....	143
__vhaddu2.....	143
__vhaddu4.....	143
__vmaxs2.....	144
__vmaxs4.....	144
__vmaxu2.....	144
__vmaxu4.....	145
__vmins2.....	145
__vmins4.....	145
__vminu2.....	146
__vminu4.....	146
__vneg2.....	146
__vneg4.....	147
__vnegss2.....	147
__vnegss4.....	147
__vsads2.....	147
__vsads4.....	148
__vsadu2.....	148
__vsadu4.....	148
__vseteq2.....	149
__vseteq4.....	149
__vsetges2.....	149
__vsetges4.....	150
__vsetgeu2.....	150
__vsetgeu4.....	150
__vsetgts2.....	151
__vsetgts4.....	151
__vsetgtu2.....	151
__vsetgtu4.....	152
__vsetles2.....	152
__vsetles4.....	152
__vsetleu2.....	153
__vsetleu4.....	153
__vsetlts2.....	153
__vsetlts4.....	154
__vsetltu2.....	154
__vsetltu4.....	154
__vsetne2.....	155
__vsetne4.....	155
__vsub2.....	155
__vsub4.....	156
__vsubss2.....	156
__vsubss4.....	156

__vsubus2.....	157
__vsubus4.....	157
1.9. Half Precision Intrinsics.....	157
Half Arithmetic Functions.....	157
Half2 Arithmetic Functions.....	157
Half Comparison Functions.....	158
Half2 Comparison Functions.....	158
Half Precision Conversion And Data Movement.....	158
Half Math Functions.....	158
Half2 Math Functions.....	158
1.9.1. Half Arithmetic Functions.....	158
__hadd.....	158
__hadd_sat.....	158
__hfma.....	159
__hfma_sat.....	159
__hmul.....	159
__hmul_sat.....	160
__hneg.....	160
__hsub.....	160
__hsub_sat.....	160
h2div.....	161
hdiv.....	161
1.9.2. Half2 Arithmetic Functions.....	161
__hadd2.....	161
__hadd2_sat.....	161
__hfma2.....	162
__hfma2_sat.....	162
__hmul2.....	162
__hmul2_sat.....	163
__hneg2.....	163
__hsub2.....	163
__hsub2_sat.....	163
1.9.3. Half Comparison Functions.....	164
__heq.....	164
__hequ.....	164
__hge.....	164
__hgeu.....	165
__hgt.....	165
__hgtu.....	165
__hisinf.....	165
__hisnan.....	166
__hle.....	166
__hleu.....	166

__hlt.....	166
__hltu.....	167
__hne.....	167
__hneu.....	167
1.9.4. Half2 Comparison Functions.....	167
__hbeq2.....	168
__hbequ2.....	168
__hbge2.....	168
__hbgeu2.....	169
__hbg2.....	169
__hbg2u2.....	169
__hble2.....	170
__hbleu2.....	170
__hblt2.....	170
__hbltu2.....	171
__hbne2.....	171
__hbneu2.....	171
__heq2.....	172
__hequ2.....	172
__hge2.....	172
__hgeu2.....	172
__hgt2.....	173
__hgtu2.....	173
__hisnan2.....	173
__hle2.....	174
__hleu2.....	174
__hlt2.....	174
__hltu2.....	174
__hne2.....	175
__hneu2.....	175
1.9.5. Half Precision Conversion And Data Movement.....	175
__float22half2_rn.....	175
__float2half.....	176
__float2half2_rn.....	176
__float2half_rd.....	176
__float2half_ru.....	176
__float2half_rz.....	177
__floats2half2_rn.....	177
__half22float2.....	177
__half2float.....	177
__half2half2.....	178
__half2int_rd.....	178
__half2int_rn.....	178

__half2int_ru.....	178
__half2int_rz.....	179
__half2ll_rd.....	179
__half2ll_rn.....	179
__half2ll_ru.....	179
__half2ll_rz.....	180
__half2short_rd.....	180
__half2short_rn.....	180
__half2short_ru.....	180
__half2short_rz.....	181
__half2uint_rd.....	181
__half2uint_rn.....	181
__half2uint_ru.....	181
__half2uint_rz.....	182
__half2ull_rd.....	182
__half2ull_rn.....	182
__half2ull_ru.....	182
__half2ull_rz.....	183
__half2ushort_rd.....	183
__half2ushort_rn.....	183
__half2ushort_ru.....	183
__half2ushort_rz.....	184
__halves2half2.....	184
__high2float.....	184
__high2half.....	184
__high2half2.....	185
__highs2half2.....	185
__int2half_rd.....	185
__int2half_rn.....	186
__int2half_ru.....	186
__int2half_rz.....	186
__ll2half_rd.....	186
__ll2half_rn.....	187
__ll2half_ru.....	187
__ll2half_rz.....	187
__low2float.....	187
__low2half.....	188
__low2half2.....	188
__lowhigh2highlow.....	188
__lows2half2.....	188
__short2half_rd.....	189
__short2half_rn.....	189
__short2half_ru.....	189

__short2half_rz.....	190
__uint2half_rd.....	190
__uint2half_rn.....	190
__uint2half_ru.....	190
__uint2half_rz.....	191
__ull2half_rd.....	191
__ull2half_rn.....	191
__ull2half_ru.....	191
__ull2half_rz.....	192
__ushort2half_rd.....	192
__ushort2half_rn.....	192
__ushort2half_ru.....	192
__ushort2half_rz.....	193
1.9.6. Half Math Functions.....	193
hceil.....	193
hcos.....	193
hexp.....	194
hexp10.....	194
hexp2.....	194
hfloor.....	194
hlog.....	195
hlog10.....	195
hlog2.....	195
hrcp.....	195
hrint.....	196
hrsqrt.....	196
hsin.....	196
hsqrt.....	196
htrunc.....	197
1.9.7. Half2 Math Functions.....	197
h2ceil.....	197
h2cos.....	197
h2exp.....	197
h2exp10.....	198
h2exp2.....	198
h2floor.....	198
h2log.....	198
h2log10.....	199
h2log2.....	199
h2rcp.....	199
h2rint.....	199
h2rsqrt.....	200
h2sin.....	200

h2sqrt.....	200
h2trunc.....	200



# Chapter 1.

## MODULES

Here is a list of all modules:

- ▶ Mathematical Functions
- ▶ Single Precision Mathematical Functions
- ▶ Double Precision Mathematical Functions
- ▶ Single Precision Intrinsics
- ▶ Double Precision Intrinsics
- ▶ Integer Intrinsics
- ▶ Type Casting Intrinsics
- ▶ SIMD Intrinsics
- ▶ Half Precision Intrinsics
  - ▶ Half Arithmetic Functions
  - ▶ Half2 Arithmetic Functions
  - ▶ Half Comparison Functions
  - ▶ Half2 Comparison Functions
  - ▶ Half Precision Conversion And Data Movement
  - ▶ Half Math Functions
  - ▶ Half2 Math Functions

## 1.1. Mathematical Functions

CUDA mathematical functions are always available in device code. Some functions are also available in host code as indicated.

Note that floating-point functions are overloaded for different argument types. For example, the `log()` function has the following prototypes:

```
double log(double x);  
float log(float x);  
float logf(float x);
```

## 1.2. Single Precision Mathematical Functions

This section describes single precision mathematical functions.

### `__device__ float acosf (float x)`

Calculate the arc cosine of the input argument.

#### Returns

Result will be in radians, in the interval  $[0, \pi]$  for  $x$  inside  $[-1, +1]$ .

- ▶ `acosf(1)` returns  $+0$ .
- ▶ `acosf(x)` returns NaN for  $x$  outside  $[-1, +1]$ .

#### Description

Calculate the principal value of the arc cosine of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

### `__device__ float acoshf (float x)`

Calculate the nonnegative arc hyperbolic cosine of the input argument.

#### Returns

Result will be in the interval  $[0, +\infty]$ .

- ▶ `acoshf(1)` returns  $0$ .
- ▶ `acoshf(x)` returns NaN for  $x$  in the interval  $[-\infty, 1)$ .

#### Description

Calculate the nonnegative arc hyperbolic cosine of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float asinf (float x)`

Calculate the arc sine of the input argument.

### Returns

Result will be in radians, in the interval  $[-\pi/2, +\pi/2]$  for  $x$  inside  $[-1, +1]$ .

- ▶ `asinf(0)` returns +0.
- ▶ `asinf(x)` returns NaN for  $x$  outside  $[-1, +1]$ .

### Description

Calculate the principal value of the arc sine of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float asinhf (float x)`

Calculate the arc hyperbolic sine of the input argument.

### Returns

- ▶ `asinhf(0)` returns 1.

### Description

Calculate the arc hyperbolic sine of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float atan2f (float y, float x)`

Calculate the arc tangent of the ratio of first and second input arguments.

### Returns

Result will be in radians, in the interval  $[-\pi, +\pi]$ .

- ▶ `atan2f(0, 1)` returns +0.

**Description**

Calculate the principal value of the arc tangent of the ratio of first and second input arguments  $y / x$ . The quadrant of the result is determined by the signs of inputs  $y$  and  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

**\_\_device\_\_ float atanf (float x)**

Calculate the arc tangent of the input argument.

**Returns**

Result will be in radians, in the interval  $[-\pi/2, +\pi/2]$ .

- ▶ `atanf(0)` returns `+0`.

**Description**

Calculate the principal value of the arc tangent of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

**\_\_device\_\_ float atanhf (float x)**

Calculate the arc hyperbolic tangent of the input argument.

**Returns**

- ▶ `atanhf( ± 0 )` returns `± 0`.
- ▶ `atanhf( ± 1 )` returns `± ∞`.
- ▶ `atanhf(x)` returns NaN for  $x$  outside interval  $[-1, 1]$ .

**Description**

Calculate the arc hyperbolic tangent of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float cbrtf (float x)`

Calculate the cube root of the input argument.

### Returns

Returns  $x^{1/3}$ .

- ▶ `cbrtf( ±0 )` returns  $\pm 0$ .
- ▶ `cbrtf( ±∞ )` returns  $\pm \infty$ .

### Description

Calculate the cube root of  $x$ ,  $x^{1/3}$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float ceilf (float x)`

Calculate ceiling of the input argument.

### Returns

Returns  $\lceil x \rceil$  expressed as a floating-point number.

- ▶ `ceilf( ±0 )` returns  $\pm 0$ .
- ▶ `ceilf( ±∞ )` returns  $\pm \infty$ .

### Description

Compute the smallest integer value not less than  $x$ .

## `__device__ float copysignf (float x, float y)`

Create value with given magnitude, copying sign of second value.

### Returns

Returns a value with the magnitude of  $x$  and the sign of  $y$ .

### Description

Create a floating-point value with the magnitude  $x$  and the sign of  $y$ .

## `__device__ float cosf (float x)`

Calculate the cosine of the input argument.

### Returns

- ▶ `cosf(0)` returns 1.
- ▶ `cosf(  $\pm \infty$  )` returns NaN.

### Description

Calculate the cosine of the input argument  $x$  (measured in radians).



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This function is affected by the `--use_fast_math` compiler flag. See the CUDA C Programming Guide, Appendix D.2, Table 8 for a complete list of functions affected.

## `__device__ float coshf (float x)`

Calculate the hyperbolic cosine of the input argument.

### Returns

- ▶ `coshf(0)` returns 1.
- ▶ `coshf(  $\pm \infty$  )` returns NaN.

### Description

Calculate the hyperbolic cosine of the input argument  $x$ .



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float cospif (float x)`

Calculate the cosine of the input argument  $x\pi$ .

### Returns

- ▶ `cospif(  $\pm 0$  )` returns 1.
- ▶ `cospif(  $\pm \infty$  )` returns NaN.

**Description**

Calculate the cosine of  $x \times \pi$  (measured in radians), where  $x$  is the input argument.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

**\_\_device\_\_ float cyl\_bessel\_i0f (float x)**

Calculate the value of the regular modified cylindrical Bessel function of order 0 for the input argument.

**Returns**

Returns the value of the regular modified cylindrical Bessel function of order 0.

**Description**

Calculate the value of the regular modified cylindrical Bessel function of order 0 for the input argument  $x$ ,  $I_0(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

**\_\_device\_\_ float cyl\_bessel\_i1f (float x)**

Calculate the value of the regular modified cylindrical Bessel function of order 1 for the input argument.

**Returns**

Returns the value of the regular modified cylindrical Bessel function of order 1.

**Description**

Calculate the value of the regular modified cylindrical Bessel function of order 1 for the input argument  $x$ ,  $I_1(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float erfcf (float x)`

Calculate the complementary error function of the input argument.

### Returns

- ▶ `erfcf( -  $\infty$  )` returns 2.
- ▶ `erfcf( +  $\infty$  )` returns +0.

### Description

Calculate the complementary error function of the input argument  $x$ ,  $1 - \text{erf}(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float erfcinvf (float y)`

Calculate the inverse complementary error function of the input argument.

### Returns

- ▶ `erfcinvf(0)` returns  $+\infty$ .
- ▶ `erfcinvf(2)` returns  $-\infty$ .

### Description

Calculate the inverse complementary error function of the input argument  $y$ , for  $y$  in the interval  $[0, 2]$ . The inverse complementary error function find the value  $x$  that satisfies the equation  $y = \text{erfc}(x)$ , for  $0 \leq y \leq 2$ , and  $-\infty \leq x \leq \infty$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float erfcxf (float x)`

Calculate the scaled complementary error function of the input argument.

### Returns

- ▶ `erfcxf( -  $\infty$  )` returns  $+\infty$
- ▶ `erfcxf( +  $\infty$  )` returns +0
- ▶ `erfcxf( $x$ )` returns  $+\infty$  if the correctly calculated value is outside the single floating point range.



**Description**

Calculate the scaled complementary error function of the input argument  $x$ ,  $e^{x^2} \cdot \operatorname{erfc}(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

**\_\_device\_\_ float erff (float x)**

Calculate the error function of the input argument.

**Returns**

- ▶  $\operatorname{erff}(\pm 0)$  returns  $\pm 0$ .
- ▶  $\operatorname{erff}(\pm \infty)$  returns  $\pm 1$ .

**Description**

Calculate the value of the error function for the input argument  $x$ ,  $\frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

**\_\_device\_\_ float erfinvf (float y)**

Calculate the inverse error function of the input argument.

**Returns**

- ▶  $\operatorname{erfinvf}(1)$  returns  $+\infty$ .
- ▶  $\operatorname{erfinvf}(-1)$  returns  $-\infty$ .

**Description**

Calculate the inverse error function of the input argument  $y$ , for  $y$  in the interval  $[-1, 1]$ . The inverse error function finds the value  $x$  that satisfies the equation  $y = \operatorname{erf}(x)$ , for  $-1 \leq y \leq 1$ , and  $-\infty \leq x \leq \infty$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float exp10f (float x)`

Calculate the base 10 exponential of the input argument.

### Returns

Returns  $10^x$ .

### Description

Calculate the base 10 exponential of the input argument  $x$ .



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This function is affected by the `--use_fast_math` compiler flag. See the CUDA C Programming Guide, Appendix D.2, Table 8 for a complete list of functions affected.

## `__device__ float exp2f (float x)`

Calculate the base 2 exponential of the input argument.

### Returns

Returns  $2^x$ .

### Description

Calculate the base 2 exponential of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float expf (float x)`

Calculate the base  $e$  exponential of the input argument.

### Returns

Returns  $e^x$ .

### Description

Calculate the base  $e$  exponential of the input argument  $x$ ,  $e^x$ .



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This function is affected by the `--use_fast_math` compiler flag. See the CUDA C Programming Guide, Appendix D.2, Table 8 for a complete list of functions affected.

## `__device__ float expm1f (float x)`

Calculate the base  $e$  exponential of the input argument, minus 1.

### Returns

Returns  $e^x - 1$ .

### Description

Calculate the base  $e$  exponential of the input argument  $x$ , minus 1.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float fabsf (float x)`

Calculate the absolute value of its argument.

### Returns

Returns the absolute value of its argument.

- ▶ `fabs(  $\pm\infty$  )` returns  $+\infty$ .
- ▶ `fabs(  $\pm 0$  )` returns 0.

### Description

Calculate the absolute value of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float fdimf (float x, float y)`

Compute the positive difference between  $x$  and  $y$ .

### Returns

Returns the positive difference between  $x$  and  $y$ .

- ▶ `fdimf(x, y)` returns  $x - y$  if  $x > y$ .
- ▶ `fdimf(x, y)` returns  $+0$  if  $x \leq y$ .

### Description

Compute the positive difference between  $x$  and  $y$ . The positive difference is  $x - y$  when  $x > y$  and  $+0$  otherwise.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float fdividef (float x, float y)`

Divide two floating point values.

### Returns

Returns  $x / y$ .

### Description

Compute  $x$  divided by  $y$ . If `--use_fast_math` is specified, use `__fdividef()` for higher performance, otherwise use normal division.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This function is affected by the `--use_fast_math` compiler flag. See the CUDA C Programming Guide, Appendix D.2, Table 8 for a complete list of functions affected.

## `__device__ float floorf (float x)`

Calculate the largest integer less than or equal to  $x$ .

### Returns

Returns  $\log_e(1+x)$  expressed as a floating-point number.

- ▶ `floorf(  $\pm\infty$  )` returns  $\pm\infty$ .
- ▶ `floorf(  $\pm 0$  )` returns  $\pm 0$ .

### Description

Calculate the largest integer value which is less than or equal to  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float fmaf (float x, float y, float z)`

Compute  $x \times y + z$  as a single operation.

### Returns

Returns the rounded value of  $x \times y + z$  as a single operation.

- ▶ `fmaf(  $\pm\infty$ ,  $\pm 0$ ,  $z$  )` returns NaN.
- ▶ `fmaf(  $\pm 0$ ,  $\pm\infty$ ,  $z$  )` returns NaN.
- ▶ `fmaf( $x$ ,  $y$ ,  $-\infty$  )` returns NaN if  $x \times y$  is an exact  $+\infty$ .
- ▶ `fmaf( $x$ ,  $y$ ,  $+\infty$  )` returns NaN if  $x \times y$  is an exact  $-\infty$ .

### Description

Compute the value of  $x \times y + z$  as a single ternary operation. After computing the value to infinite precision, the value is rounded once.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float fmaxf (float x, float y)`

Determine the maximum numeric value of the arguments.

### Returns

Returns the maximum numeric values of the arguments  $x$  and  $y$ .

- ▶ If both arguments are NaN, returns NaN.
- ▶ If one argument is NaN, returns the numeric argument.

## Description

Determines the maximum numeric value of the arguments  $x$  and  $y$ . Treats NaN arguments as missing data. If one argument is a NaN and the other is legitimate numeric value, the numeric value is chosen.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float fminf (float x, float y)`

Determine the minimum numeric value of the arguments.

## Returns

Returns the minimum numeric values of the arguments  $x$  and  $y$ .

- ▶ If both arguments are NaN, returns NaN.
- ▶ If one argument is NaN, returns the numeric argument.

## Description

Determines the minimum numeric value of the arguments  $x$  and  $y$ . Treats NaN arguments as missing data. If one argument is a NaN and the other is legitimate numeric value, the numeric value is chosen.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float fmodf (float x, float y)`

Calculate the floating-point remainder of  $x / y$ .

## Returns

- ▶ Returns the floating point remainder of  $x / y$ .
- ▶ `fmodf(  $\pm 0$ ,  $y$ )` returns  $\pm 0$  if  $y$  is not zero.
- ▶ `fmodf( $x$ ,  $y$ )` returns NaN and raised an invalid floating point exception if  $x$  is  $\pm \infty$  or  $y$  is zero.
- ▶ `fmodf( $x$ ,  $y$ )` returns zero if  $y$  is zero or the result would overflow.
- ▶ `fmodf( $x$ ,  $\pm \infty$ )` returns  $x$  if  $x$  is finite.
- ▶ `fmodf( $x$ , 0)` returns NaN.

## Description

Calculate the floating-point remainder of  $x / y$ . The absolute value of the computed value is always less than  $y$ 's absolute value and will have the same sign as  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float frexpf (float x, int *nptr)`

Extract mantissa and exponent of a floating-point value.

## Returns

Returns the fractional component  $m$ .

- ▶ `frexpf(0, nptr)` returns 0 for the fractional component and zero for the integer component.
- ▶ `frexpf( $\pm 0$ , nptr)` returns  $\pm 0$  and stores zero in the location pointed to by `nptr`.
- ▶ `frexpf( $\pm \infty$ , nptr)` returns  $\pm \infty$  and stores an unspecified value in the location to which `nptr` points.
- ▶ `frexpf(NaN, y)` returns a NaN and stores an unspecified value in the location to which `nptr` points.

## Description

Decomposes the floating-point value  $x$  into a component  $m$  for the normalized fraction element and another term  $n$  for the exponent. The absolute value of  $m$  will be greater than or equal to 0.5 and less than 1.0 or it will be equal to 0;  $x = m \cdot 2^n$ . The integer exponent  $n$  will be stored in the location to which `nptr` points.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float __CRTDECL hypotf (float x, float y)`

Calculate the square root of the sum of squares of two arguments.

## Returns

Returns the length of the hypotenuse  $\sqrt{x^2 + y^2}$ . If the correct value would overflow, returns  $+\infty$ . If the correct value would underflow, returns 0.

## Description

Calculates the length of the hypotenuse of a right triangle whose two sides have lengths  $x$  and  $y$  without undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ int ilogbf (float x)`

Compute the unbiased integer exponent of the argument.

### Returns

- ▶ If successful, returns the unbiased exponent of the argument.
- ▶ `ilogbf(0)` returns `INT_MIN`.
- ▶ `ilogbf(NaN)` returns `NaN`.
- ▶ `ilogbf(x)` returns `INT_MAX` if  $x$  is  $\infty$  or the correct value is greater than `INT_MAX`.
- ▶ `ilogbf(x)` return `INT_MIN` if the correct value is less than `INT_MIN`.

## Description

Calculates the unbiased integer exponent of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ __RETURN_TYPE isfinite (float a)`

Determine whether argument is finite.

### Returns

- ▶ With Visual Studio 2013 host compiler: `__RETURN_TYPE` is 'bool'. Returns true if and only if  $a$  is a finite value.
- ▶ With other host compilers: `__RETURN_TYPE` is 'int'. Returns a nonzero value if and only if  $a$  is a finite value.

## Description

Determine whether the floating-point value  $a$  is a finite value (zero, subnormal, or normal and not infinity or NaN).



## `__device__ __RETURN_TYPE isinf (float a)`

Determine whether argument is infinite.

### Returns

- ▶ With Visual Studio 2013 host compiler: `__RETURN_TYPE` is 'bool'. Returns true if and only if `a` is a infinite value.
- ▶ With other host compilers: `__RETURN_TYPE` is 'int'. Returns a nonzero value if and only if `a` is a infinite value.

### Description

Determine whether the floating-point value `a` is an infinite value (positive or negative).

## `__device__ __RETURN_TYPE isnan (float a)`

Determine whether argument is a NaN.

### Returns

- ▶ With Visual Studio 2013 host compiler: `__RETURN_TYPE` is 'bool'. Returns true if and only if `a` is a NaN value.
- ▶ With other host compilers: `__RETURN_TYPE` is 'int'. Returns a nonzero value if and only if `a` is a NaN value.

### Description

Determine whether the floating-point value `a` is a NaN.

## `__device__ float j0f (float x)`

Calculate the value of the Bessel function of the first kind of order 0 for the input argument.

### Returns

Returns the value of the Bessel function of the first kind of order 0.

- ▶ `j0f( ±∞ )` returns +0.
- ▶ `j0f(NaN)` returns NaN.

### Description

Calculate the value of the Bessel function of the first kind of order 0 for the input argument `x`,  $J_0(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float j1f (float x)`

Calculate the value of the Bessel function of the first kind of order 1 for the input argument.

### Returns

Returns the value of the Bessel function of the first kind of order 1.

- ▶ `j1f( ± 0 )` returns  $\pm 0$ .
- ▶ `j1f( ± ∞ )` returns  $+0$ .
- ▶ `j1f(NaN)` returns NaN.

### Description

Calculate the value of the Bessel function of the first kind of order 1 for the input argument  $x$ ,  $J_1(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float jnf (int n, float x)`

Calculate the value of the Bessel function of the first kind of order  $n$  for the input argument.

### Returns

Returns the value of the Bessel function of the first kind of order  $n$ .

- ▶ `jnf(n, NaN)` returns NaN.
- ▶ `jnf(n, x)` returns NaN for  $n < 0$ .
- ▶ `jnf(n, + ∞ )` returns  $+0$ .

### Description

Calculate the value of the Bessel function of the first kind of order  $n$  for the input argument  $x$ ,  $J_n(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float ldexpf (float x, int exp)`

Calculate the value of  $x \cdot 2^{\text{exp}}$ .

### Returns

- `ldexpf(x)` returns  $\pm \infty$  if the correctly calculated value is outside the single floating point range.

### Description

Calculate the value of  $x \cdot 2^{\text{exp}}$  of the input arguments `x` and `exp`.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float lgammaf (float x)`

Calculate the natural logarithm of the absolute value of the gamma function of the input argument.

### Returns

- `lgammaf(1)` returns +0.
- `lgammaf(2)` returns +0.
- `lgammaf(x)` returns  $\pm \infty$  if the correctly calculated value is outside the single floating point range.
- `lgammaf(x)` returns  $+\infty$  if  $x \leq 0$ .
- `lgammaf(  $-\infty$  )` returns  $-\infty$ .
- `lgammaf(  $+\infty$  )` returns  $+\infty$ .

### Description

Calculate the natural logarithm of the absolute value of the gamma function of the input argument `x`, namely the value of  $\log_e \int_0^\infty e^{-t} t^{x-1} dt$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ long long int llrintf (float x)`

Round input to nearest integer value.

### Returns

Returns rounded integer value.

### Description

Round  $x$  to the nearest integer value, with halfway cases rounded towards zero. If the result is outside the range of the return type, the result is undefined.

## `__device__ long long int llroundf (float x)`

Round to nearest integer value.

### Returns

Returns rounded integer value.

### Description

Round  $x$  to the nearest integer value, with halfway cases rounded away from zero. If the result is outside the range of the return type, the result is undefined.



This function may be slower than alternate rounding methods. See [llrintf\(\)](#).

## `__device__ float log10f (float x)`

Calculate the base 10 logarithm of the input argument.

### Returns

- ▶  $\log_{10}f(\pm 0)$  returns  $-\infty$ .
- ▶  $\log_{10}f(1)$  returns  $+0$ .
- ▶  $\log_{10}f(x)$  returns NaN for  $x < 0$ .
- ▶  $\log_{10}f(+\infty)$  returns  $+\infty$ .

### Description

Calculate the base 10 logarithm of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## \_\_device\_\_ float log1pf (float x)

Calculate the value of  $\log_e(1+x)$ .

### Returns

- ▶  $\log1pf(\pm 0)$  returns  $-\infty$ .
- ▶  $\log1pf(-1)$  returns  $+0$ .
- ▶  $\log1pf(x)$  returns NaN for  $x < -1$ .
- ▶  $\log1pf(+\infty)$  returns  $+\infty$ .

### Description

Calculate the value of  $\log_e(1+x)$  of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## \_\_device\_\_ float log2f (float x)

Calculate the base 2 logarithm of the input argument.

### Returns

- ▶  $\log2f(\pm 0)$  returns  $-\infty$ .
- ▶  $\log2f(1)$  returns  $+0$ .
- ▶  $\log2f(x)$  returns NaN for  $x < 0$ .
- ▶  $\log2f(+\infty)$  returns  $+\infty$ .

### Description

Calculate the base 2 logarithm of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## \_\_device\_\_ float logbf (float x)

Calculate the floating point representation of the exponent of the input argument.

### Returns

- ▶  $\logbf \pm 0$  returns  $-\infty$
- ▶  $\logbf +\infty$  returns  $+\infty$

**Description**

Calculate the floating point representation of the exponent of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

**\_\_device\_\_ float logf (float x)**

Calculate the natural logarithm of the input argument.

**Returns**

- ▶  $\text{logf}(\pm 0)$  returns  $-\infty$ .
- ▶  $\text{logf}(1)$  returns  $+0$ .
- ▶  $\text{logf}(x)$  returns NaN for  $x < 0$ .
- ▶  $\text{logf}(+\infty)$  returns  $+\infty$ .

**Description**

Calculate the natural logarithm of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

**\_\_device\_\_ long int lrintf (float x)**

Round input to nearest integer value.

**Returns**

Returns rounded integer value.

**Description**

Round  $x$  to the nearest integer value, with halfway cases rounded towards zero. If the result is outside the range of the return type, the result is undefined.

**\_\_device\_\_ long int lroundf (float x)**

Round to nearest integer value.

**Returns**

Returns rounded integer value.

## Description

Round  $x$  to the nearest integer value, with halfway cases rounded away from zero. If the result is outside the range of the return type, the result is undefined.



This function may be slower than alternate rounding methods. See [lrintf\(\)](#).

## \_\_device\_\_ float modff (float x, float \*iptr)

Break down the input argument into fractional and integral parts.

### Returns

- ▶ `modff(  $\pm x$ , iptr)` returns a result with the same sign as  $x$ .
- ▶ `modff(  $\pm \infty$ , iptr)` returns  $\pm 0$  and stores  $\pm \infty$  in the object pointed to by `iptr`.
- ▶ `modff(NaN, iptr)` stores a NaN in the object pointed to by `iptr` and returns a NaN.

## Description

Break down the argument  $x$  into fractional and integral parts. The integral part is stored in the argument `iptr`. Fractional and integral parts are given the same sign as the argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## \_\_device\_\_ float nanf (const char \*tagp)

Returns "Not a Number" value.

### Returns

- ▶ `nanf(tagp)` returns NaN.

## Description

Return a representation of a quiet NaN. Argument `tagp` selects one of the possible representations.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float nearbyintf (float x)`

Round the input argument to the nearest integer.

### Returns

- ▶ `nearbyintf(  $\pm 0$  )` returns  $\pm 0$ .
- ▶ `nearbyintf(  $\pm \infty$  )` returns  $\pm \infty$ .

### Description

Round argument  $x$  to an integer value in single precision floating-point format.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float nextafterf (float x, float y)`

Return next representable single-precision floating-point value after argument.

### Returns

- ▶ `nextafterf(  $\pm \infty$ ,  $y$  )` returns  $\pm \infty$ .

### Description

Calculate the next representable single-precision floating-point value following  $x$  in the direction of  $y$ . For example, if  $y$  is greater than  $x$ , `nextafterf()` returns the smallest representable number greater than  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float norm3df (float a, float b, float c)`

Calculate the square root of the sum of squares of three coordinates of the argument.

### Returns

Returns the length of the 3D  $\sqrt{p.x^2 + p.y^2 + p.z^2}$ . If the correct value would overflow, returns  $+\infty$ . If the correct value would underflow, returns 0.



**Description**

Calculates the length of three dimensional vector  $p$  in euclidean space without undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float norm4df (float a, float b, float c, float d)`

Calculate the square root of the sum of squares of four coordinates of the argument.

**Returns**

Returns the length of the 4D vector  $\sqrt{p.x^2 + p.y^2 + p.z^2 + p.t^2}$ . If the correct value would overflow, returns  $+\infty$ . If the correct value would underflow, returns 0.

**Description**

Calculates the length of four dimensional vector  $p$  in euclidean space without undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float normcdf (float y)`

Calculate the standard normal cumulative distribution function.

**Returns**

- ▶ `normcdf(  $+\infty$  )` returns 1
- ▶ `normcdf(  $-\infty$  )` returns +0

**Description**

Calculate the cumulative distribution function of the standard normal distribution for input argument  $y$ ,  $\Phi(y)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## \_\_device\_\_ float normcdfinvf (float y)

Calculate the inverse of the standard normal cumulative distribution function.

### Returns

- ▶ `normcdfinvf(0)` returns  $-\infty$ .
- ▶ `normcdfinvf(1)` returns  $+\infty$ .
- ▶ `normcdfinvf(x)` returns NaN if  $x$  is not in the interval  $[0,1]$ .

### Description

Calculate the inverse of the standard normal cumulative distribution function for input argument  $y$ ,  $\Phi^{-1}(y)$ . The function is defined for input values in the interval  $(0, 1)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## \_\_device\_\_ float normf (int dim, const float \*a)

Calculate the square root of the sum of squares of any number of coordinates.

### Returns

Returns the length of the vector  $\sqrt{p.1^2 + p.2^2 + \dots + p.\text{dim}^2}$ . If the correct value would overflow, returns  $+\infty$ . If the correct value would underflow, returns 0.

### Description

Calculates the length of a vector  $p$ , dimension of which is passed as an argument without undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## \_\_device\_\_ float powf (float x, float y)

Calculate the value of first argument to the power of second argument.

### Returns

- ▶ `powf(±0, y)` returns  $\pm\infty$  for  $y$  an integer less than 0.
- ▶ `powf(±0, y)` returns  $\pm 0$  for  $y$  an odd integer greater than 0.
- ▶ `powf(±0, y)` returns +0 for  $y > 0$  and not an odd integer.

- ▶ `powf(-1,  $\pm\infty$ )` returns 1.
- ▶ `powf(+1,  $y$ )` returns 1 for any  $y$ , even a NaN.
- ▶ `powf( $x$ ,  $\pm 0$ )` returns 1 for any  $x$ , even a NaN.
- ▶ `powf( $x$ ,  $y$ )` returns a NaN for finite  $x < 0$  and finite non-integer  $y$ .
- ▶ `powf( $x$ ,  $-\infty$ )` returns  $+\infty$  for  $|x| < 1$ .
- ▶ `powf( $x$ ,  $-\infty$ )` returns +0 for  $|x| > 1$ .
- ▶ `powf( $x$ ,  $+\infty$ )` returns +0 for  $|x| < 1$ .
- ▶ `powf( $x$ ,  $+\infty$ )` returns  $+\infty$  for  $|x| > 1$ .
- ▶ `powf( $-\infty$ ,  $y$ )` returns -0 for  $y$  an odd integer less than 0.
- ▶ `powf( $-\infty$ ,  $y$ )` returns +0 for  $y < 0$  and not an odd integer.
- ▶ `powf( $-\infty$ ,  $y$ )` returns  $-\infty$  for  $y$  an odd integer greater than 0.
- ▶ `powf( $-\infty$ ,  $y$ )` returns  $+\infty$  for  $y > 0$  and not an odd integer.
- ▶ `powf( $+\infty$ ,  $y$ )` returns +0 for  $y < 0$ .
- ▶ `powf( $+\infty$ ,  $y$ )` returns  $+\infty$  for  $y > 0$ .

### Description

Calculate the value of  $x$  to the power of  $y$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float rcbrtf (float x)`

Calculate reciprocal cube root function.

### Returns

- ▶ `rcbrt( $\pm 0$ )` returns  $\pm\infty$ .
- ▶ `rcbrt( $\pm\infty$ )` returns  $\pm 0$ .

### Description

Calculate reciprocal cube root function of  $x$



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float remainderf (float x, float y)`

Compute single-precision floating-point remainder.

### Returns

- ▶ `remainderf(x, 0)` returns NaN.
- ▶ `remainderf(  $\pm\infty$ , y)` returns NaN.
- ▶ `remainderf(x,  $\pm\infty$ )` returns x for finite x.

### Description

Compute single-precision floating-point remainder  $r$  of dividing  $x$  by  $y$  for nonzero  $y$ . Thus  $r = x - ny$ . The value  $n$  is the integer value nearest  $\frac{x}{y}$ . In the case when  $|n - \frac{x}{y}| = \frac{1}{2}$ , the even  $n$  value is chosen.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float remquof (float x, float y, int *quo)`

Compute single-precision floating-point remainder and part of quotient.

### Returns

Returns the remainder.

- ▶ `remquof(x, 0, quo)` returns NaN.
- ▶ `remquof(  $\pm\infty$ , y, quo)` returns NaN.
- ▶ `remquof(x,  $\pm\infty$ , quo)` returns x.

### Description

Compute a double-precision floating-point remainder in the same way as the `remainderf()` function. Argument `quo` returns part of quotient upon division of  $x$  by  $y$ . Value `quo` has the same sign as  $\frac{x}{y}$  and may not be the exact quotient but agrees with the exact quotient in the low order 3 bits.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float rhypotf (float x, float y)`

Calculate one over the square root of the sum of squares of two arguments.

### Returns

Returns one over the length of the hypotenuse  $\frac{1}{\sqrt{x^2 + y^2}}$ . If the square root would overflow, returns 0. If the square root would underflow, returns  $+\infty$ .

### Description

Calculates one over the length of the hypotenuse of a right triangle whose two sides have lengths  $x$  and  $y$  without undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float rintf (float x)`

Round input to nearest integer value in floating-point.

### Returns

Returns rounded integer value.

### Description

Round  $x$  to the nearest integer value in floating-point format, with halfway cases rounded towards zero.

## `__device__ float rnorm3df (float a, float b, float c)`

Calculate one over the square root of the sum of squares of three coordinates of the argument.

### Returns

Returns one over the length of the 3D vector  $\frac{1}{\sqrt{p.x^2 + p.y^2 + p.z^2}}$ . If the square root would overflow, returns 0. If the square root would underflow, returns  $+\infty$ .

### Description

Calculates one over the length of three dimension vector  $p$  in euclidean space without undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float rnorm4df (float a, float b, float c, float d)`

Calculate one over the square root of the sum of squares of four coordinates of the argument.

### Returns

Returns one over the length of the 3D vector  $\frac{1}{\sqrt{p.x^2 + p.y^2 + p.z^2 + p.z^2}}$ . If the square root would overflow, returns 0. If the square root would underflow, returns  $+\infty$ .

### Description

Calculates one over the length of four dimension vector  $p$  in euclidean space without undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float rnormf (int dim, const float *a)`

Calculate the reciprocal of square root of the sum of squares of any number of coordinates.

### Returns

Returns one over the length of the vector  $\frac{1}{\sqrt{p.1^2 + p.2^2 + \dots + p.dim^2}}$ . If the square root would overflow, returns 0. If the square root would underflow, returns  $+\infty$ .

### Description

Calculates one over the length of vector  $p$ , dimension of which is passed as an argument, in euclidean space without undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float roundf (float x)`

Round to nearest integer value in floating-point.

### Returns

Returns rounded integer value.

### Description

Round  $x$  to the nearest integer value in floating-point format, with halfway cases rounded away from zero.



This function may be slower than alternate rounding methods. See [rintf\(\)](#).

## `__device__ float rsqrtf (float x)`

Calculate the reciprocal of the square root of the input argument.

### Returns

Returns  $1/\sqrt{x}$ .

- ▶ `rsqrtf( +∞ )` returns `+0`.
- ▶ `rsqrtf( ±0 )` returns  $\pm\infty$ .
- ▶ `rsqrtf(x)` returns NaN if  $x$  is less than 0.

### Description

Calculate the reciprocal of the nonnegative square root of  $x$ ,  $1/\sqrt{x}$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float scalblnf (float x, long int n)`

Scale floating-point input by integer power of two.

### Returns

Returns  $x * 2^n$ .

- ▶ `scalblnf( ±0 , n )` returns  $\pm 0$ .
- ▶ `scalblnf(x, 0)` returns  $x$ .

- ▶ `scalbnf(  $\pm\infty$ , n)` returns  $\pm\infty$ .

### Description

Scale  $x$  by  $2^n$  by efficient manipulation of the floating-point exponent.

## `__device__ float scalbnf (float x, int n)`

Scale floating-point input by integer power of two.

### Returns

Returns  $x * 2^n$ .

- ▶ `scalbnf(  $\pm 0$ , n)` returns  $\pm 0$ .
- ▶ `scalbnf(x, 0)` returns  $x$ .
- ▶ `scalbnf(  $\pm\infty$ , n)` returns  $\pm\infty$ .

### Description

Scale  $x$  by  $2^n$  by efficient manipulation of the floating-point exponent.

## `__device__ __RETURN_TYPE signbit (float a)`

Return the sign bit of the input.

### Returns

Reports the sign bit of all values including infinities, zeros, and NaNs.

- ▶ With Visual Studio 2013 host compiler: `__RETURN_TYPE` is 'bool'. Returns true if and only if  $a$  is negative.
- ▶ With other host compilers: `__RETURN_TYPE` is 'int'. Returns a nonzero value if and only if  $a$  is negative.

### Description

Determine whether the floating-point value  $a$  is negative.

## `__device__ void sincosf (float x, float *sptr, float *cptr)`

Calculate the sine and cosine of the first input argument.

### Returns

- ▶ none



## Description

Calculate the sine and cosine of the first input argument  $x$  (measured in radians). The results for sine and cosine are written into the second argument, `sptr`, and, respectively, third argument, `cptr`.

### See also:

[sinf\(\)](#) and [cosf\(\)](#).



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- This function is affected by the `--use_fast_math` compiler flag. See the CUDA C Programming Guide, Appendix D.2, Table 8 for a complete list of functions affected.

## `__device__ void sincospif (float x, float *sptr, float *cptr)`

Calculate the sine and cosine of the first input argument  $x \times \pi$ .

### Returns

- none

## Description

Calculate the sine and cosine of the first input argument,  $x$  (measured in radians),  $x \times \pi$ . The results for sine and cosine are written into the second argument, `sptr`, and, respectively, third argument, `cptr`.

### See also:

[sinpif\(\)](#) and [cospif\(\)](#).



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float sinf (float x)`

Calculate the sine of the input argument.

### Returns

- `sinf(  $\pm 0$  )` returns  $\pm 0$ .

- ▶  $\text{sinf}(\pm\infty)$  returns NaN.

### Description

Calculate the sine of the input argument  $x$  (measured in radians).



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This function is affected by the `--use_fast_math` compiler flag. See the CUDA C Programming Guide, Appendix D.2, Table 8 for a complete list of functions affected.

## `__device__ float sinh (float x)`

Calculate the hyperbolic sine of the input argument.

### Returns

- ▶  $\text{sinh}(\pm 0)$  returns  $\pm 0$ .
- ▶  $\text{sinh}(\pm\infty)$  returns NaN.

### Description

Calculate the hyperbolic sine of the input argument  $x$ .



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float sinpif (float x)`

Calculate the sine of the input argument  $x \times \pi$ .

### Returns

- ▶  $\text{sinpif}(\pm 0)$  returns  $\pm 0$ .
- ▶  $\text{sinpif}(\pm\infty)$  returns NaN.

### Description

Calculate the sine of  $x \times \pi$  (measured in radians), where  $x$  is the input argument.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float sqrtf (float x)`

Calculate the square root of the input argument.

### Returns

Returns  $\sqrt{x}$ .

- ▶ `sqrtf( ± 0 )` returns  $\pm 0$ .
- ▶ `sqrtf( + ∞ )` returns  $+\infty$ .
- ▶ `sqrtf(x)` returns NaN if  $x$  is less than 0.

### Description

Calculate the nonnegative square root of  $x$ ,  $\sqrt{x}$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float tanf (float x)`

Calculate the tangent of the input argument.

### Returns

- ▶ `tanf( ± 0 )` returns  $\pm 0$ .
- ▶ `tanf( ± ∞ )` returns NaN.

### Description

Calculate the tangent of the input argument  $x$  (measured in radians).



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This function is affected by the `--use_fast_math` compiler flag. See the CUDA C Programming Guide, Appendix D.2, Table 8 for a complete list of functions affected.

## `__device__ float tanhf (float x)`

Calculate the hyperbolic tangent of the input argument.

### Returns

- ▶ `tanhf( ± 0 )` returns  $\pm 0$ .

### Description

Calculate the hyperbolic tangent of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float tgammaf (float x)`

Calculate the gamma function of the input argument.

### Returns

- ▶ `tgammaf( ± 0 )` returns  $\pm \infty$ .
- ▶ `tgammaf(2)` returns  $+0$ .
- ▶ `tgammaf(x)` returns  $\pm \infty$  if the correctly calculated value is outside the single floating point range.
- ▶ `tgammaf(x)` returns NaN if  $x < 0$ .
- ▶ `tgammaf( - ∞ )` returns NaN.
- ▶ `tgammaf( + ∞ )` returns  $+\infty$ .

### Description

Calculate the gamma function of the input argument  $x$ , namely the value of  $\int_0^{\infty} e^{-t} t^{x-1} dt$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float truncf (float x)`

Truncate input argument to the integral part.

### Returns

Returns truncated integer value.

**Description**

Round  $x$  to the nearest integer value that does not exceed  $x$  in magnitude.

**\_\_device\_\_ float y0f (float x)**

Calculate the value of the Bessel function of the second kind of order 0 for the input argument.

**Returns**

Returns the value of the Bessel function of the second kind of order 0.

- ▶  $y0f(0)$  returns  $-\infty$ .
- ▶  $y0f(x)$  returns NaN for  $x < 0$ .
- ▶  $y0f(+\infty)$  returns +0.
- ▶  $y0f(\text{NaN})$  returns NaN.

**Description**

Calculate the value of the Bessel function of the second kind of order 0 for the input argument  $x$ ,  $Y_0(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

**\_\_device\_\_ float y1f (float x)**

Calculate the value of the Bessel function of the second kind of order 1 for the input argument.

**Returns**

Returns the value of the Bessel function of the second kind of order 1.

- ▶  $y1f(0)$  returns  $-\infty$ .
- ▶  $y1f(x)$  returns NaN for  $x < 0$ .
- ▶  $y1f(+\infty)$  returns +0.
- ▶  $y1f(\text{NaN})$  returns NaN.

**Description**

Calculate the value of the Bessel function of the second kind of order 1 for the input argument  $x$ ,  $Y_1(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float ynf (int n, float x)`

Calculate the value of the Bessel function of the second kind of order  $n$  for the input argument.

### Returns

Returns the value of the Bessel function of the second kind of order  $n$ .

- ▶ `ynf(n, x)` returns NaN for  $n < 0$ .
- ▶ `ynf(n, 0)` returns  $-\infty$ .
- ▶ `ynf(n, x)` returns NaN for  $x < 0$ .
- ▶ `ynf(n,  $+\infty$ )` returns  $+0$ .
- ▶ `ynf(n, NaN)` returns NaN.

### Description

Calculate the value of the Bessel function of the second kind of order  $n$  for the input argument  $x$ ,  $Y_n(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## 1.3. Double Precision Mathematical Functions

This section describes double precision mathematical functions.

### `__device__ double acos (double x)`

Calculate the arc cosine of the input argument.

### Returns

Result will be in radians, in the interval  $[0, \pi]$  for  $x$  inside  $[-1, +1]$ .

- ▶ `acos(1)` returns  $+0$ .
- ▶ `acos(x)` returns NaN for  $x$  outside  $[-1, +1]$ .

**Description**

Calculate the principal value of the arc cosine of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

**\_\_device\_\_ double acosh (double x)**

Calculate the nonnegative arc hyperbolic cosine of the input argument.

**Returns**

Result will be in the interval  $[0, +\infty]$ .

- ▶ `acosh(1)` returns 0.
- ▶ `acosh(x)` returns NaN for  $x$  in the interval  $[-\infty, 1)$ .

**Description**

Calculate the nonnegative arc hyperbolic cosine of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

**\_\_device\_\_ double asin (double x)**

Calculate the arc sine of the input argument.

**Returns**

Result will be in radians, in the interval  $[-\pi/2, +\pi/2]$  for  $x$  inside  $[-1, +1]$ .

- ▶ `asin(0)` returns +0.
- ▶ `asin(x)` returns NaN for  $x$  outside  $[-1, +1]$ .

**Description**

Calculate the principal value of the arc sine of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double asinh (double x)`

Calculate the arc hyperbolic sine of the input argument.

### Returns

- `asinh(0)` returns 1.

### Description

Calculate the arc hyperbolic sine of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double atan (double x)`

Calculate the arc tangent of the input argument.

### Returns

Result will be in radians, in the interval  $[-\pi/2, +\pi/2]$ .

- `atan(0)` returns +0.

### Description

Calculate the principal value of the arc tangent of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double atan2 (double y, double x)`

Calculate the arc tangent of the ratio of first and second input arguments.

### Returns

Result will be in radians, in the interval  $[-\pi, +\pi]$ .

- `atan2(0, 1)` returns +0.

### Description

Calculate the principal value of the arc tangent of the ratio of first and second input arguments  $y/x$ . The quadrant of the result is determined by the signs of inputs  $y$  and  $x$ .





For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double atanh (double x)`

Calculate the arc hyperbolic tangent of the input argument.

### Returns

- ▶ `atanh( ± 0 )` returns  $\pm 0$ .
- ▶ `atanh( ± 1 )` returns  $\pm \infty$ .
- ▶ `atanh(x)` returns NaN for  $x$  outside interval  $[-1, 1]$ .

### Description

Calculate the arc hyperbolic tangent of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double cbrt (double x)`

Calculate the cube root of the input argument.

### Returns

Returns  $x^{1/3}$ .

- ▶ `cbrt( ± 0 )` returns  $\pm 0$ .
- ▶ `cbrt( ± ∞ )` returns  $\pm \infty$ .

### Description

Calculate the cube root of  $x$ ,  $x^{1/3}$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ __CUDA_MATH_CRTIMP double ceil (double x)`

Calculate ceiling of the input argument.

### Returns

Returns  $\lceil x \rceil$  expressed as a floating-point number.

- ▶ `ceil(  $\pm 0$  )` returns  $\pm 0$ .
- ▶ `ceil(  $\pm \infty$  )` returns  $\pm \infty$ .

### Description

Compute the smallest integer value not less than  $x$ .

## `__device__ double copysign (double x, double y)`

Create value with given magnitude, copying sign of second value.

### Returns

Returns a value with the magnitude of  $x$  and the sign of  $y$ .

### Description

Create a floating-point value with the magnitude  $x$  and the sign of  $y$ .

## `__device__ double cos (double x)`

Calculate the cosine of the input argument.

### Returns

- ▶ `cos(  $\pm 0$  )` returns 1.
- ▶ `cos(  $\pm \infty$  )` returns NaN.

### Description

Calculate the cosine of the input argument  $x$  (measured in radians).



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double cosh (double x)`

Calculate the hyperbolic cosine of the input argument.

### Returns

- ▶ `cosh(0)` returns 1.
- ▶ `cosh(  $\pm \infty$  )` returns  $+\infty$ .

### Description

Calculate the hyperbolic cosine of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double cospi (double x)`

Calculate the cosine of the input argument  $x \times \pi$ .

### Returns

- ▶ `cospi(  $\pm 0$  )` returns 1.
- ▶ `cospi(  $\pm \infty$  )` returns NaN.

### Description

Calculate the cosine of  $x \times \pi$  (measured in radians), where  $x$  is the input argument.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ __CUDA_MATH_CRTIMP double cyl_bessel_i0 (double x)`

Calculate the value of the regular modified cylindrical Bessel function of order 0 for the input argument.

### Returns

Returns the value of the regular modified cylindrical Bessel function of order 0.

**Description**

Calculate the value of the regular modified cylindrical Bessel function of order 0 for the input argument  $x$ ,  $I_0(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ __CUDA_MATH_CRTIMP double cyl_bessel_i1(double x)`

Calculate the value of the regular modified cylindrical Bessel function of order 1 for the input argument.

**Returns**

Returns the value of the regular modified cylindrical Bessel function of order 1.

**Description**

Calculate the value of the regular modified cylindrical Bessel function of order 1 for the input argument  $x$ ,  $I_1(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double erf(double x)`

Calculate the error function of the input argument.

**Returns**

- ▶ `erf( ± 0 )` returns  $\pm 0$ .
- ▶ `erf( ± ∞ )` returns  $\pm 1$ .

**Description**

Calculate the value of the error function for the input argument  $x$ ,  $\frac{2}{\sqrt{\pi}} \int_0^x e^{-t^2} dt$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double erfc (double x)`

Calculate the complementary error function of the input argument.

### Returns

- ▶ `erfc( -  $\infty$  )` returns 2.
- ▶ `erfc( +  $\infty$  )` returns +0.

### Description

Calculate the complementary error function of the input argument  $x$ ,  $1 - \text{erf}(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double erfcinv (double y)`

Calculate the inverse complementary error function of the input argument.

### Returns

- ▶ `erfcinv(0)` returns  $+\infty$ .
- ▶ `erfcinv(2)` returns  $-\infty$ .

### Description

Calculate the inverse complementary error function of the input argument  $y$ , for  $y$  in the interval  $[0, 2]$ . The inverse complementary error function find the value  $x$  that satisfies the equation  $y = \text{erfc}(x)$ , for  $0 \leq y \leq 2$ , and  $-\infty \leq x \leq \infty$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double erfcx (double x)`

Calculate the scaled complementary error function of the input argument.

### Returns

- ▶ `erfcx( -  $\infty$  )` returns  $+\infty$
- ▶ `erfcx( +  $\infty$  )` returns +0
- ▶ `erfcx(x)` returns  $+\infty$  if the correctly calculated value is outside the double floating point range.

**Description**

Calculate the scaled complementary error function of the input argument  $x$ ,  $e^{x^2} \cdot \operatorname{erfc}(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

**\_\_device\_\_ double erfinv (double y)**

Calculate the inverse error function of the input argument.

**Returns**

- ▶  $\operatorname{erfinv}(1)$  returns  $+\infty$ .
- ▶  $\operatorname{erfinv}(-1)$  returns  $-\infty$ .

**Description**

Calculate the inverse error function of the input argument  $y$ , for  $y$  in the interval  $[-1, 1]$ . The inverse error function finds the value  $x$  that satisfies the equation  $y = \operatorname{erf}(x)$ , for  $-1 \leq y \leq 1$ , and  $-\infty \leq x \leq \infty$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

**\_\_device\_\_ double exp (double x)**

Calculate the base  $e$  exponential of the input argument.

**Returns**

Returns  $e^x$ .

**Description**

Calculate the base  $e$  exponential of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double exp10 (double x)`

Calculate the base 10 exponential of the input argument.

### Returns

Returns  $10^x$ .

### Description

Calculate the base 10 exponential of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double exp2 (double x)`

Calculate the base 2 exponential of the input argument.

### Returns

Returns  $2^x$ .

### Description

Calculate the base 2 exponential of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double expm1 (double x)`

Calculate the base  $e$  exponential of the input argument, minus 1.

### Returns

Returns  $e^x - 1$ .

### Description

Calculate the base  $e$  exponential of the input argument  $x$ , minus 1.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double fabs (double x)`

Calculate the absolute value of the input argument.

### Returns

Returns the absolute value of the input argument.

- ▶ `fabs(  $\pm \infty$  )` returns  $+\infty$ .
- ▶ `fabs(  $\pm 0$  )` returns 0.

### Description

Calculate the absolute value of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double fdim (double x, double y)`

Compute the positive difference between  $x$  and  $y$ .

### Returns

Returns the positive difference between  $x$  and  $y$ .

- ▶ `fdim( $x$ ,  $y$ )` returns  $x - y$  if  $x > y$ .
- ▶ `fdim( $x$ ,  $y$ )` returns  $+0$  if  $x \leq y$ .

### Description

Compute the positive difference between  $x$  and  $y$ . The positive difference is  $x - y$  when  $x > y$  and  $+0$  otherwise.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ __CUDA_MATH_CRTIMP double floor (double x)`

Calculate the largest integer less than or equal to  $x$ .

### Returns

Returns  $\log_e(1+x)$  expressed as a floating-point number.



- ▶ `floor(  $\pm\infty$  )` returns  $\pm\infty$ .
- ▶ `floor(  $\pm 0$  )` returns  $\pm 0$ .

### Description

Calculates the largest integer value which is less than or equal to  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double fma (double x, double y, double z)`

Compute  $x \times y + z$  as a single operation.

### Returns

Returns the rounded value of  $x \times y + z$  as a single operation.

- ▶ `fma(  $\pm\infty$ ,  $\pm 0$ ,  $z$  )` returns NaN.
- ▶ `fma(  $\pm 0$ ,  $\pm\infty$ ,  $z$  )` returns NaN.
- ▶ `fma( $x$ ,  $y$ ,  $-\infty$ )` returns NaN if  $x \times y$  is an exact  $+\infty$ .
- ▶ `fma( $x$ ,  $y$ ,  $+\infty$ )` returns NaN if  $x \times y$  is an exact  $-\infty$ .

### Description

Compute the value of  $x \times y + z$  as a single ternary operation. After computing the value to infinite precision, the value is rounded once.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double fmax (double, double)`

Determine the maximum numeric value of the arguments.

### Returns

Returns the maximum numeric values of the arguments  $x$  and  $y$ .

- ▶ If both arguments are NaN, returns NaN.
- ▶ If one argument is NaN, returns the numeric argument.

## Description

Determines the maximum numeric value of the arguments  $x$  and  $y$ . Treats NaN arguments as missing data. If one argument is a NaN and the other is legitimate numeric value, the numeric value is chosen.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double fmin (double x, double y)`

Determine the minimum numeric value of the arguments.

## Returns

Returns the minimum numeric values of the arguments  $x$  and  $y$ .

- ▶ If both arguments are NaN, returns NaN.
- ▶ If one argument is NaN, returns the numeric argument.

## Description

Determines the minimum numeric value of the arguments  $x$  and  $y$ . Treats NaN arguments as missing data. If one argument is a NaN and the other is legitimate numeric value, the numeric value is chosen.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double fmod (double x, double y)`

Calculate the floating-point remainder of  $x / y$ .

## Returns

- ▶ Returns the floating point remainder of  $x / y$ .
- ▶ `fmod(  $\pm 0$ ,  $y$  )` returns  $\pm 0$  if  $y$  is not zero.
- ▶ `fmod( $x$ ,  $y$ )` returns NaN and raised an invalid floating point exception if  $x$  is  $\pm \infty$  or  $y$  is zero.
- ▶ `fmod( $x$ ,  $y$ )` returns zero if  $y$  is zero or the result would overflow.
- ▶ `fmod( $x$ ,  $\pm \infty$ )` returns  $x$  if  $x$  is finite.
- ▶ `fmod( $x$ , 0)` returns NaN.

## Description

Calculate the floating-point remainder of  $x / y$ . The absolute value of the computed value is always less than  $y$ 's absolute value and will have the same sign as  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ __CUDA_MATH_CRTIMP double frexp (double x, int *nptr)`

Extract mantissa and exponent of a floating-point value.

## Returns

Returns the fractional component  $m$ .

- ▶ `frexp(0, nptr)` returns 0 for the fractional component and zero for the integer component.
- ▶ `frexp( $\pm 0$ , nptr)` returns  $\pm 0$  and stores zero in the location pointed to by `nptr`.
- ▶ `frexp( $\pm \infty$ , nptr)` returns  $\pm \infty$  and stores an unspecified value in the location to which `nptr` points.
- ▶ `frexp(NaN, y)` returns a NaN and stores an unspecified value in the location to which `nptr` points.

## Description

Decompose the floating-point value  $x$  into a component  $m$  for the normalized fraction element and another term  $n$  for the exponent. The absolute value of  $m$  will be greater than or equal to 0.5 and less than 1.0 or it will be equal to 0;  $x = m \cdot 2^n$ . The integer exponent  $n$  will be stored in the location to which `nptr` points.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ __ACRTIMP double hypot (double x, double y)`

Calculate the square root of the sum of squares of two arguments.

## Returns

Returns the length of the hypotenuse  $\sqrt{x^2 + y^2}$ . If the correct value would overflow, returns  $+\infty$ . If the correct value would underflow, returns 0.

## Description

Calculate the length of the hypotenuse of a right triangle whose two sides have lengths  $x$  and  $y$  without undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ int ilogb (double x)`

Compute the unbiased integer exponent of the argument.

### Returns

- ▶ If successful, returns the unbiased exponent of the argument.
- ▶ `ilogb(0)` returns `INT_MIN`.
- ▶ `ilogb(NaN)` returns `NaN`.
- ▶ `ilogb(x)` returns `INT_MAX` if  $x$  is  $\infty$  or the correct value is greater than `INT_MAX`.
- ▶ `ilogb(x)` return `INT_MIN` if the correct value is less than `INT_MIN`.

## Description

Calculates the unbiased integer exponent of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ __RETURN_TYPE isfinite (double a)`

Determine whether argument is finite.

### Returns

- ▶ With Visual Studio 2013 host compiler: `__RETURN_TYPE` is 'bool'. Returns true if and only if  $a$  is a finite value.
- ▶ With other host compilers: `__RETURN_TYPE` is 'int'. Returns a nonzero value if and only if  $a$  is a finite value.

## Description

Determine whether the floating-point value  $a$  is a finite value (zero, subnormal, or normal and not infinity or NaN).

## `__device__ __RETURN_TYPE isinf (double a)`

Determine whether argument is infinite.

### Returns

- ▶ With Visual Studio 2013 host compiler: Returns true if and only if *a* is a infinite value.
- ▶ With other host compilers: Returns a nonzero value if and only if *a* is a infinite value.

### Description

Determine whether the floating-point value *a* is an infinite value (positive or negative).

## `__device__ __RETURN_TYPE isnan (double a)`

Determine whether argument is a NaN.

### Returns

- ▶ With Visual Studio 2013 host compiler: `__RETURN_TYPE` is 'bool'. Returns true if and only if *a* is a NaN value.
- ▶ With other host compilers: `__RETURN_TYPE` is 'int'. Returns a nonzero value if and only if *a* is a NaN value.

### Description

Determine whether the floating-point value *a* is a NaN.

## `__device__ __CUDA_MATH_CRTIMP double j0 (double x)`

Calculate the value of the Bessel function of the first kind of order 0 for the input argument.

### Returns

Returns the value of the Bessel function of the first kind of order 0.

- ▶  $j_0(\pm\infty)$  returns +0.
- ▶  $j_0(\text{NaN})$  returns NaN.

### Description

Calculate the value of the Bessel function of the first kind of order 0 for the input argument *x*,  $J_0(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ __CUDA_MATH_CRTIMP double j1 (double x)`

Calculate the value of the Bessel function of the first kind of order 1 for the input argument.

### Returns

Returns the value of the Bessel function of the first kind of order 1.

- ▶ `j1( ± 0 )` returns  $\pm 0$ .
- ▶ `j1( ± ∞ )` returns  $+0$ .
- ▶ `j1(NaN)` returns NaN.

### Description

Calculate the value of the Bessel function of the first kind of order 1 for the input argument  $x$ ,  $J_1(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ __CUDA_MATH_CRTIMP double jn (int n, double x)`

Calculate the value of the Bessel function of the first kind of order  $n$  for the input argument.

### Returns

Returns the value of the Bessel function of the first kind of order  $n$ .

- ▶ `jn(n, NaN)` returns NaN.
- ▶ `jn(n, x)` returns NaN for  $n < 0$ .
- ▶ `jn(n, + ∞ )` returns  $+0$ .

### Description

Calculate the value of the Bessel function of the first kind of order  $n$  for the input argument  $x$ ,  $J_n(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ __CUDA_MATH_CRTIMP double ldexp (double x, int exp)`

Calculate the value of  $x \cdot 2^{\text{exp}}$ .

### Returns

- `ldexp(x)` returns  $\pm \infty$  if the correctly calculated value is outside the double floating point range.

### Description

Calculate the value of  $x \cdot 2^{\text{exp}}$  of the input arguments `x` and `exp`.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double lgamma (double x)`

Calculate the natural logarithm of the absolute value of the gamma function of the input argument.

### Returns

- `lgamma(1)` returns +0.
- `lgamma(2)` returns +0.
- `lgamma(x)` returns  $\pm \infty$  if the correctly calculated value is outside the double floating point range.
- `lgamma(x)` returns  $+\infty$  if  $x \leq 0$ .
- `lgamma(  $-\infty$  )` returns  $-\infty$ .
- `lgamma(  $+\infty$  )` returns  $+\infty$ .

### Description

Calculate the natural logarithm of the absolute value of the gamma function of the input argument `x`, namely the value of  $\log_e \left| \int_0^\infty e^{-t} t^{x-1} dt \right|$



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ long long int llrint (double x)`

Round input to nearest integer value.

### Returns

Returns rounded integer value.

### Description

Round  $x$  to the nearest integer value, with halfway cases rounded towards zero. If the result is outside the range of the return type, the result is undefined.

## `__device__ long long int llround (double x)`

Round to nearest integer value.

### Returns

Returns rounded integer value.

### Description

Round  $x$  to the nearest integer value, with halfway cases rounded away from zero. If the result is outside the range of the return type, the result is undefined.



This function may be slower than alternate rounding methods. See [llrint\(\)](#).

## `__device__ double log (double x)`

Calculate the base  $e$  logarithm of the input argument.

### Returns

- ▶  $\log(\pm 0)$  returns  $-\infty$ .
- ▶  $\log(1)$  returns  $+0$ .
- ▶  $\log(x)$  returns NaN for  $x < 0$ .
- ▶  $\log(+\infty)$  returns  $+\infty$ .

### Description

Calculate the base  $e$  logarithm of the input argument  $x$ .





For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double log10 (double x)`

Calculate the base 10 logarithm of the input argument.

### Returns

- ▶  $\log_{10}(\pm 0)$  returns  $-\infty$ .
- ▶  $\log_{10}(1)$  returns  $+0$ .
- ▶  $\log_{10}(x)$  returns NaN for  $x < 0$ .
- ▶  $\log_{10}(+\infty)$  returns  $+\infty$ .

### Description

Calculate the base 10 logarithm of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double log1p (double x)`

Calculate the value of  $\log_e(1+x)$ .

### Returns

- ▶  $\log_{1p}(\pm 0)$  returns  $-\infty$ .
- ▶  $\log_{1p}(-1)$  returns  $+0$ .
- ▶  $\log_{1p}(x)$  returns NaN for  $x < -1$ .
- ▶  $\log_{1p}(+\infty)$  returns  $+\infty$ .

### Description

Calculate the value of  $\log_e(1+x)$  of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## \_\_device\_\_ double log2 (double x)

Calculate the base 2 logarithm of the input argument.

### Returns

- ▶  $\log_2(\pm 0)$  returns  $-\infty$ .
- ▶  $\log_2(1)$  returns  $+0$ .
- ▶  $\log_2(x)$  returns NaN for  $x < 0$ .
- ▶  $\log_2(+\infty)$  returns  $+\infty$ .

### Description

Calculate the base 2 logarithm of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## \_\_device\_\_ double logb (double x)

Calculate the floating point representation of the exponent of the input argument.

### Returns

- ▶  $\log_b \pm 0$  returns  $-\infty$
- ▶  $\log_b \pm \infty$  returns  $+\infty$

### Description

Calculate the floating point representation of the exponent of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## \_\_device\_\_ long int lrint (double x)

Round input to nearest integer value.

### Returns

Returns rounded integer value.

**Description**

Round  $x$  to the nearest integer value, with halfway cases rounded towards zero. If the result is outside the range of the return type, the result is undefined.

**\_\_device\_\_ long int lround (double x)**

Round to nearest integer value.

**Returns**

Returns rounded integer value.

**Description**

Round  $x$  to the nearest integer value, with halfway cases rounded away from zero. If the result is outside the range of the return type, the result is undefined.



This function may be slower than alternate rounding methods. See [lrint\(\)](#).

**\_\_device\_\_ \_\_CUDA\_MATH\_CRTIMP double modf (double x, double \*iptr)**

Break down the input argument into fractional and integral parts.

**Returns**

- ▶ `modf(  $\pm x$ , iptr)` returns a result with the same sign as  $x$ .
- ▶ `modf(  $\pm \infty$ , iptr)` returns  $\pm 0$  and stores  $\pm \infty$  in the object pointed to by `iptr`.
- ▶ `modf(NaN, iptr)` stores a NaN in the object pointed to by `iptr` and returns a NaN.

**Description**

Break down the argument  $x$  into fractional and integral parts. The integral part is stored in the argument `iptr`. Fractional and integral parts are given the same sign as the argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double nan (const char *tagp)`

Returns "Not a Number" value.

### Returns

- ▶ `nan(tagp)` returns NaN.

### Description

Return a representation of a quiet NaN. Argument `tagp` selects one of the possible representations.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double nearbyint (double x)`

Round the input argument to the nearest integer.

### Returns

- ▶ `nearbyint( ±0 )` returns  $\pm 0$ .
- ▶ `nearbyint( ±∞ )` returns  $\pm \infty$ .

### Description

Round argument  $x$  to an integer value in double precision floating-point format.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double nextafter (double x, double y)`

Return next representable double-precision floating-point value after argument.

### Returns

- ▶ `nextafter( ±∞ , y)` returns  $\pm \infty$ .

### Description

Calculate the next representable double-precision floating-point value following  $x$  in the direction of  $y$ . For example, if  $y$  is greater than  $x$ , `nextafter()` returns the smallest representable number greater than  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double norm (int dim, const double *t)`

Calculate the square root of the sum of squares of any number of coordinates.

### Returns

Returns the length of the dim-D vector  $\sqrt{p.1^2 + p.2^2 + \dots + p.\text{dim}^2}$ . If the correct value would overflow, returns  $+\infty$ . If the correct value would underflow, returns 0. If two of the input arguments is 0, returns remaining argument

### Description

Calculate the length of a vector  $p$ , dimension of which is passed as an argument without undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ __CUDA_MATH_CRTIMP double norm3d (double a, double b, double c)`

Calculate the square root of the sum of squares of three coordinates of the argument.

### Returns

Returns the length of 3D vector  $\sqrt{p.x^2 + p.y^2 + p.z^2}$ . If the correct value would overflow, returns  $+\infty$ . If the correct value would underflow, returns 0.

### Description

Calculate the length of three dimensional vector  $p$  in euclidean space without undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ __CUDA_MATH_CRTIMP double norm4d` (double a, double b, double c, double d)

Calculate the square root of the sum of squares of four coordinates of the argument.

### Returns

Returns the length of 4D vector  $\sqrt{p.x^2 + p.y^2 + p.z^2 + p.t^2}$ . If the correct value would overflow, returns  $+\infty$ . If the correct value would underflow, returns 0.

### Description

Calculate the length of four dimensional vector  $p$  in euclidean space without undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double normcdf` (double y)

Calculate the standard normal cumulative distribution function.

### Returns

- ▶ `normcdf(  $+\infty$  )` returns 1
- ▶ `normcdf(  $-\infty$  )` returns +0

### Description

Calculate the cumulative distribution function of the standard normal distribution for input argument  $y$ ,  $\Phi(y)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double normcdfinv` (double y)

Calculate the inverse of the standard normal cumulative distribution function.

### Returns

- ▶ `normcdfinv(0)` returns  $-\infty$ .
- ▶ `normcdfinv(1)` returns  $+\infty$ .
- ▶ `normcdfinv(x)` returns NaN if  $x$  is not in the interval  $[0,1]$ .

## Description

Calculate the inverse of the standard normal cumulative distribution function for input argument  $y$ ,  $\Phi^{-1}(y)$ . The function is defined for input values in the interval  $(0, 1)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double pow (double x, double y)`

Calculate the value of first argument to the power of second argument.

## Returns

- ▶ `pow( ±0, y)` returns  $\pm\infty$  for  $y$  an integer less than 0.
- ▶ `pow( ±0, y)` returns  $\pm 0$  for  $y$  an odd integer greater than 0.
- ▶ `pow( ±0, y)` returns +0 for  $y > 0$  and not an odd integer.
- ▶ `pow(-1, ±∞)` returns 1.
- ▶ `pow(+1, y)` returns 1 for any  $y$ , even a NaN.
- ▶ `pow(x, ±0)` returns 1 for any  $x$ , even a NaN.
- ▶ `pow(x, y)` returns a NaN for finite  $x < 0$  and finite non-integer  $y$ .
- ▶ `pow(x, -∞)` returns  $+\infty$  for  $|x| < 1$ .
- ▶ `pow(x, -∞)` returns +0 for  $|x| > 1$ .
- ▶ `pow(x, +∞)` returns +0 for  $|x| < 1$ .
- ▶ `pow(x, +∞)` returns  $+\infty$  for  $|x| > 1$ .
- ▶ `pow(-∞, y)` returns -0 for  $y$  an odd integer less than 0.
- ▶ `pow(-∞, y)` returns +0 for  $y < 0$  and not an odd integer.
- ▶ `pow(-∞, y)` returns  $-\infty$  for  $y$  an odd integer greater than 0.
- ▶ `pow(-∞, y)` returns  $+\infty$  for  $y > 0$  and not an odd integer.
- ▶ `pow(+∞, y)` returns +0 for  $y < 0$ .
- ▶ `pow(+∞, y)` returns  $+\infty$  for  $y > 0$ .

## Description

Calculate the value of  $x$  to the power of  $y$



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double rcbrt (double x)`

Calculate reciprocal cube root function.

### Returns

- ▶ `rcbrt(  $\pm 0$  )` returns  $\pm \infty$ .
- ▶ `rcbrt(  $\pm \infty$  )` returns  $\pm 0$ .

### Description

Calculate reciprocal cube root function of  $x$



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double remainder (double x, double y)`

Compute double-precision floating-point remainder.

### Returns

- ▶ `remainder( $x$ , 0)` returns NaN.
- ▶ `remainder(  $\pm \infty$ ,  $y$  )` returns NaN.
- ▶ `remainder( $x$ ,  $\pm \infty$  )` returns  $x$  for finite  $x$ .

### Description

Compute double-precision floating-point remainder  $r$  of dividing  $x$  by  $y$  for nonzero  $y$ .

Thus  $r = x - ny$ . The value  $n$  is the integer value nearest  $\frac{x}{y}$ . In the case when  $|n - \frac{x}{y}| = \frac{1}{2}$ , the even  $n$  value is chosen.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double remquo (double x, double y, int *quo)`

Compute double-precision floating-point remainder and part of quotient.

### Returns

Returns the remainder.



- ▶ `remquo(x, 0, quo)` returns NaN.
- ▶ `remquo( $\pm\infty$ , y, quo)` returns NaN.
- ▶ `remquo(x,  $\pm\infty$ , quo)` returns x.

### Description

Compute a double-precision floating-point remainder in the same way as the `remainder()` function. Argument `quo` returns part of quotient upon division of `x` by `y`. Value `quo` has the same sign as  $\frac{x}{y}$  and may not be the exact quotient but agrees with the exact quotient in the low order 3 bits.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double rhypot (double x, double y)`

Calculate one over the square root of the sum of squares of two arguments.

### Returns

Returns one over the length of the hypotenuse  $\frac{1}{\sqrt{x^2 + y^2}}$ . If the square root would overflow, returns 0. If the square root would underflow, returns  $+\infty$ .

### Description

Calculate one over the length of the hypotenuse of a right triangle whose two sides have lengths `x` and `y` without undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double rint (double x)`

Round to nearest integer value in floating-point.

### Returns

Returns rounded integer value.

### Description

Round `x` to the nearest integer value in floating-point format, with halfway cases rounded to the nearest even integer value.

## `__device__ double rnorm (int dim, const double *t)`

Calculate the reciprocal of square root of the sum of squares of any number of coordinates.

### Returns

Returns one over the length of the vector  $\frac{1}{\sqrt{p.1^2 + p.2^2 + \dots + p.\text{dim}^2}}$ . If the square root would overflow, returns 0. If the square root would underflow, returns  $+\infty$ .

### Description

Calculates one over the length of vector  $p$ , dimension of which is passed as an argument, in euclidean space without undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double rnorm3d (double a, double b, double c)`

Calculate one over the square root of the sum of squares of three coordinates of the argument.

### Returns

Returns one over the length of the 3D vector  $\frac{1}{\sqrt{p.x^2 + p.y^2 + p.z^2}}$ . If the square root would overflow, returns 0. If the square root would underflow, returns  $+\infty$ .

### Description

Calculate one over the length of three dimensional vector  $p$  in euclidean space undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double rnorm4d (double a, double b, double c, double d)`

Calculate one over the square root of the sum of squares of four coordinates of the argument.

### Returns

Returns one over the length of the 3D vector  $\frac{1}{\sqrt{p.x^2 + p.y^2 + p.z^2 + p.t^2}}$ . If the square root would overflow, returns 0. If the square root would underflow, returns  $+\infty$ .

### Description

Calculate one over the length of four dimensional vector  $p$  in euclidean space undue overflow or underflow.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double round (double x)`

Round to nearest integer value in floating-point.

### Returns

Returns rounded integer value.

### Description

Round  $x$  to the nearest integer value in floating-point format, with halfway cases rounded away from zero.



This function may be slower than alternate rounding methods. See [rint\(\)](#).

## `__device__ double rsqrt (double x)`

Calculate the reciprocal of the square root of the input argument.

### Returns

Returns  $1/\sqrt{x}$ .

- `rsqrt( +∞ )` returns `+0`.

- ▶ `rsqrt( ± 0 )` returns  $\pm \infty$ .
- ▶ `rsqrt(x)` returns NaN if  $x$  is less than 0.

### Description

Calculate the reciprocal of the nonnegative square root of  $x$ ,  $1/\sqrt{x}$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double scalbln (double x, long int n)`

Scale floating-point input by integer power of two.

### Returns

Returns  $x * 2^n$ .

- ▶ `scalbln( ± 0 , n)` returns  $\pm 0$ .
- ▶ `scalbln(x, 0)` returns  $x$ .
- ▶ `scalbln( ± ∞ , n)` returns  $\pm \infty$ .

### Description

Scale  $x$  by  $2^n$  by efficient manipulation of the floating-point exponent.

## `__device__ double scalbn (double x, int n)`

Scale floating-point input by integer power of two.

### Returns

Returns  $x * 2^n$ .

- ▶ `scalbn( ± 0 , n)` returns  $\pm 0$ .
- ▶ `scalbn(x, 0)` returns  $x$ .
- ▶ `scalbn( ± ∞ , n)` returns  $\pm \infty$ .

### Description

Scale  $x$  by  $2^n$  by efficient manipulation of the floating-point exponent.

## `__device__ __RETURN_TYPE signbit (double a)`

Return the sign bit of the input.

### Returns

Reports the sign bit of all values including infinities, zeros, and NaNs.

- ▶ With Visual Studio 2013 host compiler: `__RETURN_TYPE` is 'bool'. Returns true if and only if `a` is negative.
- ▶ With other host compilers: `__RETURN_TYPE` is 'int'. Returns a nonzero value if and only if `a` is negative.

### Description

Determine whether the floating-point value `a` is negative.

## `__device__ double sin (double x)`

Calculate the sine of the input argument.

### Returns

- ▶ `sin( ±0 )` returns `±0`.
- ▶ `sin( ±∞ )` returns NaN.

### Description

Calculate the sine of the input argument `x` (measured in radians).



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ void sincos (double x, double *sptr, double *cptr)`

Calculate the sine and cosine of the first input argument.

### Returns

- ▶ none

**Description**

Calculate the sine and cosine of the first input argument  $x$  (measured in radians). The results for sine and cosine are written into the second argument, `sptr`, and, respectively, third argument, `cptr`.

**See also:**

[sin\(\)](#) and [cos\(\)](#).



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## **`__device__ void sincospi (double x, double *sptr, double *cptr)`**

Calculate the sine and cosine of the first input argument  $x \times \pi$ .

**Returns**

- none

**Description**

Calculate the sine and cosine of the first input argument,  $x$  (measured in radians),  $x \times \pi$ . The results for sine and cosine are written into the second argument, `sptr`, and, respectively, third argument, `cptr`.

**See also:**

[sinpi\(\)](#) and [cospi\(\)](#).



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## **`__device__ double sinh (double x)`**

Calculate the hyperbolic sine of the input argument.

**Returns**

- `sinh(  $\pm 0$  )` returns  $\pm 0$ .

**Description**

Calculate the hyperbolic sine of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double sinpi (double x)`

Calculate the sine of the input argument  $x \times \pi$ .

### Returns

- ▶ `sinpi( ± 0 )` returns  $\pm 0$ .
- ▶ `sinpi( ± ∞ )` returns NaN.

### Description

Calculate the sine of  $x \times \pi$  (measured in radians), where  $x$  is the input argument.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double sqrt (double x)`

Calculate the square root of the input argument.

### Returns

Returns  $\sqrt{x}$ .

- ▶ `sqrt( ± 0 )` returns  $\pm 0$ .
- ▶ `sqrt( + ∞ )` returns  $+\infty$ .
- ▶ `sqrt(x)` returns NaN if  $x$  is less than 0.

### Description

Calculate the nonnegative square root of  $x$ ,  $\sqrt{x}$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double tan (double x)`

Calculate the tangent of the input argument.

### Returns

- ▶  $\tan(\pm 0)$  returns  $\pm 0$ .
- ▶  $\tan(\pm \infty)$  returns NaN.

### Description

Calculate the tangent of the input argument  $x$  (measured in radians).



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double tanh (double x)`

Calculate the hyperbolic tangent of the input argument.

### Returns

- ▶  $\tanh(\pm 0)$  returns  $\pm 0$ .

### Description

Calculate the hyperbolic tangent of the input argument  $x$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double tgamma (double x)`

Calculate the gamma function of the input argument.

### Returns

- ▶  $\text{tgamma}(\pm 0)$  returns  $\pm \infty$ .
- ▶  $\text{tgamma}(2)$  returns  $+0$ .
- ▶  $\text{tgamma}(x)$  returns  $\pm \infty$  if the correctly calculated value is outside the double floating point range.
- ▶  $\text{tgamma}(x)$  returns NaN if  $x < 0$ .
- ▶  $\text{tgamma}(-\infty)$  returns NaN.
- ▶  $\text{tgamma}(+\infty)$  returns  $+\infty$ .



**Description**

Calculate the gamma function of the input argument  $x$ , namely the value of  $\int_0^{\infty} e^{-t} t^{x-1} dt$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

**\_\_device\_\_ double trunc (double x)**

Truncate input argument to the integral part.

**Returns**

Returns truncated integer value.

**Description**

Round  $x$  to the nearest integer value that does not exceed  $x$  in magnitude.

**\_\_device\_\_ \_\_CUDA\_MATH\_CRTIMP double y0 (double x)**

Calculate the value of the Bessel function of the second kind of order 0 for the input argument.

**Returns**

Returns the value of the Bessel function of the second kind of order 0.

- ▶  $y_0(0)$  returns  $-\infty$ .
- ▶  $y_0(x)$  returns NaN for  $x < 0$ .
- ▶  $y_0(+\infty)$  returns +0.
- ▶  $y_0(\text{NaN})$  returns NaN.

**Description**

Calculate the value of the Bessel function of the second kind of order 0 for the input argument  $x$ ,  $Y_0(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ __CUDA_MATH_CRTIMP double y1 (double x)`

Calculate the value of the Bessel function of the second kind of order 1 for the input argument.

### Returns

Returns the value of the Bessel function of the second kind of order 1.

- ▶ `y1(0)` returns  $-\infty$ .
- ▶ `y1(x)` returns NaN for  $x < 0$ .
- ▶ `y1(+\infty)` returns +0.
- ▶ `y1(NaN)` returns NaN.

### Description

Calculate the value of the Bessel function of the second kind of order 1 for the input argument  $x$ ,  $Y_1(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ __CUDA_MATH_CRTIMP double yn (int n, double x)`

Calculate the value of the Bessel function of the second kind of order  $n$  for the input argument.

### Returns

Returns the value of the Bessel function of the second kind of order  $n$ .

- ▶ `yn(n, x)` returns NaN for  $n < 0$ .
- ▶ `yn(n, 0)` returns  $-\infty$ .
- ▶ `yn(n, x)` returns NaN for  $x < 0$ .
- ▶ `yn(n, +\infty)` returns +0.
- ▶ `yn(n, NaN)` returns NaN.

### Description

Calculate the value of the Bessel function of the second kind of order  $n$  for the input argument  $x$ ,  $Y_n(x)$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## 1.4. Single Precision Intrinsics

This section describes single precision intrinsic functions that are only supported in device code.

### `__device__ float __cosf (float x)`

Calculate the fast approximate cosine of the input argument.

#### Returns

Returns the approximate cosine of  $x$ .

#### Description

Calculate the fast approximate cosine of the input argument  $x$ , measured in radians.



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.2, Table 9.
- Input and output in the denormal range is flushed to sign preserving 0.0.

### `__device__ float __exp10f (float x)`

Calculate the fast approximate base 10 exponential of the input argument.

#### Returns

Returns an approximation to  $10^x$ .

#### Description

Calculate the fast approximate base 10 exponential of the input argument  $x$ ,  $10^x$ .



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.2, Table 9.
- Most input and output values around denormal range are flushed to sign preserving 0.0.

## `__device__ float __expf (float x)`

Calculate the fast approximate base  $e$  exponential of the input argument.

### Returns

Returns an approximation to  $e^x$ .

### Description

Calculate the fast approximate base  $e$  exponential of the input argument  $x$ ,  $e^x$ .



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.2, Table 9.
- Most input and output values around denormal range are flushed to sign preserving 0.0.

## `__device__ float __fadd_rd (float x, float y)`

Add two floating point values in round-down mode.

### Returns

Returns  $x + y$ .

### Description

Compute the sum of  $x$  and  $y$  in round-down (to negative infinity) mode.



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- This operation will never be merged into a single multiply-add instruction.

## `__device__ float __fadd_rn (float x, float y)`

Add two floating point values in round-to-nearest-even mode.

### Returns

Returns  $x + y$ .

### Description

Compute the sum of  $x$  and  $y$  in round-to-nearest-even rounding mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ float __fadd_ru (float x, float y)`

Add two floating point values in round-up mode.

### Returns

Returns  $x + y$ .

### Description

Compute the sum of  $x$  and  $y$  in round-up (to positive infinity) mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ float __fadd_rz (float x, float y)`

Add two floating point values in round-towards-zero mode.

### Returns

Returns  $x + y$ .

### Description

Compute the sum of  $x$  and  $y$  in round-towards-zero mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ float __fddiv_rd (float x, float y)`

Divide two floating point values in round-down mode.

### Returns

Returns  $x / y$ .

**Description**

Divide two floating point values  $x$  by  $y$  in round-down (to negative infinity) mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## **\_\_device\_\_ float \_\_fdiv\_rn (float x, float y)**

Divide two floating point values in round-to-nearest-even mode.

**Returns**

Returns  $x / y$ .

**Description**

Divide two floating point values  $x$  by  $y$  in round-to-nearest-even mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## **\_\_device\_\_ float \_\_fdiv\_ru (float x, float y)**

Divide two floating point values in round-up mode.

**Returns**

Returns  $x / y$ .

**Description**

Divide two floating point values  $x$  by  $y$  in round-up (to positive infinity) mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## **\_\_device\_\_ float \_\_fdiv\_rz (float x, float y)**

Divide two floating point values in round-towards-zero mode.

**Returns**

Returns  $x / y$ .

**Description**

Divide two floating point values  $x$  by  $y$  in round-towards-zero mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float __fdividef (float x, float y)`

Calculate the fast approximate division of the input arguments.

**Returns**

Returns  $x / y$ .

- ▶ `__fdividef(  $\infty$ ,  $y$  )` returns NaN for  $2^{126} < y < 2^{128}$ .
- ▶ `__fdividef( $x$ ,  $y$ )` returns 0 for  $2^{126} < y < 2^{128}$  and  $x \neq \infty$ .

**Description**

Calculate the fast approximate division of  $x$  by  $y$ .



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.2, Table 9.

## `__device__ float __fmaf_rd (float x, float y, float z)`

Compute  $x \times y + z$  as a single operation, in round-down mode.

**Returns**

Returns the rounded value of  $x \times y + z$  as a single operation.

- ▶ `fmaf(  $\pm \infty$ ,  $\pm 0$ ,  $z$  )` returns NaN.
- ▶ `fmaf(  $\pm 0$ ,  $\pm \infty$ ,  $z$  )` returns NaN.
- ▶ `fmaf( $x$ ,  $y$ ,  $-\infty$ )` returns NaN if  $x \times y$  is an exact  $+\infty$ .
- ▶ `fmaf( $x$ ,  $y$ ,  $+\infty$ )` returns NaN if  $x \times y$  is an exact  $-\infty$ .

**Description**

Computes the value of  $x \times y + z$  as a single ternary operation, rounding the result once in round-down (to negative infinity) mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float __fmaf_rn (float x, float y, float z)`

Compute  $x \times y + z$  as a single operation, in round-to-nearest-even mode.

### Returns

Returns the rounded value of  $x \times y + z$  as a single operation.

- ▶ `fmaf(  $\pm\infty$ ,  $\pm 0$ , z )` returns NaN.
- ▶ `fmaf(  $\pm 0$ ,  $\pm\infty$ , z )` returns NaN.
- ▶ `fmaf(x, y,  $-\infty$ )` returns NaN if  $x \times y$  is an exact  $+\infty$ .
- ▶ `fmaf(x, y,  $+\infty$ )` returns NaN if  $x \times y$  is an exact  $-\infty$ .

### Description

Computes the value of  $x \times y + z$  as a single ternary operation, rounding the result once in round-to-nearest-even mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float __fmaf_ru (float x, float y, float z)`

Compute  $x \times y + z$  as a single operation, in round-up mode.

### Returns

Returns the rounded value of  $x \times y + z$  as a single operation.

- ▶ `fmaf(  $\pm\infty$ ,  $\pm 0$ , z )` returns NaN.
- ▶ `fmaf(  $\pm 0$ ,  $\pm\infty$ , z )` returns NaN.
- ▶ `fmaf(x, y,  $-\infty$ )` returns NaN if  $x \times y$  is an exact  $+\infty$ .
- ▶ `fmaf(x, y,  $+\infty$ )` returns NaN if  $x \times y$  is an exact  $-\infty$ .

### Description

Computes the value of  $x \times y + z$  as a single ternary operation, rounding the result once in round-up (to positive infinity) mode.





For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float __fmaf_rz (float x, float y, float z)`

Compute  $x \times y + z$  as a single operation, in round-towards-zero mode.

### Returns

Returns the rounded value of  $x \times y + z$  as a single operation.

- ▶ `fmaf(  $\pm\infty$ ,  $\pm 0$ , z )` returns NaN.
- ▶ `fmaf(  $\pm 0$ ,  $\pm\infty$ , z )` returns NaN.
- ▶ `fmaf(x, y,  $-\infty$ )` returns NaN if  $x \times y$  is an exact  $+\infty$ .
- ▶ `fmaf(x, y,  $+\infty$ )` returns NaN if  $x \times y$  is an exact  $-\infty$ .

### Description

Computes the value of  $x \times y + z$  as a single ternary operation, rounding the result once in round-towards-zero mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float __fmul_rd (float x, float y)`

Multiply two floating point values in round-down mode.

### Returns

Returns  $x * y$ .

### Description

Compute the product of  $x$  and  $y$  in round-down (to negative infinity) mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ float __fmul_rn (float x, float y)`

Multiply two floating point values in round-to-nearest-even mode.

### Returns

Returns  $x * y$ .

### Description

Compute the product of  $x$  and  $y$  in round-to-nearest-even mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ float __fmul_ru (float x, float y)`

Multiply two floating point values in round-up mode.

### Returns

Returns  $x * y$ .

### Description

Compute the product of  $x$  and  $y$  in round-up (to positive infinity) mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ float __fmul_rz (float x, float y)`

Multiply two floating point values in round-towards-zero mode.

### Returns

Returns  $x * y$ .

### Description

Compute the product of  $x$  and  $y$  in round-towards-zero mode.



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- This operation will never be merged into a single multiply-add instruction.

## `__device__ float __frcp_rd (float x)`

Compute  $\frac{1}{x}$  in round-down mode.

### Returns

Returns  $\frac{1}{x}$ .

### Description

Compute the reciprocal of  $x$  in round-down (to negative infinity) mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float __frcp_rn (float x)`

Compute  $\frac{1}{x}$  in round-to-nearest-even mode.

### Returns

Returns  $\frac{1}{x}$ .

### Description

Compute the reciprocal of  $x$  in round-to-nearest-even mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float __frcp_ru (float x)`

Compute  $\frac{1}{x}$  in round-up mode.

### Returns

Returns  $\frac{1}{x}$ .

**Description**

Compute the reciprocal of  $x$  in round-up (to positive infinity) mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float __frcp_rz (float x)`

Compute  $\frac{1}{x}$  in round-towards-zero mode.

**Returns**

Returns  $\frac{1}{x}$ .

**Description**

Compute the reciprocal of  $x$  in round-towards-zero mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float __frsqrtn (float x)`

Compute  $1/\sqrt{x}$  in round-to-nearest-even mode.

**Returns**

Returns  $1/\sqrt{x}$ .

**Description**

Compute the reciprocal square root of  $x$  in round-to-nearest-even mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

**\_\_device\_\_ float \_\_fsqrt\_rd (float x)**

Compute  $\sqrt{x}$  in round-down mode.

**Returns**

Returns  $\sqrt{x}$ .

**Description**

Compute the square root of  $x$  in round-down (to negative infinity) mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

**\_\_device\_\_ float \_\_fsqrt\_rn (float x)**

Compute  $\sqrt{x}$  in round-to-nearest-even mode.

**Returns**

Returns  $\sqrt{x}$ .

**Description**

Compute the square root of  $x$  in round-to-nearest-even mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

**\_\_device\_\_ float \_\_fsqrt\_ru (float x)**

Compute  $\sqrt{x}$  in round-up mode.

**Returns**

Returns  $\sqrt{x}$ .

**Description**

Compute the square root of  $x$  in round-up (to positive infinity) mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float __fsqrt_rz (float x)`

Compute  $\sqrt{x}$  in round-towards-zero mode.

### Returns

Returns  $\sqrt{x}$ .

### Description

Compute the square root of  $x$  in round-towards-zero mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.

## `__device__ float __fsub_rd (float x, float y)`

Subtract two floating point values in round-down mode.

### Returns

Returns  $x - y$ .

### Description

Compute the difference of  $x$  and  $y$  in round-down (to negative infinity) mode.



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- This operation will never be merged into a single multiply-add instruction.

## `__device__ float __fsub_rn (float x, float y)`

Subtract two floating point values in round-to-nearest-even mode.

### Returns

Returns  $x - y$ .

**Description**

Compute the difference of  $x$  and  $y$  in round-to-nearest-even rounding mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This operation will never be merged into a single multiply-add instruction.

## **`__device__ float __fsub_ru (float x, float y)`**

Subtract two floating point values in round-up mode.

**Returns**

Returns  $x - y$ .

**Description**

Compute the difference of  $x$  and  $y$  in round-up (to positive infinity) mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This operation will never be merged into a single multiply-add instruction.

## **`__device__ float __fsub_rz (float x, float y)`**

Subtract two floating point values in round-towards-zero mode.

**Returns**

Returns  $x - y$ .

**Description**

Compute the difference of  $x$  and  $y$  in round-towards-zero mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 6.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ float __log10f (float x)`

Calculate the fast approximate base 10 logarithm of the input argument.

### Returns

Returns an approximation to  $\log_{10}(x)$ .

### Description

Calculate the fast approximate base 10 logarithm of the input argument  $x$ .



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.2, Table 9.
- ▶ Most input and output values around denormal range are flushed to sign preserving 0.0.

## `__device__ float __log2f (float x)`

Calculate the fast approximate base 2 logarithm of the input argument.

### Returns

Returns an approximation to  $\log_2(x)$ .

### Description

Calculate the fast approximate base 2 logarithm of the input argument  $x$ .



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.2, Table 9.
- ▶ Input and output in the denormal range is flushed to sign preserving 0.0.

## `__device__ float __logf (float x)`

Calculate the fast approximate base  $e$  logarithm of the input argument.

### Returns

Returns an approximation to  $\log_e(x)$ .

### Description

Calculate the fast approximate base  $e$  logarithm of the input argument  $x$ .





- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.2, Table 9.
- ▶ Most input and output values around denormal range are flushed to sign preserving 0.0.

## `__device__ float __powf (float x, float y)`

Calculate the fast approximate of  $x^y$ .

### Returns

Returns an approximation to  $x^y$ .

### Description

Calculate the fast approximate of  $x$ , the first input argument, raised to the power of  $y$ , the second input argument,  $x^y$ .



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.2, Table 9.
- ▶ Most input and output values around denormal range are flushed to sign preserving 0.0.

## `__device__ float __saturatef (float x)`

Clamp the input argument to  $[+0.0, 1.0]$ .

### Returns

- ▶ `__saturatef(x)` returns 0 if  $x < 0$ .
- ▶ `__saturatef(x)` returns 1 if  $x > 1$ .
- ▶ `__saturatef(x)` returns  $x$  if  $0 \leq x \leq 1$ .
- ▶ `__saturatef(NaN)` returns 0.

### Description

Clamp the input argument  $x$  to be within the interval  $[+0.0, 1.0]$ .

## `__device__ void __sincosf (float x, float *sptr, float *cptr)`

Calculate the fast approximate of sine and cosine of the first input argument.

### Returns

- ▶ none

### Description

Calculate the fast approximate of sine and cosine of the first input argument  $x$  (measured in radians). The results for sine and cosine are written into the second argument, `sptr`, and, respectively, third argument, `cptr`.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.2, Table 9.
- ▶ Denorm input/output is flushed to sign preserving 0.0.

## `__device__ float __sinf (float x)`

Calculate the fast approximate sine of the input argument.

### Returns

Returns the approximate sine of  $x$ .

### Description

Calculate the fast approximate sine of the input argument  $x$ , measured in radians.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.2, Table 9.
- ▶ Input and output in the denormal range is flushed to sign preserving 0.0.

## `__device__ float __tanf (float x)`

Calculate the fast approximate tangent of the input argument.

### Returns

Returns the approximate tangent of  $x$ .

**Description**

Calculate the fast approximate tangent of the input argument  $x$ , measured in radians.



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.2, Table 9.
- The result is computed as the fast divide of `__sinf()` by `__cosf()`. Denormal input and output are flushed to sign-preserving 0.0 at each step of the computation.

## 1.5. Double Precision Intrinsic

This section describes double precision intrinsic functions that are only supported in device code.

### `__device__ double __dadd_rd (double x, double y)`

Add two floating point values in round-down mode.

**Returns**

Returns  $x + y$ .

**Description**

Adds two floating point values  $x$  and  $y$  in round-down (to negative infinity) mode.



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- This operation will never be merged into a single multiply-add instruction.

### `__device__ double __dadd_rn (double x, double y)`

Add two floating point values in round-to-nearest-even mode.

**Returns**

Returns  $x + y$ .

**Description**

Adds two floating point values  $x$  and  $y$  in round-to-nearest-even mode.



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- This operation will never be merged into a single multiply-add instruction.

## `__device__ double __dadd_ru (double x, double y)`

Add two floating point values in round-up mode.

### Returns

Returns  $x + y$ .

### Description

Adds two floating point values  $x$  and  $y$  in round-up (to positive infinity) mode.



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- This operation will never be merged into a single multiply-add instruction.

## `__device__ double __dadd_rz (double x, double y)`

Add two floating point values in round-towards-zero mode.

### Returns

Returns  $x + y$ .

### Description

Adds two floating point values  $x$  and  $y$  in round-towards-zero mode.



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- This operation will never be merged into a single multiply-add instruction.

## `__device__ double __ddiv_rd (double x, double y)`

Divide two floating point values in round-down mode.

### Returns

Returns  $x / y$ .

**Description**

Divides two floating point values  $x$  by  $y$  in round-down (to negative infinity) mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ Requires compute capability  $\geq 2.0$ .

**\_\_device\_\_ double \_\_ddiv\_rn (double x, double y)**

Divide two floating point values in round-to-nearest-even mode.

**Returns**

Returns  $x / y$ .

**Description**

Divides two floating point values  $x$  by  $y$  in round-to-nearest-even mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ Requires compute capability  $\geq 2.0$ .

**\_\_device\_\_ double \_\_ddiv\_ru (double x, double y)**

Divide two floating point values in round-up mode.

**Returns**

Returns  $x / y$ .

**Description**

Divides two floating point values  $x$  by  $y$  in round-up (to positive infinity) mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ Requires compute capability  $\geq 2.0$ .

## `__device__ double __ddiv_rz (double x, double y)`

Divide two floating point values in round-towards-zero mode.

### Returns

Returns  $x / y$ .

### Description

Divides two floating point values  $x$  by  $y$  in round-towards-zero mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ Requires compute capability  $\geq 2.0$ .

## `__device__ double __dmul_rd (double x, double y)`

Multiply two floating point values in round-down mode.

### Returns

Returns  $x * y$ .

### Description

Multiplies two floating point values  $x$  and  $y$  in round-down (to negative infinity) mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ double __dmul_rn (double x, double y)`

Multiply two floating point values in round-to-nearest-even mode.

### Returns

Returns  $x * y$ .

### Description

Multiplies two floating point values  $x$  and  $y$  in round-to-nearest-even mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ double __dmul_ru (double x, double y)`

Multiply two floating point values in round-up mode.

### Returns

Returns  $x * y$ .

### Description

Multiplies two floating point values  $x$  and  $y$  in round-up (to positive infinity) mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ double __dmul_rz (double x, double y)`

Multiply two floating point values in round-towards-zero mode.

### Returns

Returns  $x * y$ .

### Description

Multiplies two floating point values  $x$  and  $y$  in round-towards-zero mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ double __drcp_rd (double x)`

Compute  $\frac{1}{x}$  in round-down mode.

### Returns

Returns  $\frac{1}{x}$ .

**Description**

Compute the reciprocal of  $x$  in round-down (to negative infinity) mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ Requires compute capability  $\geq 2.0$ .

**\_\_device\_\_ double \_\_drcp\_rn (double x)**

Compute  $\frac{1}{x}$  in round-to-nearest-even mode.

**Returns**

Returns  $\frac{1}{x}$ .

**Description**

Compute the reciprocal of  $x$  in round-to-nearest-even mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ Requires compute capability  $\geq 2.0$ .

**\_\_device\_\_ double \_\_drcp\_ru (double x)**

Compute  $\frac{1}{x}$  in round-up mode.

**Returns**

Returns  $\frac{1}{x}$ .

**Description**

Compute the reciprocal of  $x$  in round-up (to positive infinity) mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ Requires compute capability  $\geq 2.0$ .



## `__device__ double __drcp_rz (double x)`

Compute  $\frac{1}{x}$  in round-towards-zero mode.

### Returns

Returns  $\frac{1}{x}$ .

### Description

Compute the reciprocal of  $x$  in round-towards-zero mode.



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- Requires compute capability  $\geq 2.0$ .

## `__device__ double __dsqrt_rd (double x)`

Compute  $\sqrt{x}$  in round-down mode.

### Returns

Returns  $\sqrt{x}$ .

### Description

Compute the square root of  $x$  in round-down (to negative infinity) mode.



- For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- Requires compute capability  $\geq 2.0$ .

## `__device__ double __dsqrt_rn (double x)`

Compute  $\sqrt{x}$  in round-to-nearest-even mode.

### Returns

Returns  $\sqrt{x}$ .

### Description

Compute the square root of  $x$  in round-to-nearest-even mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ Requires compute capability  $\geq 2.0$ .

## `__device__ double __dsqrt_ru (double x)`

Compute  $\sqrt{x}$  in round-up mode.

### Returns

Returns  $\sqrt{x}$ .

### Description

Compute the square root of  $x$  in round-up (to positive infinity) mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ Requires compute capability  $\geq 2.0$ .

## `__device__ double __dsqrt_rz (double x)`

Compute  $\sqrt{x}$  in round-towards-zero mode.

### Returns

Returns  $\sqrt{x}$ .

### Description

Compute the square root of  $x$  in round-towards-zero mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ Requires compute capability  $\geq 2.0$ .

## `__device__ double __dsub_rd (double x, double y)`

Subtract two floating point values in round-down mode.

### Returns

Returns  $x - y$ .

### Description

Subtracts two floating point values  $x$  and  $y$  in round-down (to negative infinity) mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ double __dsub_rn (double x, double y)`

Subtract two floating point values in round-to-nearest-even mode.

### Returns

Returns  $x - y$ .

### Description

Subtracts two floating point values  $x$  and  $y$  in round-to-nearest-even mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ double __dsub_ru (double x, double y)`

Subtract two floating point values in round-up mode.

### Returns

Returns  $x - y$ .

### Description

Subtracts two floating point values  $x$  and  $y$  in round-up (to positive infinity) mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ double __dsub_rz (double x, double y)`

Subtract two floating point values in round-towards-zero mode.

### Returns

Returns  $x - y$ .

### Description

Subtracts two floating point values  $x$  and  $y$  in round-towards-zero mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.
- ▶ This operation will never be merged into a single multiply-add instruction.

## `__device__ double __fma_rd (double x, double y, double z)`

Compute  $x \times y + z$  as a single operation in round-down mode.

### Returns

Returns the rounded value of  $x \times y + z$  as a single operation.

- ▶ `fmaf(  $\pm\infty$ ,  $\pm 0$ ,  $z$  )` returns NaN.
- ▶ `fmaf(  $\pm 0$ ,  $\pm\infty$ ,  $z$  )` returns NaN.
- ▶ `fmaf( $x$ ,  $y$ ,  $-\infty$  )` returns NaN if  $x \times y$  is an exact  $+\infty$
- ▶ `fmaf( $x$ ,  $y$ ,  $+\infty$  )` returns NaN if  $x \times y$  is an exact  $-\infty$

### Description

Computes the value of  $x \times y + z$  as a single ternary operation, rounding the result once in round-down (to negative infinity) mode.



- ▶ For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double __fma_rn (double x, double y, double z)`

Compute  $x \times y + z$  as a single operation in round-to-nearest-even mode.

### Returns

Returns the rounded value of  $x \times y + z$  as a single operation.

- ▶ `fmaf(  $\pm\infty$ ,  $\pm 0$ ,  $z$  )` returns NaN.
- ▶ `fmaf(  $\pm 0$ ,  $\pm\infty$ ,  $z$  )` returns NaN.
- ▶ `fmaf( $x$ ,  $y$ ,  $-\infty$  )` returns NaN if  $x \times y$  is an exact  $+\infty$
- ▶ `fmaf( $x$ ,  $y$ ,  $+\infty$  )` returns NaN if  $x \times y$  is an exact  $-\infty$

### Description

Computes the value of  $x \times y + z$  as a single ternary operation, rounding the result once in round-to-nearest-even mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double __fma_ru (double x, double y, double z)`

Compute  $x \times y + z$  as a single operation in round-up mode.

### Returns

Returns the rounded value of  $x \times y + z$  as a single operation.

- ▶ `fmaf(  $\pm\infty$ ,  $\pm 0$ ,  $z$  )` returns NaN.
- ▶ `fmaf(  $\pm 0$ ,  $\pm\infty$ ,  $z$  )` returns NaN.
- ▶ `fmaf( $x$ ,  $y$ ,  $-\infty$  )` returns NaN if  $x \times y$  is an exact  $+\infty$
- ▶ `fmaf( $x$ ,  $y$ ,  $+\infty$  )` returns NaN if  $x \times y$  is an exact  $-\infty$

### Description

Computes the value of  $x \times y + z$  as a single ternary operation, rounding the result once in round-up (to positive infinity) mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## `__device__ double __fma_rz (double x, double y, double z)`

Compute  $x \times y + z$  as a single operation in round-towards-zero mode.

### Returns

Returns the rounded value of  $x \times y + z$  as a single operation.

- ▶ `fmaf(  $\pm\infty$ ,  $\pm 0$ , z )` returns NaN.
- ▶ `fmaf(  $\pm 0$ ,  $\pm\infty$ , z )` returns NaN.
- ▶ `fmaf(x, y,  $-\infty$ )` returns NaN if  $x \times y$  is an exact  $+\infty$
- ▶ `fmaf(x, y,  $+\infty$ )` returns NaN if  $x \times y$  is an exact  $-\infty$

### Description

Computes the value of  $x \times y + z$  as a single ternary operation, rounding the result once in round-towards-zero mode.



For accuracy information for this function see the CUDA C Programming Guide, Appendix D.1, Table 7.

## 1.6. Integer Intrinsics

This section describes integer intrinsic functions that are only supported in device code.

### `__device__ unsigned int __brev (unsigned int x)`

Reverse the bit order of a 32 bit unsigned integer.

### Returns

Returns the bit-reversed value of  $x$ . i.e. bit  $N$  of the return value corresponds to bit  $31-N$  of  $x$ .

### Description

Reverses the bit order of the 32 bit unsigned integer  $x$ .

## `__device__ unsigned long long int __brevll (unsigned long long int x)`

Reverse the bit order of a 64 bit unsigned integer.

### Returns

Returns the bit-reversed value of  $x$ . i.e. bit  $N$  of the return value corresponds to bit  $63-N$  of  $x$ .

### Description

Reverses the bit order of the 64 bit unsigned integer  $x$ .

## `__device__ unsigned int __byte_perm (unsigned int x, unsigned int y, unsigned int s)`

Return selected bytes from two 32 bit unsigned integers.

### Returns

The returned value  $r$  is computed to be: `result[n] := input[selector[n]]` where `result[n]` is the  $n$ th byte of  $r$ .

### Description

`byte_perm(x,y,s)` returns a 32-bit integer consisting of four bytes from eight input bytes provided in the two input integers  $x$  and  $y$ , as specified by a selector,  $s$ .

The input bytes are indexed as follows: `input[0] = x<7:0>` `input[1] = x<15:8>` `input[2] = x<23:16>` `input[3] = x<31:24>` `input[4] = y<7:0>` `input[5] = y<15:8>` `input[6] = y<23:16>` `input[7] = y<31:24>` The selector indices are as follows (the upper 16-bits of the selector are not used): `selector[0] = s<2:0>` `selector[1] = s<6:4>` `selector[2] = s<10:8>` `selector[3] = s<14:12>`

## `__device__ int __clz (int x)`

Return the number of consecutive high-order zero bits in a 32 bit integer.

### Returns

Returns a value between 0 and 32 inclusive representing the number of zero bits.

### Description

Count the number of consecutive leading zero bits, starting at the most significant bit (bit 31) of  $x$ .

## `__device__ int __clzll (long long int x)`

Count the number of consecutive high-order zero bits in a 64 bit integer.

### Returns

Returns a value between 0 and 64 inclusive representing the number of zero bits.

### Description

Count the number of consecutive leading zero bits, starting at the most significant bit (bit 63) of  $x$ .

## `__device__ int __ffs (int x)`

Find the position of the least significant bit set to 1 in a 32 bit integer.

### Returns

Returns a value between 0 and 32 inclusive representing the position of the first bit set.

- ▶ `__ffs(0)` returns 0.

### Description

Find the position of the first (least significant) bit set to 1 in  $x$ , where the least significant bit position is 1.

## `__device__ int __ffsll (long long int x)`

Find the position of the least significant bit set to 1 in a 64 bit integer.

### Returns

Returns a value between 0 and 64 inclusive representing the position of the first bit set.

- ▶ `__ffsll(0)` returns 0.

### Description

Find the position of the first (least significant) bit set to 1 in  $x$ , where the least significant bit position is 1.



## `__device__ int __hadd (int, int)`

Compute average of signed input arguments, avoiding overflow in the intermediate sum.

### Returns

Returns a signed integer value representing the signed average value of the two inputs.

### Description

Compute average of signed input arguments  $x$  and  $y$  as  $(x + y) >> 1$ , avoiding overflow in the intermediate sum.

## `__device__ int __mul24 (int x, int y)`

Calculate the least significant 32 bits of the product of the least significant 24 bits of two integers.

### Returns

Returns the least significant 32 bits of the product  $x * y$ .

### Description

Calculate the least significant 32 bits of the product of the least significant 24 bits of  $x$  and  $y$ . The high order 8 bits of  $x$  and  $y$  are ignored.

## `__device__ long long int __mul64hi (long long int x, long long int y)`

Calculate the most significant 64 bits of the product of the two 64 bit integers.

### Returns

Returns the most significant 64 bits of the product  $x * y$ .

### Description

Calculate the most significant 64 bits of the 128-bit product  $x * y$ , where  $x$  and  $y$  are 64-bit integers.

## `__device__ int __mulhi (int x, int y)`

Calculate the most significant 32 bits of the product of the two 32 bit integers.

### Returns

Returns the most significant 32 bits of the product  $x * y$ .

### Description

Calculate the most significant 32 bits of the 64-bit product  $x * y$ , where  $x$  and  $y$  are 32-bit integers.

## `__device__ int __popc (unsigned int x)`

Count the number of bits that are set to 1 in a 32 bit integer.

### Returns

Returns a value between 0 and 32 inclusive representing the number of set bits.

### Description

Count the number of bits that are set to 1 in  $x$ .

## `__device__ int __popcll (unsigned long long int x)`

Count the number of bits that are set to 1 in a 64 bit integer.

### Returns

Returns a value between 0 and 64 inclusive representing the number of set bits.

### Description

Count the number of bits that are set to 1 in  $x$ .

## `__device__ int __rhadd (int, int)`

Compute rounded average of signed input arguments, avoiding overflow in the intermediate sum.

### Returns

Returns a signed integer value representing the signed rounded average value of the two inputs.

**Description**

Compute average of signed input arguments  $x$  and  $y$  as  $(x + y + 1) \gg 1$ , avoiding overflow in the intermediate sum.

**\_\_device\_\_ unsigned int \_\_sad (int x, int y, unsigned int z)**

Calculate  $|x - y| + z$ , the sum of absolute difference.

**Returns**

Returns  $|x - y| + z$ .

**Description**

Calculate  $|x - y| + z$ , the 32-bit sum of the third argument  $z$  plus and the absolute value of the difference between the first argument,  $x$ , and second argument,  $y$ .

Inputs  $x$  and  $y$  are signed 32-bit integers, input  $z$  is a 32-bit unsigned integer.

**\_\_device\_\_ unsigned int \_\_uhadd (unsigned int, unsigned int)**

Compute average of unsigned input arguments, avoiding overflow in the intermediate sum.

**Returns**

Returns an unsigned integer value representing the unsigned average value of the two inputs.

**Description**

Compute average of unsigned input arguments  $x$  and  $y$  as  $(x + y) \gg 1$ , avoiding overflow in the intermediate sum.

**\_\_device\_\_ unsigned int \_\_umul24 (unsigned int x, unsigned int y)**

Calculate the least significant 32 bits of the product of the least significant 24 bits of two unsigned integers.

**Returns**

Returns the least significant 32 bits of the product  $x * y$ .

**Description**

Calculate the least significant 32 bits of the product of the least significant 24 bits of  $x$  and  $y$ . The high order 8 bits of  $x$  and  $y$  are ignored.

**\_\_device\_\_ unsigned long long int \_\_umul64hi (unsigned long long int x, unsigned long long int y)**

Calculate the most significant 64 bits of the product of the two 64 unsigned bit integers.

**Returns**

Returns the most significant 64 bits of the product  $x * y$ .

**Description**

Calculate the most significant 64 bits of the 128-bit product  $x * y$ , where  $x$  and  $y$  are 64-bit unsigned integers.

**\_\_device\_\_ unsigned int \_\_umulhi (unsigned int x, unsigned int y)**

Calculate the most significant 32 bits of the product of the two 32 bit unsigned integers.

**Returns**

Returns the most significant 32 bits of the product  $x * y$ .

**Description**

Calculate the most significant 32 bits of the 64-bit product  $x * y$ , where  $x$  and  $y$  are 32-bit unsigned integers.

**\_\_device\_\_ unsigned int \_\_urhadd (unsigned int, unsigned int)**

Compute rounded average of unsigned input arguments, avoiding overflow in the intermediate sum.

**Returns**

Returns an unsigned integer value representing the unsigned rounded average value of the two inputs.

**Description**

Compute average of unsigned input arguments  $x$  and  $y$  as  $(x + y + 1) >> 1$ , avoiding overflow in the intermediate sum.

**\_\_device\_\_ unsigned int \_\_usad (unsigned int x, unsigned int y, unsigned int z)**

Calculate  $|x - y| + z$ , the sum of absolute difference.

#### Returns

Returns  $|x - y| + z$ .

#### Description

Calculate  $|x - y| + z$ , the 32-bit sum of the third argument  $z$  plus and the absolute value of the difference between the first argument,  $x$ , and second argument,  $y$ .

Inputs  $x$ ,  $y$ , and  $z$  are unsigned 32-bit integers.

## 1.7. Type Casting Ininsics

This section describes type casting intrinsic functions that are only supported in device code.

**\_\_device\_\_ float \_\_double2float\_rd (double x)**

Convert a double to a float in round-down mode.

#### Returns

Returns converted value.

#### Description

Convert the double-precision floating point value  $x$  to a single-precision floating point value in round-down (to negative infinity) mode.

**\_\_device\_\_ float \_\_double2float\_rn (double x)**

Convert a double to a float in round-to-nearest-even mode.

#### Returns

Returns converted value.

#### Description

Convert the double-precision floating point value  $x$  to a single-precision floating point value in round-to-nearest-even mode.

## `__device__ float __double2float_ru (double x)`

Convert a double to a float in round-up mode.

### Returns

Returns converted value.

### Description

Convert the double-precision floating point value  $x$  to a single-precision floating point value in round-up (to positive infinity) mode.

## `__device__ float __double2float_rz (double x)`

Convert a double to a float in round-towards-zero mode.

### Returns

Returns converted value.

### Description

Convert the double-precision floating point value  $x$  to a single-precision floating point value in round-towards-zero mode.

## `__device__ int __double2hiint (double x)`

Reinterpret high 32 bits in a double as a signed integer.

### Returns

Returns reinterpreted value.

### Description

Reinterpret the high 32 bits in the double-precision floating point value  $x$  as a signed integer.

## `__device__ int __double2int_rd (double x)`

Convert a double to a signed int in round-down mode.

### Returns

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to a signed integer value in round-down (to negative infinity) mode.

**`__device__ int __double2int_rn (double x)`**

Convert a double to a signed int in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to a signed integer value in round-to-nearest-even mode.

**`__device__ int __double2int_ru (double x)`**

Convert a double to a signed int in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to a signed integer value in round-up (to positive infinity) mode.

**`__device__ int __double2int_rz (double)`**

Convert a double to a signed int in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to a signed integer value in round-towards-zero mode.

**\_\_device\_\_ long long int \_\_double2ll\_rd (double x)**

Convert a double to a signed 64-bit int in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to a signed 64-bit integer value in round-down (to negative infinity) mode.

**\_\_device\_\_ long long int \_\_double2ll\_rn (double x)**

Convert a double to a signed 64-bit int in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to a signed 64-bit integer value in round-to-nearest-even mode.

**\_\_device\_\_ long long int \_\_double2ll\_ru (double x)**

Convert a double to a signed 64-bit int in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to a signed 64-bit integer value in round-up (to positive infinity) mode.

**\_\_device\_\_ long long int \_\_double2ll\_rz (double)**

Convert a double to a signed 64-bit int in round-towards-zero mode.

**Returns**

Returns converted value.



**Description**

Convert the double-precision floating point value  $x$  to a signed 64-bit integer value in round-towards-zero mode.

**`__device__ int __double2loint (double x)`**

Reinterpret low 32 bits in a double as a signed integer.

**Returns**

Returns reinterpreted value.

**Description**

Reinterpret the low 32 bits in the double-precision floating point value  $x$  as a signed integer.

**`__device__ unsigned int __double2uint_rd (double x)`**

Convert a double to an unsigned int in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to an unsigned integer value in round-down (to negative infinity) mode.

**`__device__ unsigned int __double2uint_rn (double x)`**

Convert a double to an unsigned int in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to an unsigned integer value in round-to-nearest-even mode.

**\_\_device\_\_ unsigned int \_\_double2uint\_ru (double x)**

Convert a double to an unsigned int in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to an unsigned integer value in round-up (to positive infinity) mode.

**\_\_device\_\_ unsigned int \_\_double2uint\_rz (double)**

Convert a double to an unsigned int in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to an unsigned integer value in round-towards-zero mode.

**\_\_device\_\_ unsigned long long int \_\_double2ull\_rd (double x)**

Convert a double to an unsigned 64-bit int in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to an unsigned 64-bit integer value in round-down (to negative infinity) mode.

**\_\_device\_\_ unsigned long long int \_\_double2ull\_rn (double x)**

Convert a double to an unsigned 64-bit int in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to an unsigned 64-bit integer value in round-to-nearest-even mode.

**\_\_device\_\_ unsigned long long int \_\_double2ull\_ru (double x)**

Convert a double to an unsigned 64-bit int in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to an unsigned 64-bit integer value in round-up (to positive infinity) mode.

**\_\_device\_\_ unsigned long long int \_\_double2ull\_rz (double)**

Convert a double to an unsigned 64-bit int in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the double-precision floating point value  $x$  to an unsigned 64-bit integer value in round-towards-zero mode.

**\_\_device\_\_ long long int \_\_double\_as\_longlong (double x)**

Reinterpret bits in a double as a 64-bit signed integer.

**Returns**

Returns reinterpreted value.

**Description**

Reinterpret the bits in the double-precision floating point value  $x$  as a signed 64-bit integer.

## `__device__ unsigned short __float2half_rn (float x)`

Convert a single-precision float to a half-precision float in round-to-nearest-even mode.

### Returns

Returns converted value.

### Description

Convert the single-precision float value  $x$  to a half-precision floating point value represented in `unsigned short` format, in round-to-nearest-even mode.

## `__device__ int __float2int_rd (float x)`

Convert a float to a signed integer in round-down mode.

### Returns

Returns converted value.

### Description

Convert the single-precision floating point value  $x$  to a signed integer in round-down (to negative infinity) mode.

## `__device__ int __float2int_rn (float x)`

Convert a float to a signed integer in round-to-nearest-even mode.

### Returns

Returns converted value.

### Description

Convert the single-precision floating point value  $x$  to a signed integer in round-to-nearest-even mode.

## `__device__ int __float2int_ru (float)`

Convert a float to a signed integer in round-up mode.

### Returns

Returns converted value.

**Description**

Convert the single-precision floating point value  $x$  to a signed integer in round-up (to positive infinity) mode.

**\_\_device\_\_ int \_\_float2int\_rz (float x)**

Convert a float to a signed integer in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the single-precision floating point value  $x$  to a signed integer in round-towards-zero mode.

**\_\_device\_\_ long long int \_\_float2ll\_rd (float x)**

Convert a float to a signed 64-bit integer in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the single-precision floating point value  $x$  to a signed 64-bit integer in round-down (to negative infinity) mode.

**\_\_device\_\_ long long int \_\_float2ll\_rn (float x)**

Convert a float to a signed 64-bit integer in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the single-precision floating point value  $x$  to a signed 64-bit integer in round-to-nearest-even mode.

**\_\_device\_\_ long long int \_\_float2ll\_ru (float x)**

Convert a float to a signed 64-bit integer in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the single-precision floating point value  $x$  to a signed 64-bit integer in round-up (to positive infinity) mode.

**\_\_device\_\_ long long int \_\_float2ll\_rz (float x)**

Convert a float to a signed 64-bit integer in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the single-precision floating point value  $x$  to a signed 64-bit integer in round-towards-zero mode.

**\_\_device\_\_ unsigned int \_\_float2uint\_rd (float x)**

Convert a float to an unsigned integer in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the single-precision floating point value  $x$  to an unsigned integer in round-down (to negative infinity) mode.

**\_\_device\_\_ unsigned int \_\_float2uint\_rn (float x)**

Convert a float to an unsigned integer in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the single-precision floating point value  $x$  to an unsigned integer in round-to-nearest-even mode.

**\_\_device\_\_ unsigned int \_\_float2uint\_ru (float x)**

Convert a float to an unsigned integer in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the single-precision floating point value  $x$  to an unsigned integer in round-up (to positive infinity) mode.

**\_\_device\_\_ unsigned int \_\_float2uint\_rz (float x)**

Convert a float to an unsigned integer in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the single-precision floating point value  $x$  to an unsigned integer in round-towards-zero mode.

**\_\_device\_\_ unsigned long long int \_\_float2ull\_rd (float x)**

Convert a float to an unsigned 64-bit integer in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the single-precision floating point value  $x$  to an unsigned 64-bit integer in round-down (to negative infinity) mode.

**\_\_device\_\_ unsigned long long int \_\_float2ull\_rn (float x)**

Convert a float to an unsigned 64-bit integer in round-to-nearest-even mode.

#### Returns

Returns converted value.

#### Description

Convert the single-precision floating point value  $x$  to an unsigned 64-bit integer in round-to-nearest-even mode.

**\_\_device\_\_ unsigned long long int \_\_float2ull\_ru (float x)**

Convert a float to an unsigned 64-bit integer in round-up mode.

#### Returns

Returns converted value.

#### Description

Convert the single-precision floating point value  $x$  to an unsigned 64-bit integer in round-up (to positive infinity) mode.

**\_\_device\_\_ unsigned long long int \_\_float2ull\_rz (float x)**

Convert a float to an unsigned 64-bit integer in round-towards-zero mode.

#### Returns

Returns converted value.

#### Description

Convert the single-precision floating point value  $x$  to an unsigned 64-bit integer in round-towards\_zero mode.



**\_\_device\_\_ int \_\_float\_as\_int (float x)**

Reinterpret bits in a float as a signed integer.

**Returns**

Returns reinterpreted value.

**Description**

Reinterpret the bits in the single-precision floating point value  $x$  as a signed integer.

**\_\_device\_\_ unsigned int \_\_float\_as\_uint (float x)**

Reinterpret bits in a float as a unsigned integer.

**Returns**

Returns reinterpreted value.

**Description**

Reinterpret the bits in the single-precision floating point value  $x$  as a unsigned integer.

**\_\_device\_\_ float \_\_half2float (unsigned short x)**

Convert a half-precision float to a single-precision float in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value  $x$  represented in `unsigned short` format to a single-precision floating point value.

**\_\_device\_\_ double \_\_hiloint2double (int hi, int lo)**

Reinterpret high and low 32-bit integer values as a double.

**Returns**

Returns reinterpreted value.

**Description**

Reinterpret the integer value of `hi` as the high 32 bits of a double-precision floating point value and the integer value of `lo` as the low 32 bits of the same double-precision floating point value.

**\_\_device\_\_ double \_\_int2double\_rn (int x)**

Convert a signed int to a double.

**Returns**

Returns converted value.

**Description**

Convert the signed integer value `x` to a double-precision floating point value.

**\_\_device\_\_ float \_\_int2float\_rd (int x)**

Convert a signed integer to a float in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the signed integer value `x` to a single-precision floating point value in round-down (to negative infinity) mode.

**\_\_device\_\_ float \_\_int2float\_rn (int x)**

Convert a signed integer to a float in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the signed integer value `x` to a single-precision floating point value in round-to-nearest-even mode.

## `__device__ float __int2float_ru (int x)`

Convert a signed integer to a float in round-up mode.

### Returns

Returns converted value.

### Description

Convert the signed integer value  $x$  to a single-precision floating point value in round-up (to positive infinity) mode.

## `__device__ float __int2float_rz (int x)`

Convert a signed integer to a float in round-towards-zero mode.

### Returns

Returns converted value.

### Description

Convert the signed integer value  $x$  to a single-precision floating point value in round-towards-zero mode.

## `__device__ float __int_as_float (int x)`

Reinterpret bits in an integer as a float.

### Returns

Returns reinterpreted value.

### Description

Reinterpret the bits in the signed integer value  $x$  as a single-precision floating point value.

## `__device__ double __ll2double_rd (long long int x)`

Convert a signed 64-bit int to a double in round-down mode.

### Returns

Returns converted value.

**Description**

Convert the signed 64-bit integer value  $x$  to a double-precision floating point value in round-down (to negative infinity) mode.

**`__device__ double __ll2double_rn (long long int x)`**

Convert a signed 64-bit int to a double in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the signed 64-bit integer value  $x$  to a double-precision floating point value in round-to-nearest-even mode.

**`__device__ double __ll2double_ru (long long int x)`**

Convert a signed 64-bit int to a double in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the signed 64-bit integer value  $x$  to a double-precision floating point value in round-up (to positive infinity) mode.

**`__device__ double __ll2double_rz (long long int x)`**

Convert a signed 64-bit int to a double in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the signed 64-bit integer value  $x$  to a double-precision floating point value in round-towards-zero mode.

## `__device__ float __ll2float_rd (long long int x)`

Convert a signed integer to a float in round-down mode.

### Returns

Returns converted value.

### Description

Convert the signed integer value  $x$  to a single-precision floating point value in round-down (to negative infinity) mode.

## `__device__ float __ll2float_rn (long long int x)`

Convert a signed 64-bit integer to a float in round-to-nearest-even mode.

### Returns

Returns converted value.

### Description

Convert the signed 64-bit integer value  $x$  to a single-precision floating point value in round-to-nearest-even mode.

## `__device__ float __ll2float_ru (long long int x)`

Convert a signed integer to a float in round-up mode.

### Returns

Returns converted value.

### Description

Convert the signed integer value  $x$  to a single-precision floating point value in round-up (to positive infinity) mode.

## `__device__ float __ll2float_rz (long long int x)`

Convert a signed integer to a float in round-towards-zero mode.

### Returns

Returns converted value.

**Description**

Convert the signed integer value  $x$  to a single-precision floating point value in round-towards-zero mode.

**\_\_device\_\_ double \_\_longlong\_as\_double (long long int x)**

Reinterpret bits in a 64-bit signed integer as a double.

**Returns**

Returns reinterpreted value.

**Description**

Reinterpret the bits in the 64-bit signed integer value  $x$  as a double-precision floating point value.

**\_\_device\_\_ double \_\_uint2double\_rn (unsigned int x)**

Convert an unsigned int to a double.

**Returns**

Returns converted value.

**Description**

Convert the unsigned integer value  $x$  to a double-precision floating point value.

**\_\_device\_\_ float \_\_uint2float\_rd (unsigned int x)**

Convert an unsigned integer to a float in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned integer value  $x$  to a single-precision floating point value in round-down (to negative infinity) mode.

## `__device__ float __uint2float_rn (unsigned int x)`

Convert an unsigned integer to a float in round-to-nearest-even mode.

### Returns

Returns converted value.

### Description

Convert the unsigned integer value  $x$  to a single-precision floating point value in round-to-nearest-even mode.

## `__device__ float __uint2float_ru (unsigned int x)`

Convert an unsigned integer to a float in round-up mode.

### Returns

Returns converted value.

### Description

Convert the unsigned integer value  $x$  to a single-precision floating point value in round-up (to positive infinity) mode.

## `__device__ float __uint2float_rz (unsigned int x)`

Convert an unsigned integer to a float in round-towards-zero mode.

### Returns

Returns converted value.

### Description

Convert the unsigned integer value  $x$  to a single-precision floating point value in round-towards-zero mode.

## `__device__ float __uint_as_float (unsigned int x)`

Reinterpret bits in an unsigned integer as a float.

### Returns

Returns reinterpreted value.

**Description**

Reinterpret the bits in the unsigned integer value  $x$  as a single-precision floating point value.

**\_\_device\_\_ double \_\_ull2double\_rd (unsigned long long int x)**

Convert an unsigned 64-bit int to a double in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned 64-bit integer value  $x$  to a double-precision floating point value in round-down (to negative infinity) mode.

**\_\_device\_\_ double \_\_ull2double\_rn (unsigned long long int x)**

Convert an unsigned 64-bit int to a double in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned 64-bit integer value  $x$  to a double-precision floating point value in round-to-nearest-even mode.

**\_\_device\_\_ double \_\_ull2double\_ru (unsigned long long int x)**

Convert an unsigned 64-bit int to a double in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned 64-bit integer value  $x$  to a double-precision floating point value in round-up (to positive infinity) mode.



## `__device__ double __ull2double_rz (unsigned long long int x)`

Convert an unsigned 64-bit int to a double in round-towards-zero mode.

### Returns

Returns converted value.

### Description

Convert the unsigned 64-bit integer value  $x$  to a double-precision floating point value in round-towards-zero mode.

## `__device__ float __ull2float_rd (unsigned long long int x)`

Convert an unsigned integer to a float in round-down mode.

### Returns

Returns converted value.

### Description

Convert the unsigned integer value  $x$  to a single-precision floating point value in round-down (to negative infinity) mode.

## `__device__ float __ull2float_rn (unsigned long long int x)`

Convert an unsigned integer to a float in round-to-nearest-even mode.

### Returns

Returns converted value.

### Description

Convert the unsigned integer value  $x$  to a single-precision floating point value in round-to-nearest-even mode.

## `__device__ float __ull2float_ru (unsigned long long int x)`

Convert an unsigned integer to a float in round-up mode.

### Returns

Returns converted value.

### Description

Convert the unsigned integer value  $x$  to a single-precision floating point value in round-up (to positive infinity) mode.

## `__device__ float __ull2float_rz (unsigned long long int x)`

Convert an unsigned integer to a float in round-towards-zero mode.

### Returns

Returns converted value.

### Description

Convert the unsigned integer value  $x$  to a single-precision floating point value in round-towards-zero mode.

## 1.8. SIMD Intrinsics

This section describes SIMD intrinsic functions that are only supported in device code.

## `__device__ unsigned int __vabs2 (unsigned int a)`

Computes per-halfword absolute value.

### Returns

Returns computed value.

### Description

Splits 4 bytes of argument into 2 parts, each consisting of 2 bytes, then computes absolute value for each of parts. Result is stored as unsigned int and returned.

**\_\_device\_\_ unsigned int \_\_vabs4 (unsigned int a)**

Computes per-byte absolute value.

**Returns**

Returns computed value.

**Description**

Splits argument by bytes. Computes absolute value of each byte. Result is stored as unsigned int.

**\_\_device\_\_ unsigned int \_\_vabsdiffs2 (unsigned int a, unsigned int b)**

Computes per-halfword sum of absolute difference of signed integer.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of each into 2 parts, each consisting of 2 bytes. For corresponding parts function computes absolute difference. Result is stored as unsigned int and returned.

**\_\_device\_\_ unsigned int \_\_vabsdiffs4 (unsigned int a, unsigned int b)**

Computes per-byte absolute difference of signed integer.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of each into 4 parts, each consisting of 1 byte. For corresponding parts function computes absolute difference. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vabsdiffu2 (unsigned int a, unsigned int b)`

Performs per-halfword absolute difference of unsigned integer computation:  $|a - b|$ .

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function computes absolute difference. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vabsdiffu4 (unsigned int a, unsigned int b)`

Computes per-byte absolute difference of unsigned integer.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function computes absolute difference. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vabsss2 (unsigned int a)`

Computes per-halfword absolute value with signed saturation.

### Returns

Returns computed value.

### Description

Splits 4 bytes of argument into 2 parts, each consisting of 2 bytes, then computes absolute value with signed saturation for each of parts. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vabsss4 (unsigned int a)`

Computes per-byte absolute value with signed saturation.

### Returns

Returns computed value.

### Description

Splits 4 bytes of argument into 4 parts, each consisting of 1 byte, then computes absolute value with signed saturation for each of parts. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vadd2 (unsigned int a, unsigned int b)`

Performs per-halfword (un)signed addition, with wrap-around:  $a + b$ .

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes, then performs unsigned addition on corresponding parts. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vadd4 (unsigned int a, unsigned int b)`

Performs per-byte (un)signed addition.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte, then performs unsigned addition on corresponding parts. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vaddss2 (unsigned int a, unsigned int b)`

Performs per-halfword addition with signed saturation.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes, then performs addition with signed saturation on corresponding parts. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vaddss4 (unsigned int a, unsigned int b)`

Performs per-byte addition with signed saturation.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte, then performs addition with signed saturation on corresponding parts. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vaddus2 (unsigned int a, unsigned int b)`

Performs per-halfword addition with unsigned saturation.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes, then performs addition with unsigned saturation on corresponding parts.

## `__device__ unsigned int __vaddus4 (unsigned int a, unsigned int b)`

Performs per-byte addition with unsigned saturation.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte, then performs addition with unsigned saturation on corresponding parts.

## `__device__ unsigned int __vavgs2 (unsigned int a, unsigned int b)`

Performs per-halfword signed rounded average computation.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. then computes signed rounded average of corresponding parts. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vavgs4 (unsigned int a, unsigned int b)`

Computes per-byte signed rounded average.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. then computes signed rounded average of corresponding parts. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vavgu2 (unsigned int a, unsigned int b)`

Performs per-halfword unsigned rounded average computation.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. then computes unsigned rounded average of corresponding parts. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vavgu4 (unsigned int a, unsigned int b)`

Performs per-byte unsigned rounded average.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. then computes unsigned rounded average of corresponding parts. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vcmpeq2 (unsigned int a, unsigned int b)`

Performs per-halfword (un)signed comparison.

### Returns

Returns 0xffff computed value.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts result is ffff if they are equal, and 0000 otherwise. For example `__vcmpeq2(0x1234aba5, 0x1234aba6)` returns 0xffff0000.



## `__device__ unsigned int __vcmpeq4 (unsigned int a, unsigned int b)`

Performs per-byte (un)signed comparison.

### Returns

Returns 0xff if a = b, else returns 0.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts result is ff if they are equal, and 00 otherwise. For example `__vcmpeq4(0x1234aba5, 0x1234aba6)` returns 0xffffffff00.

## `__device__ unsigned int __vcmpges2 (unsigned int a, unsigned int b)`

Performs per-halfword signed comparison: a >= b ? 0xffff : 0.

### Returns

Returns 0xffff if a >= b, else returns 0.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts result is ffff if 'a' part >= 'b' part, and 0000 otherwise. For example `__vcmpges2(0x1234aba5, 0x1234aba6)` returns 0xffff0000.

## `__device__ unsigned int __vcmpges4 (unsigned int a, unsigned int b)`

Performs per-byte signed comparison.

### Returns

Returns 0xff if a >= b, else returns 0.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts result is ff if 'a' part >= 'b' part, and 00 otherwise. For example `__vcmpges4(0x1234aba5, 0x1234aba6)` returns 0xffffffff00.

## `__device__ unsigned int __vcmpgeu2 (unsigned int a, unsigned int b)`

Performs per-halfword unsigned comparison:  $a \geq b ? 0xffff : 0$ .

### Returns

Returns 0xffff if  $a \geq b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts result is ffff if 'a' part  $\geq$  'b' part, and 0000 otherwise. For example `__vcmpgeu2(0x1234aba5, 0x1234aba6)` returns 0xffff0000.

## `__device__ unsigned int __vcmpgeu4 (unsigned int a, unsigned int b)`

Performs per-byte unsigned comparison.

### Returns

Returns 0xff if  $a = b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts result is ff if 'a' part  $\geq$  'b' part, and 00 otherwise. For example `__vcmpgeu4(0x1234aba5, 0x1234aba6)` returns 0xffffff00.

## `__device__ unsigned int __vcmpgts2 (unsigned int a, unsigned int b)`

Performs per-halfword signed comparison:  $a > b ? 0xffff : 0$ .

### Returns

Returns 0xffff if  $a > b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts result is ffff if 'a' part  $>$  'b' part, and 0000 otherwise. For example `__vcmpgts2(0x1234aba5, 0x1234aba6)` returns 0x00000000.

## `__device__ unsigned int __vcmpgts4 (unsigned int a, unsigned int b)`

Performs per-byte signed comparison.

### Returns

Returns 0xff if  $a > b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts result is ff if 'a' part > 'b' part, and 00 otherwise. For example `__vcmpgts4(0x1234aba5, 0x1234aba6)` returns 0x00000000.

## `__device__ unsigned int __vcmpgtu2 (unsigned int a, unsigned int b)`

Performs per-halfword unsigned comparison:  $a > b ? 0xffff : 0$ .

### Returns

Returns 0xffff if  $a > b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts result is ffff if 'a' part > 'b' part, and 0000 otherwise. For example `__vcmpgtu2(0x1234aba5, 0x1234aba6)` returns 0x00000000.

## `__device__ unsigned int __vcmpgtu4 (unsigned int a, unsigned int b)`

Performs per-byte unsigned comparison.

### Returns

Returns 0xff if  $a > b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts result is ff if 'a' part > 'b' part, and 00 otherwise. For example `__vcmpgtu4(0x1234aba5, 0x1234aba6)` returns 0x00000000.

## `__device__ unsigned int __vcmples2 (unsigned int a, unsigned int b)`

Performs per-halfword signed comparison:  $a \leq b ? 0xffff : 0$ .

### Returns

Returns 0xffff if  $a \leq b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts result is ffff if 'a' part  $\leq$  'b' part, and 0000 otherwise. For example `__vcmples2(0x1234aba5, 0x1234aba6)` returns 0xffffffff.

## `__device__ unsigned int __vcmples4 (unsigned int a, unsigned int b)`

Performs per-byte signed comparison.

### Returns

Returns 0xff if  $a \leq b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts result is ff if 'a' part  $\leq$  'b' part, and 00 otherwise. For example `__vcmples4(0x1234aba5, 0x1234aba6)` returns 0xffffffff.

## `__device__ unsigned int __vcmpleu2 (unsigned int a, unsigned int b)`

Performs per-halfword unsigned comparison:  $a \leq b ? 0xffff : 0$ .

### Returns

Returns 0xffff if  $a \leq b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts result is ffff if 'a' part  $\leq$  'b' part, and 0000 otherwise. For example `__vcmpleu2(0x1234aba5, 0x1234aba6)` returns 0xffffffff.

## `__device__ unsigned int __vcmpleu4 (unsigned int a, unsigned int b)`

Performs per-byte unsigned comparison.

### Returns

Returns 0xff if  $a \leq b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts result is ff if 'a' part  $\leq$  'b' part, and 00 otherwise. For example `__vcmpleu4(0x1234aba5, 0x1234aba6)` returns 0xffffffff.

## `__device__ unsigned int __vcmplts2 (unsigned int a, unsigned int b)`

Performs per-halfword signed comparison:  $a < b ? 0xffff : 0$ .

### Returns

Returns 0xffff if  $a < b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts result is ffff if 'a' part  $<$  'b' part, and 0000 otherwise. For example `__vcmplts2(0x1234aba5, 0x1234aba6)` returns 0x0000ffff.

## `__device__ unsigned int __vcmplts4 (unsigned int a, unsigned int b)`

Performs per-byte signed comparison.

### Returns

Returns 0xff if  $a < b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts result is ff if 'a' part  $<$  'b' part, and 00 otherwise. For example `__vcmplts4(0x1234aba5, 0x1234aba6)` returns 0x000000ff.

## `__device__ unsigned int __vcmpltu2 (unsigned int a, unsigned int b)`

Performs per-halfword unsigned comparison:  $a < b ? 0xffff : 0$ .

### Returns

Returns 0xffff if  $a < b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts result is ffff if 'a' part < 'b' part, and 0000 otherwise. For example `__vcmpltu2(0x1234aba5, 0x1234aba6)` returns 0x0000ffff.

## `__device__ unsigned int __vcmpltu4 (unsigned int a, unsigned int b)`

Performs per-byte unsigned comparison.

### Returns

Returns 0xff if  $a < b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts result is ff if 'a' part < 'b' part, and 00 otherwise. For example `__vcmpltu4(0x1234aba5, 0x1234aba6)` returns 0x000000ff.

## `__device__ unsigned int __vcmpne2 (unsigned int a, unsigned int b)`

Performs per-halfword (un)signed comparison:  $a != b ? 0xffff : 0$ .

### Returns

Returns 0xffff if  $a != b$ , else returns 0.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts result is ffff if 'a' part != 'b' part, and 0000 otherwise. For example `__vcmplts2(0x1234aba5, 0x1234aba6)` returns 0x0000ffff.

## `__device__ unsigned int __vcmpne4 (unsigned int a, unsigned int b)`

Performs per-byte (un)signed comparison.

### Returns

Returns 0xff if a != b, else returns 0.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts result is ff if 'a' part != 'b' part, and 00 otherwise. For example `__vcmplt4(0x1234aba5, 0x1234aba6)` returns 0x000000ff.

## `__device__ unsigned int __vhaddu2 (unsigned int a, unsigned int b)`

Performs per-halfword unsigned average computation.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. then computes unsigned average of corresponding parts. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vhaddu4 (unsigned int a, unsigned int b)`

Computes per-byte unsigned average.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. then computes unsigned average of corresponding parts. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vmaxs2 (unsigned int a, unsigned int b)`

Performs per-halfword signed maximum computation.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function computes signed maximum. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vmaxs4 (unsigned int a, unsigned int b)`

Computes per-byte signed maximum.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function computes signed maximum. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vmaxu2 (unsigned int a, unsigned int b)`

Performs per-halfword unsigned maximum computation.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function computes unsigned maximum. Result is stored as unsigned int and returned.



## `__device__ unsigned int __vmaxu4 (unsigned int a, unsigned int b)`

Computes per-byte unsigned maximum.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function computes unsigned maximum. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vmins2 (unsigned int a, unsigned int b)`

Performs per-halfword signed minimum computation.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function computes signed minimum. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vmins4 (unsigned int a, unsigned int b)`

Computes per-byte signed minimum.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function computes signed minimum. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vminu2 (unsigned int a, unsigned int b)`

Performs per-halfword unsigned minimum computation.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function computes unsigned minimum. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vminu4 (unsigned int a, unsigned int b)`

Computes per-byte unsigned minimum.

### Returns

Returns computed value.

### Description

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function computes unsigned minimum. Result is stored as unsigned int and returned.

## `__device__ unsigned int __vneg2 (unsigned int a)`

Computes per-halfword negation.

### Returns

Returns computed value.

### Description

Splits 4 bytes of argument into 2 parts, each consisting of 2 bytes. For each part function computes negation. Result is stored as unsigned int and returned.

**\_\_device\_\_ unsigned int \_\_vneg4 (unsigned int a)**

Performs per-byte negation.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of argument into 4 parts, each consisting of 1 byte. For each part function computes negation. Result is stored as unsigned int and returned.

**\_\_device\_\_ unsigned int \_\_vnegss2 (unsigned int a)**

Computes per-halfword negation with signed saturation.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of argument into 2 parts, each consisting of 2 bytes. For each part function computes negation. Result is stored as unsigned int and returned.

**\_\_device\_\_ unsigned int \_\_vnegss4 (unsigned int a)**

Performs per-byte negation with signed saturation.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of argument into 4 parts, each consisting of 1 byte. For each part function computes negation. Result is stored as unsigned int and returned.

**\_\_device\_\_ unsigned int \_\_vsads2 (unsigned int a, unsigned int b)**

Performs per-halfword sum of absolute difference of signed.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts functions computes absolute difference and sum it up. Result is stored as unsigned int and returned.

**\_\_device\_\_ unsigned int \_\_vsads4 (unsigned int a, unsigned int b)**

Computes per-byte sum of abs difference of signed.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts functions computes absolute difference and sum it up. Result is stored as unsigned int and returned.

**\_\_device\_\_ unsigned int \_\_vsadu2 (unsigned int a, unsigned int b)**

Computes per-halfword sum of abs diff of unsigned.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function computes absolute differences, and returns sum of those differences.

**\_\_device\_\_ unsigned int \_\_vsadu4 (unsigned int a, unsigned int b)**

Computes per-byte sum af abs difference of unsigned.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function computes absolute differences, and returns sum of those differences.

**\_\_device\_\_ unsigned int \_\_vseteq2 (unsigned int a, unsigned int b)**

Performs per-halfword (un)signed comparison.

**Returns**

Returns 1 if a = b, else returns 0.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function performs comparison 'a' part == 'b' part. If both equalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vseteq4 (unsigned int a, unsigned int b)**

Performs per-byte (un)signed comparison.

**Returns**

Returns 1 if a = b, else returns 0.

**Description**

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function performs comparison 'a' part == 'b' part. If both equalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetges2 (unsigned int a, unsigned int b)**

Performs per-halfword signed comparison.

**Returns**

Returns 1 if a >= b, else returns 0.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function performs comparison 'a' part >= 'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetges4 (unsigned int a, unsigned int b)**

Performs per-byte signed comparison.

**Returns**

Returns 1 if a >= b, else returns 0.

**Description**

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function performs comparison 'a' part >= 'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetgeu2 (unsigned int a, unsigned int b)**

Performs per-halfword unsigned minimum unsigned comparison.

**Returns**

Returns 1 if a >= b, else returns 0.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function performs comparison 'a' part >= 'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetgeu4 (unsigned int a, unsigned int b)**

Performs per-byte unsigned comparison.

**Returns**

Returns 1 if a >= b, else returns 0.

**Description**

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function performs comparison 'a' part  $\geq$  'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetgts2 (unsigned int a, unsigned int b)**

Performs per-halfword signed comparison.

**Returns**

Returns 1 if  $a > b$ , else returns 0.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function performs comparison 'a' part  $>$  'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetgts4 (unsigned int a, unsigned int b)**

Performs per-byte signed comparison.

**Returns**

Returns 1 if  $a > b$ , else returns 0.

**Description**

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function performs comparison 'a' part  $>$  'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetgtu2 (unsigned int a, unsigned int b)**

Performs per-halfword unsigned comparison.

**Returns**

Returns 1 if  $a > b$ , else returns 0.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function performs comparison 'a' part > 'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetgtu4 (unsigned int a, unsigned int b)**

Performs per-byte unsigned comparison.

**Returns**

Returns 1 if  $a > b$ , else returns 0.

**Description**

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function performs comparison 'a' part > 'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetles2 (unsigned int a, unsigned int b)**

Performs per-halfword unsigned minimum computation.

**Returns**

Returns 1 if  $a \leq b$ , else returns 0.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function performs comparison 'a' part  $\leq$  'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetles4 (unsigned int a, unsigned int b)**

Performs per-byte signed comparison.

**Returns**

Returns 1 if  $a \leq b$ , else returns 0.



**Description**

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function performs comparison 'a' part <= 'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetleu2 (unsigned int a, unsigned int b)**

Performs per-halfword signed comparison.

**Returns**

Returns 1 if a <= b, else returns 0.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function performs comparison 'a' part <= 'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetleu4 (unsigned int a, unsigned int b)**

Performs per-byte unsigned comparison.

**Returns**

Returns 1 if a <= b, else returns 0.

**Description**

Splits 4 bytes of each argument into 4 part, each consisting of 1 byte. For corresponding parts function performs comparison 'a' part <= 'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetlts2 (unsigned int a, unsigned int b)**

Performs per-halfword signed comparison.

**Returns**

Returns 1 if a < b, else returns 0.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function performs comparison 'a' part <= 'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetlts4 (unsigned int a, unsigned int b)**

Performs per-byte signed comparison.

**Returns**

Returns 1 if a < b, else returns 0.

**Description**

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function performs comparison 'a' part <= 'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetltu2 (unsigned int a, unsigned int b)**

Performs per-halfword unsigned comparison.

**Returns**

Returns 1 if a < b, else returns 0.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function performs comparison 'a' part <= 'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetltu4 (unsigned int a, unsigned int b)**

Performs per-byte unsigned comparison.

**Returns**

Returns 1 if a < b, else returns 0.

**Description**

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function performs comparison 'a' part <= 'b' part. If both inequalities are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetne2 (unsigned int a, unsigned int b)**

Performs per-halfword (un)signed comparison.

**Returns**

Returns 1 if a != b, else returns 0.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts function performs comparison 'a' part != 'b' part. If both conditions are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsetne4 (unsigned int a, unsigned int b)**

Performs per-byte (un)signed comparison.

**Returns**

Returns 1 if a != b, else returns 0.

**Description**

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts function performs comparison 'a' part != 'b' part. If both conditions are satisfied, function returns 1.

**\_\_device\_\_ unsigned int \_\_vsub2 (unsigned int a, unsigned int b)**

Performs per-halfword (un)signed subtraction, with wrap-around.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts functions performs subtraction. Result is stored as unsigned int and returned.

**\_\_device\_\_ unsigned int \_\_vsub4 (unsigned int a, unsigned int b)**

Performs per-byte subtraction.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts functions performs subtraction. Result is stored as unsigned int and returned.

**\_\_device\_\_ unsigned int \_\_vsubss2 (unsigned int a, unsigned int b)**

Performs per-halfword (un)signed subtraction, with signed saturation.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts functions performs subtraction with signed saturation. Result is stored as unsigned int and returned.

**\_\_device\_\_ unsigned int \_\_vsubss4 (unsigned int a, unsigned int b)**

Performs per-byte subtraction with signed saturation.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts functions performs subtraction with signed saturation. Result is stored as unsigned int and returned.

**\_\_device\_\_ unsigned int \_\_vsubus2 (unsigned int a, unsigned int b)**

Performs per-halfword subtraction with unsigned saturation.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of each argument into 2 parts, each consisting of 2 bytes. For corresponding parts functions performs subtraction with unsigned saturation. Result is stored as unsigned int and returned.

**\_\_device\_\_ unsigned int \_\_vsubus4 (unsigned int a, unsigned int b)**

Performs per-byte subtraction with unsigned saturation.

**Returns**

Returns computed value.

**Description**

Splits 4 bytes of each argument into 4 parts, each consisting of 1 byte. For corresponding parts functions performs subtraction with unsigned saturation. Result is stored as unsigned int and returned.

## 1.9. Half Precision Intrinsics

This section describes half precision intrinsic functions that are only supported in device code.

### Half Arithmetic Functions

### Half2 Arithmetic Functions

## Half Comparison Functions

## Half2 Comparison Functions

## Half Precision Conversion And Data Movement

## Half Math Functions

## Half2 Math Functions

### 1.9.1. Half Arithmetic Functions

#### Half Precision Intrinsics

**`__device__ __half __hadd (const __half a, const __half b)`**

Performs `half` addition in round-to-nearest-even mode.

#### Returns

Returns the `half` result of adding `a` and `b`.

#### Description

Performs `half` addition of inputs `a` and `b`, in round-to-nearest-even mode.

**`__device__ __half __hadd_sat (const __half a, const __half b)`**

Performs `half` addition in round-to-nearest-even mode, with saturation to `[0.0, 1.0]`.

#### Returns

Returns the `half` result of adding `a` and `b` with saturation.

#### Description

Performs `half` add of inputs `a` and `b`, in round-to-nearest-even mode, and clamps the result to range `[0.0, 1.0]`. NaN results are flushed to `+0.0`.

`__device__ __half __hfma (const __half a, const __half b, const __half c)`

Performs `half` fused multiply-add in round-to-nearest-even mode.

### Returns

Returns the `half` result of the fused multiply-add operation on `a`, `b`, and `c`.

### Description

Performs `half` multiply on inputs `a` and `b`, then performs a `half` add of the result with `c`, rounding the result once in round-to-nearest-even mode.

`__device__ __half __hfma_sat (const __half a, const __half b, const __half c)`

Performs `half` fused multiply-add in round-to-nearest-even mode, with saturation to `[0.0, 1.0]`.

### Returns

Returns the `half` result of the fused multiply-add operation on `a`, `b`, and `c` with saturation.

### Description

Performs `half` multiply on inputs `a` and `b`, then performs a `half` add of the result with `c`, rounding the result once in round-to-nearest-even mode, and clamps the result to range `[0.0, 1.0]`. NaN results are flushed to `+0.0`.

`__device__ __half __hmul (const __half a, const __half b)`

Performs `half` multiplication in round-to-nearest-even mode.

### Returns

Returns the `half` result of multiplying `a` and `b`.

### Description

Performs `half` multiplication of inputs `a` and `b`, in round-to-nearest mode.

**\_\_device\_\_ \_\_half \_\_hmul\_sat (const \_\_half a, const \_\_half b)**

Performs `half` multiplication in round-to-nearest-even mode, with saturation to [0.0, 1.0].

**Returns**

Returns the `half` result of multiplying `a` and `b` with saturation.

**Description**

Performs `half` multiplication of inputs `a` and `b`, in round-to-nearest mode, and clamps the result to range [0.0, 1.0]. NaN results are flushed to +0.0.

**\_\_device\_\_ \_\_half \_\_hneg (const \_\_half a)**

Negates input `half` number and returns the result.

**Returns**

Returns negated `half` input `a`.

**Description**

Negates input `half` number and returns the result.

**\_\_device\_\_ \_\_half \_\_hsub (const \_\_half a, const \_\_half b)**

Performs `half` subtraction in round-to-nearest-even mode.

**Returns**

Returns the `half` result of subtraction `b` from `a`.

**Description**

Subtracts `half` input `b` from input `a` in round-to-nearest mode.

**\_\_device\_\_ \_\_half \_\_hsub\_sat (const \_\_half a, const \_\_half b)**

Performs `half` subtraction in round-to-nearest-even mode, with saturation to [0.0, 1.0].

**Returns**

Returns the `half` result of subtraction `b` from `a` with saturation.

**Description**

Subtracts `half` input `b` from input `a` in round-to-nearest mode, and clamps the result to range [0.0, 1.0]. NaN results are flushed to +0.0.



**\_\_device\_\_ \_\_half2 h2div (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector division in round-to-nearest-even mode.

### Returns

Returns the `half2` vector result of division `a` by `b`.

### Description

Divides `half2` input vector `a` by input vector `b` in round-to-nearest mode.

**\_\_device\_\_ \_\_half hdiv (const \_\_half a, const \_\_half b)**

Performs `half` division in round-to-nearest-even mode.

### Returns

Returns the `half` result of division `a` by `b`.

### Description

Divides `half` input `a` by input `b` in round-to-nearest mode.

## 1.9.2. Half2 Arithmetic Functions

Half Precision Intrinsics

**\_\_device\_\_ \_\_half2 \_\_hadd2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector addition in round-to-nearest-even mode.

### Returns

Returns the `half2` vector result of adding vectors `a` and `b`.

### Description

Performs `half2` vector add of inputs `a` and `b`, in round-to-nearest mode.

**\_\_device\_\_ \_\_half2 \_\_hadd2\_sat (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector addition in round-to-nearest-even mode, with saturation to `[0.0, 1.0]`.

### Returns

Returns the `half2` vector result of adding vectors `a` and `b` with saturation.

**Description**

Performs `half2` vector add of inputs `a` and `b`, in round-to-nearest mode, and clamps the results to range `[0.0, 1.0]`. NaN results are flushed to `+0.0`.

```
__device__ __half2 __hfma2 (const __half2 a, const __half2 b, const
__half2 c)
```

Performs `half2` vector fused multiply-add in round-to-nearest-even mode.

**Returns**

Returns the `half2` vector result of the fused multiply-add operation on vectors `a`, `b`, and `c`.

**Description**

Performs `half2` vector multiply on inputs `a` and `b`, then performs a `half2` vector add of the result with `c`, rounding the result once in round-to-nearest-even mode.

```
__device__ __half2 __hfma2_sat (const __half2 a, const __half2 b,
const __half2 c)
```

Performs `half2` vector fused multiply-add in round-to-nearest-even mode, with saturation to `[0.0, 1.0]`.

**Returns**

Returns the `half2` vector result of the fused multiply-add operation on vectors `a`, `b`, and `c` with saturation.

**Description**

Performs `half2` vector multiply on inputs `a` and `b`, then performs a `half2` vector add of the result with `c`, rounding the result once in round-to-nearest-even mode, and clamps the results to range `[0.0, 1.0]`. NaN results are flushed to `+0.0`.

```
__device__ __half2 __hmul2 (const __half2 a, const __half2 b)
```

Performs `half2` vector multiplication in round-to-nearest-even mode.

**Returns**

Returns the `half2` vector result of multiplying vectors `a` and `b`.

**Description**

Performs `half2` vector multiplication of inputs `a` and `b`, in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half2 \_\_hmul2\_sat (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector multiplication in round-to-nearest-even mode, with saturation to [0.0, 1.0].

### Returns

Returns the `half2` vector result of multiplying vectors `a` and `b` with saturation.

### Description

Performs `half2` vector multiplication of inputs `a` and `b`, in round-to-nearest-even mode, and clamps the results to range [0.0, 1.0]. NaN results are flushed to +0.0.

**\_\_device\_\_ \_\_half2 \_\_hneg2 (const \_\_half2 a)**

Negates both halves of the input `half2` number and returns the result.

### Returns

Returns `half2` number with both halves negated.

### Description

Negates both halves of the input `half2` number `a` and returns the result.

**\_\_device\_\_ \_\_half2 \_\_hsub2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector subtraction in round-to-nearest-even mode.

### Returns

Returns the `half2` vector result of subtraction vector `b` from `a`.

### Description

Subtracts `half2` input vector `b` from input vector `a` in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half2 \_\_hsub2\_sat (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector subtraction in round-to-nearest-even mode, with saturation to [0.0, 1.0].

### Returns

Returns the `half2` vector result of subtraction vector `b` from `a` with saturation.

**Description**

Subtracts `half2` input vector `b` from input vector `a` in round-to-nearest-even mode, and clamps the results to range `[0.0, 1.0]`. NaN results are flushed to `+0.0`.

## 1.9.3. Half Comparison Functions

Half Precision Intrinsics

`__device__ bool __heq (const __half a, const __half b)`

Performs `half` if-equal comparison.

**Returns**

Returns boolean result of if-equal comparison of `a` and `b`.

**Description**

Performs `half` if-equal comparison of inputs `a` and `b`. NaN inputs generate false results.

`__device__ bool __hequ (const __half a, const __half b)`

Performs `half` unordered if-equal comparison.

**Returns**

Returns boolean result of unordered if-equal comparison of `a` and `b`.

**Description**

Performs `half` if-equal comparison of inputs `a` and `b`. NaN inputs generate true results.

`__device__ bool __hge (const __half a, const __half b)`

Performs `half` greater-equal comparison.

**Returns**

Returns boolean result of greater-equal comparison of `a` and `b`.

**Description**

Performs `half` greater-equal comparison of inputs `a` and `b`. NaN inputs generate false results.

**\_\_device\_\_ bool \_\_hgeu (const \_\_half a, const \_\_half b)**

Performs `half` unordered greater-equal comparison.

### Returns

Returns boolean result of unordered greater-equal comparison of `a` and `b`.

### Description

Performs `half` greater-equal comparison of inputs `a` and `b`. NaN inputs generate true results.

**\_\_device\_\_ bool \_\_hgt (const \_\_half a, const \_\_half b)**

Performs `half` greater-than comparison.

### Returns

Returns boolean result of greater-than comparison of `a` and `b`.

### Description

Performs `half` greater-than comparison of inputs `a` and `b`. NaN inputs generate false results.

**\_\_device\_\_ bool \_\_hgtu (const \_\_half a, const \_\_half b)**

Performs `half` unordered greater-than comparison.

### Returns

Returns boolean result of unordered greater-than comparison of `a` and `b`.

### Description

Performs `half` greater-than comparison of inputs `a` and `b`. NaN inputs generate true results.

**\_\_device\_\_ int \_\_hisinf (const \_\_half a)**

Checks if the input `half` number is infinite.

### Returns

Returns -1 iff `a` is equal to negative infinity, 1 iff `a` is equal to positive infinity and 0 otherwise.

**Description**

Checks if the input `half` number `a` is infinite.

`__device__ bool __hisnan (const __half a)`

Determine whether `half` argument is a NaN.

**Returns**

Returns boolean true iff argument is a NaN, boolean false otherwise.

**Description**

Determine whether `half` value `a` is a NaN.

`__device__ bool __hle (const __half a, const __half b)`

Performs `half` less-equal comparison.

**Returns**

Returns boolean result of less-equal comparison of `a` and `b`.

**Description**

Performs `half` less-equal comparison of inputs `a` and `b`. NaN inputs generate false results.

`__device__ bool __hleu (const __half a, const __half b)`

Performs `half` unordered less-equal comparison.

**Returns**

Returns boolean result of unordered less-equal comparison of `a` and `b`.

**Description**

Performs `half` less-equal comparison of inputs `a` and `b`. NaN inputs generate true results.

`__device__ bool __hlt (const __half a, const __half b)`

Performs `half` less-than comparison.

**Returns**

Returns boolean result of less-than comparison of `a` and `b`.

**Description**

Performs `half` less-than comparison of inputs `a` and `b`. NaN inputs generate false results.

`__device__ bool __hltu (const __half a, const __half b)`

Performs `half` unordered less-than comparison.

**Returns**

Returns boolean result of unordered less-than comparison of `a` and `b`.

**Description**

Performs `half` less-than comparison of inputs `a` and `b`. NaN inputs generate true results.

`__device__ bool __hne (const __half a, const __half b)`

Performs `half` not-equal comparison.

**Returns**

Returns boolean result of not-equal comparison of `a` and `b`.

**Description**

Performs `half` not-equal comparison of inputs `a` and `b`. NaN inputs generate false results.

`__device__ bool __hneu (const __half a, const __half b)`

Performs `half` unordered not-equal comparison.

**Returns**

Returns boolean result of unordered not-equal comparison of `a` and `b`.

**Description**

Performs `half` not-equal comparison of inputs `a` and `b`. NaN inputs generate true results.

## 1.9.4. Half2 Comparison Functions

Half Precision Intrinsics

**\_\_device\_\_ bool \_\_hbeq2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector if-equal comparison, and returns boolean true iff both `half` results are true, boolean false otherwise.

**Returns**

Returns boolean true if both `half` results of if-equal comparison of vectors `a` and `b` are true, boolean false otherwise.

**Description**

Performs `half2` vector if-equal comparison of inputs `a` and `b`. The bool result is set to true only if both `half` if-equal comparisons evaluate to true, or false otherwise. NaN inputs generate false results.

**\_\_device\_\_ bool \_\_hbequ2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector unordered if-equal comparison, and returns boolean true iff both `half` results are true, boolean false otherwise.

**Returns**

Returns boolean true if both `half` results of unordered if-equal comparison of vectors `a` and `b` are true, boolean false otherwise.

**Description**

Performs `half2` vector if-equal comparison of inputs `a` and `b`. The bool result is set to true only if both `half` if-equal comparisons evaluate to true, or false otherwise. NaN inputs generate true results.

**\_\_device\_\_ bool \_\_hbge2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector greater-equal comparison, and returns boolean true iff both `half` results are true, boolean false otherwise.

**Returns**

Returns boolean true if both `half` results of greater-equal comparison of vectors `a` and `b` are true, boolean false otherwise.

**Description**

Performs `half2` vector greater-equal comparison of inputs `a` and `b`. The bool result is set to true only if both `half` greater-equal comparisons evaluate to true, or false otherwise. NaN inputs generate false results.



**\_\_device\_\_ bool \_\_hbgeu2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector unordered greater-equal comparison, and returns boolean true iff both `half` results are true, boolean false otherwise.

**Returns**

Returns boolean true if both `half` results of unordered greater-equal comparison of vectors `a` and `b` are true, boolean false otherwise.

**Description**

Performs `half2` vector greater-equal comparison of inputs `a` and `b`. The bool result is set to true only if both `half` greater-equal comparisons evaluate to true, or false otherwise. NaN inputs generate true results.

**\_\_device\_\_ bool \_\_hbgt2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector greater-than comparison, and returns boolean true iff both `half` results are true, boolean false otherwise.

**Returns**

Returns boolean true if both `half` results of greater-than comparison of vectors `a` and `b` are true, boolean false otherwise.

**Description**

Performs `half2` vector greater-than comparison of inputs `a` and `b`. The bool result is set to true only if both `half` greater-than comparisons evaluate to true, or false otherwise. NaN inputs generate false results.

**\_\_device\_\_ bool \_\_hbgtu2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector unordered greater-than comparison, and returns boolean true iff both `half` results are true, boolean false otherwise.

**Returns**

Returns boolean true if both `half` results of unordered greater-than comparison of vectors `a` and `b` are true, boolean false otherwise.

**Description**

Performs `half2` vector greater-than comparison of inputs `a` and `b`. The bool result is set to true only if both `half` greater-than comparisons evaluate to true, or false otherwise. NaN inputs generate true results.

**\_\_device\_\_ bool \_\_hble2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector less-equal comparison, and returns boolean true iff both `half` results are true, boolean false otherwise.

**Returns**

Returns boolean true if both `half` results of less-equal comparison of vectors `a` and `b` are true, boolean false otherwise.

**Description**

Performs `half2` vector less-equal comparison of inputs `a` and `b`. The bool result is set to true only if both `half` less-equal comparisons evaluate to true, or false otherwise. NaN inputs generate false results.

**\_\_device\_\_ bool \_\_hbleu2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector unordered less-equal comparison, and returns boolean true iff both `half` results are true, boolean false otherwise.

**Returns**

Returns boolean true if both `half` results of unordered less-equal comparison of vectors `a` and `b` are true, boolean false otherwise.

**Description**

Performs `half2` vector less-equal comparison of inputs `a` and `b`. The bool result is set to true only if both `half` less-equal comparisons evaluate to true, or false otherwise. NaN inputs generate true results.

**\_\_device\_\_ bool \_\_hblt2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector less-than comparison, and returns boolean true iff both `half` results are true, boolean false otherwise.

**Returns**

Returns boolean true if both `half` results of less-than comparison of vectors `a` and `b` are true, boolean false otherwise.

**Description**

Performs `half2` vector less-than comparison of inputs `a` and `b`. The bool result is set to true only if both `half` less-than comparisons evaluate to true, or false otherwise. NaN inputs generate false results.

**\_\_device\_\_ bool \_\_hbltu2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector unordered less-than comparison, and returns boolean true iff both `half` results are true, boolean false otherwise.

**Returns**

Returns boolean true if both `half` results of unordered less-than comparison of vectors `a` and `b` are true, boolean false otherwise.

**Description**

Performs `half2` vector less-than comparison of inputs `a` and `b`. The bool result is set to true only if both `half` less-than comparisons evaluate to true, or false otherwise. NaN inputs generate true results.

**\_\_device\_\_ bool \_\_hbne2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector not-equal comparison, and returns boolean true iff both `half` results are true, boolean false otherwise.

**Returns**

Returns boolean true if both `half` results of not-equal comparison of vectors `a` and `b` are true, boolean false otherwise.

**Description**

Performs `half2` vector not-equal comparison of inputs `a` and `b`. The bool result is set to true only if both `half` not-equal comparisons evaluate to true, or false otherwise. NaN inputs generate false results.

**\_\_device\_\_ bool \_\_hbneu2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector unordered not-equal comparison, and returns boolean true iff both `half` results are true, boolean false otherwise.

**Returns**

Returns boolean true if both `half` results of unordered not-equal comparison of vectors `a` and `b` are true, boolean false otherwise.

**Description**

Performs `half2` vector not-equal comparison of inputs `a` and `b`. The bool result is set to true only if both `half` not-equal comparisons evaluate to true, or false otherwise. NaN inputs generate true results.

**\_\_device\_\_ \_\_half2 \_\_heq2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector if-equal comparison.

### Returns

Returns the `half2` vector result of if-equal comparison of vectors `a` and `b`.

### Description

Performs `half2` vector if-equal comparison of inputs `a` and `b`. The corresponding `half` results are set to 1.0 for true, or 0.0 for false. NaN inputs generate false results.

**\_\_device\_\_ \_\_half2 \_\_hequ2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector unordered if-equal comparison.

### Returns

Returns the `half2` vector result of unordered if-equal comparison of vectors `a` and `b`.

### Description

Performs `half2` vector if-equal comparison of inputs `a` and `b`. The corresponding `half` results are set to 1.0 for true, or 0.0 for false. NaN inputs generate true results.

**\_\_device\_\_ \_\_half2 \_\_hge2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector greater-equal comparison.

### Returns

Returns the `half2` vector result of greater-equal comparison of vectors `a` and `b`.

### Description

Performs `half2` vector greater-equal comparison of inputs `a` and `b`. The corresponding `half` results are set to 1.0 for true, or 0.0 for false. NaN inputs generate false results.

**\_\_device\_\_ \_\_half2 \_\_hgeu2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector unordered greater-equal comparison.

### Returns

Returns the `half2` vector result of unordered greater-equal comparison of vectors `a` and `b`.

**Description**

Performs `half2` vector greater-equal comparison of inputs `a` and `b`. The corresponding `half` results are set to 1.0 for true, or 0.0 for false. NaN inputs generate true results.

`__device__ __half2 __hgt2 (const __half2 a, const __half2 b)`

Performs `half2` vector greater-than comparison.

**Returns**

Returns the `half2` vector result of greater-than comparison of vectors `a` and `b`.

**Description**

Performs `half2` vector greater-than comparison of inputs `a` and `b`. The corresponding `half` results are set to 1.0 for true, or 0.0 for false. NaN inputs generate false results.

`__device__ __half2 __hgtu2 (const __half2 a, const __half2 b)`

Performs `half2` vector unordered greater-than comparison.

**Returns**

Returns the `half2` vector result of unordered greater-than comparison of vectors `a` and `b`.

**Description**

Performs `half2` vector greater-than comparison of inputs `a` and `b`. The corresponding `half` results are set to 1.0 for true, or 0.0 for false. NaN inputs generate true results.

`__device__ __half2 __hisnan2 (const __half2 a)`

Determine whether `half2` argument is a NaN.

**Returns**

Returns `half2` which has the corresponding `half` results set to 1.0 for true, or 0.0 for false.

**Description**

Determine whether each half of input `half2` number `a` is a NaN.

**\_\_device\_\_ \_\_half2 \_\_hle2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector less-equal comparison.

### Returns

Returns the `half2` vector result of less-equal comparison of vectors `a` and `b`.

### Description

Performs `half2` vector less-equal comparison of inputs `a` and `b`. The corresponding `half` results are set to 1.0 for true, or 0.0 for false. NaN inputs generate false results.

**\_\_device\_\_ \_\_half2 \_\_hleu2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector unordered less-equal comparison.

### Returns

Returns the `half2` vector result of unordered less-equal comparison of vectors `a` and `b`.

### Description

Performs `half2` vector less-equal comparison of inputs `a` and `b`. The corresponding `half` results are set to 1.0 for true, or 0.0 for false. NaN inputs generate true results.

**\_\_device\_\_ \_\_half2 \_\_hlt2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector less-than comparison.

### Returns

Returns the `half2` vector result of less-than comparison of vectors `a` and `b`.

### Description

Performs `half2` vector less-than comparison of inputs `a` and `b`. The corresponding `half` results are set to 1.0 for true, or 0.0 for false. NaN inputs generate false results.

**\_\_device\_\_ \_\_half2 \_\_hltu2 (const \_\_half2 a, const \_\_half2 b)**

Performs `half2` vector unordered less-than comparison.

### Returns

Returns the `half2` vector result of unordered less-than comparison of vectors `a` and `b`.

**Description**

Performs `half2` vector less-than comparison of inputs `a` and `b`. The corresponding `half` results are set to 1.0 for true, or 0.0 for false. NaN inputs generate true results.

`__device__ __half2 __hne2 (const __half2 a, const __half2 b)`

Performs `half2` vector not-equal comparison.

**Returns**

Returns the `half2` vector result of not-equal comparison of vectors `a` and `b`.

**Description**

Performs `half2` vector not-equal comparison of inputs `a` and `b`. The corresponding `half` results are set to 1.0 for true, or 0.0 for false. NaN inputs generate false results.

`__device__ __half2 __hneu2 (const __half2 a, const __half2 b)`

Performs `half2` vector unordered not-equal comparison.

**Returns**

Returns the `half2` vector result of unordered not-equal comparison of vectors `a` and `b`.

**Description**

Performs `half2` vector not-equal comparison of inputs `a` and `b`. The corresponding `half` results are set to 1.0 for true, or 0.0 for false. NaN inputs generate true results.

## 1.9.5. Half Precision Conversion And Data Movement

Half Precision Intrinsics

`__device__ __half2 __float22half2_rn (const float2 a)`

Converts both components of `float2` number to half precision in round-to-nearest-even mode and returns `half2` with converted values.

**Returns**

Returns `half2` which has corresponding halves equal to the converted `float2` components.

**Description**

Converts both components of `float2` to half precision in round-to-nearest mode and combines the results into one `half2` number. Low 16 bits of the return value correspond to `a.x` and high 16 bits of the return value correspond to `a.y`.

**\_\_device\_\_ \_\_half \_\_float2half (const float a)**

Converts float number to half precision in round-to-nearest-even mode and returns `half` with converted value.

**Returns**

Returns `half` result with converted value.

**Description**

Converts float number `a` to half precision in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half2 \_\_float2half2\_rn (const float a)**

Converts input to half precision in round-to-nearest-even mode and populates both halves of `half2` with converted value.

**Returns**

Returns `half2` with both halves equal to the converted half precision number.

**Description**

Converts input `a` to half precision in round-to-nearest-even mode and populates both halves of `half2` with converted value.

**\_\_device\_\_ \_\_half \_\_float2half\_rd (const float a)**

Converts float number to half precision in round-down mode and returns `half` with converted value.

**Returns**

Returns `half` result with converted value.

**Description**

Converts float number `a` to half precision in round-down mode.

**\_\_device\_\_ \_\_half \_\_float2half\_ru (const float a)**

Converts float number to half precision in round-up mode and returns `half` with converted value.

**Returns**

Returns `half` result with converted value.



**Description**

Converts float number `a` to half precision in round-up mode.

**`__device__ __half __float2half_rz (const float a)`**

Converts float number to half precision in round-towards-zero mode and returns `half` with converted value.

**Returns**

Returns `half` result with converted value.

**Description**

Converts float number `a` to half precision in round-towards-zero mode.

**`__device__ __half2 __floats2half2_rn (const float a, const float b)`**

Converts both input floats to half precision in round-to-nearest-even mode and returns `half2` with converted values.

**Returns**

Returns `half2` which has corresponding halves equal to the converted input floats.

**Description**

Converts both input floats to half precision in round-to-nearest-even mode and combines the results into one `half2` number. Low 16 bits of the return value correspond to the input `a`, high 16 bits correspond to the input `b`.

**`__device__ float2 __half22float2 (const __half2 a)`**

Converts both halves of `half2` to `float2` and returns the result.

**Returns**

Returns converted `float2`.

**Description**

Converts both halves of `half2` input `a` to `float2` and returns the result.

**`__device__ float __half2float (const __half a)`**

Converts `half` number to `float`.

**Returns**

Returns `float` result with converted value.

**Description**

Converts half number `a` to float.

`__device__ __half2 __half2half2 (const __half a)`

Returns `half2` with both halves equal to the input value.

**Returns**

Returns `half2` with both halves equal to the input `a`.

**Description**

Returns `half2` number with both halves equal to the input `a half` number.

`__device__ int __half2int_rd (__half h)`

Convert a half to a signed integer in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to a signed integer in round-down mode.

`__device__ int __half2int_rn (__half h)`

Convert a half to a signed integer in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to a signed integer in round-to-nearest-even mode.

`__device__ int __half2int_ru (__half h)`

Convert a half to a signed integer in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to a signed integer in round-up mode.

`__device__ int __half2int_rz (__half h)`

Convert a half to a signed integer in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to a signed integer in round-towards-zero mode.

`__device__ long long int __half2ll_rd (__half h)`

Convert a half to a signed 64-bit integer in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to a signed 64-bit integer in round-down mode.

`__device__ long long int __half2ll_rn (__half h)`

Convert a half to a signed 64-bit integer in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to a signed 64-bit integer in round-to-nearest-even mode.

`__device__ long long int __half2ll_ru (__half h)`

Convert a half to a signed 64-bit integer in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to a signed 64-bit integer in round-up mode.

`__device__ long long int __half2ll_rz (__half h)`

Convert a half to a signed 64-bit integer in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to a signed 64-bit integer in round-towards-zero mode.

`__device__ short int __half2short_rd (__half h)`

Convert a half to a signed short integer in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to a signed short integer in round-down mode.

`__device__ short int __half2short_rn (__half h)`

Convert a half to a signed short integer in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to a signed short integer in round-to-nearest-even mode.

`__device__ short int __half2short_ru (__half h)`

Convert a half to a signed short integer in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to a signed short integer in round-up mode.

`__device__ short int __half2short_rz (__half h)`

Convert a half to a signed short integer in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to a signed short integer in round-towards-zero mode.

`__device__ unsigned int __half2uint_rd (__half h)`

Convert a half to an unsigned integer in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to an unsigned integer in round-down mode.

`__device__ unsigned int __half2uint_rn (__half h)`

Convert a half to an unsigned integer in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to an unsigned integer in round-to-nearest-even mode.

`__device__ unsigned int __half2uint_ru (__half h)`

Convert a half to an unsigned integer in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to an unsigned integer in round-up mode.

`__device__ unsigned int __half2uint_rz (__half h)`

Convert a half to an unsigned integer in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to an unsigned integer in round-towards-zero mode.

`__device__ unsigned long long int __half2ull_rd (__half h)`

Convert a half to an unsigned 64-bit integer in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to an unsigned 64-bit integer in round-down mode.

`__device__ unsigned long long int __half2ull_rn (__half h)`

Convert a half to an unsigned 64-bit integer in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to an unsigned 64-bit integer in round-to-nearest-even mode.

`__device__ unsigned long long int __half2ull_ru (__half h)`

Convert a half to an unsigned 64-bit integer in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to an unsigned 64-bit integer in round-up mode.

`__device__ unsigned long long int __half2ull_rz (__half h)`

Convert a half to an unsigned 64-bit integer in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to an unsigned 64-bit integer in round-towards-zero mode.

`__device__ unsigned short int __half2ushort_rd (__half h)`

Convert a half to an unsigned short integer in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to an unsigned short integer in round-down mode.

`__device__ unsigned short int __half2ushort_rn (__half h)`

Convert a half to an unsigned short integer in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to an unsigned short integer in round-to-nearest-even mode.

`__device__ unsigned short int __half2ushort_ru (__half h)`

Convert a half to an unsigned short integer in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to an unsigned short integer in round-up mode.

**`__device__ unsigned short int __half2ushort_rz (__half h)`**

Convert a half to an unsigned short integer in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the half-precision floating point value `h` to an unsigned short integer in round-towards-zero mode.

**`__device__ __half2 __halves2half2 (const __half a, const __half b)`**

Combines two `half` numbers into one `half2` number.

**Returns**

Returns `half2` number which has one half equal to `a` and the other to `b`.

**Description**

Combines two input `half` number `a` and `b` into one `half2` number. Input `a` is stored in low 16 bits of the return value, input `b` is stored in high 16 bits of the return value.

**`__device__ float __high2float (const __half2 a)`**

Converts high 16 bits of `half2` to float and returns the result.

**Returns**

Returns high 16 bits of `a` converted to float.

**Description**

Converts high 16 bits of `half2` input `a` to 32 bit floating point number and returns the result.

**`__device__ __half __high2half (const __half2 a)`**

Returns high 16 bits of `half2` input.

**Returns**

Returns `half` which contains high 16 bits of the input.



**Description**

Returns high 16 bits of `half2` input `a`.

`__device__ __half2 __high2half2 (const __half2 a)`

Extracts high 16 bits from `half2` input.

**Returns**

Returns `half2` with both halves equal to high 16 bits from the input.

**Description**

Extracts high 16 bits from `half2` input `a` and returns a new `half2` number which has both halves equal to the extracted bits.

`__device__ __half2 __highs2half2 (const __half2 a, const __half2 b)`

Extracts high 16 bits from each of the two `half2` inputs and combines into one `half2` number.

**Returns**

Returns `half2` which contains high 16 bits from `a` and `b`.

**Description**

Extracts high 16 bits from each of the two `half2` inputs and combines into one `half2` number. High 16 bits from input `a` is stored in low 16 bits of the return value, high 16 bits from input `b` is stored in high 16 bits of the return value.

`__device__ __half __int2half_rd (int i)`

Convert a signed integer to a half in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the signed integer value `i` to a half-precision floating point value in round-down mode.

**\_\_device\_\_ \_\_half \_\_int2half\_rn (int i)**

Convert a signed integer to a half in round-to-nearest-even mode.

### Returns

Returns converted value.

### Description

Convert the signed integer value *i* to a half-precision floating point value in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half \_\_int2half\_ru (int i)**

Convert a signed integer to a half in round-up mode.

### Returns

Returns converted value.

### Description

Convert the signed integer value *i* to a half-precision floating point value in round-up mode.

**\_\_device\_\_ \_\_half \_\_int2half\_rz (int i)**

Convert a signed integer to a half in round-towards-zero mode.

### Returns

Returns converted value.

### Description

Convert the signed integer value *i* to a half-precision floating point value in round-towards-zero mode.

**\_\_device\_\_ \_\_half \_\_ll2half\_rd (long long int i)**

Convert a signed 64-bit integer to a half in round-down mode.

### Returns

Returns converted value.

**Description**

Convert the signed 64-bit integer value `i` to a half-precision floating point value in round-down mode.

`__device__ __half __ll2half_rn (long long int i)`

Convert a signed 64-bit integer to a half in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the signed 64-bit integer value `i` to a half-precision floating point value in round-to-nearest-even mode.

`__device__ __half __ll2half_ru (long long int i)`

Convert a signed 64-bit integer to a half in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the signed 64-bit integer value `i` to a half-precision floating point value in round-up mode.

`__device__ __half __ll2half_rz (long long int i)`

Convert a signed 64-bit integer to a half in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the signed 64-bit integer value `i` to a half-precision floating point value in round-towards-zero mode.

`__device__ float __low2float (const __half2 a)`

Converts low 16 bits of `half2` to float and returns the result.

**Returns**

Returns low 16 bits of `a` converted to float.

**Description**

Converts low 16 bits of `half2` input `a` to 32 bit floating point number and returns the result.

`__device__ __half __low2half (const __half2 a)`

Returns low 16 bits of `half2` input.

**Returns**

Returns `half` which contains low 16 bits of the input.

**Description**

Returns low 16 bits of `half2` input `a`.

`__device__ __half2 __low2half2 (const __half2 a)`

Extracts low 16 bits from `half2` input.

**Returns**

Returns `half2` with both halves equal to low 16 bits from the input.

**Description**

Extracts low 16 bits from `half2` input `a` and returns a new `half2` number which has both halves equal to the extracted bits.

`__device__ __half2 __lowhigh2highlow (const __half2 a)`

Swaps both halves of the `half2` input.

**Returns**

Returns `half2` with halves swapped.

**Description**

Swaps both halves of the `half2` input and returns a new `half2` number with swapped halves.

`__device__ __half2 __lows2half2 (const __half2 a, const __half2 b)`

Extracts low 16 bits from each of the two `half2` inputs and combines into one `half2` number.

**Returns**

Returns `half2` which contains low 16 bits from `a` and `b`.

**Description**

Extracts low 16 bits from each of the two `half2` inputs and combines into one `half` number. Low 16 bits from input `a` is stored in low 16 bits of the return value, low 16 bits from input `b` is stored in high 16 bits of the return value.

`__device__ __half __short2half_rd (short int i)`

Convert a signed short integer to a half in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the signed short integer value `i` to a half-precision floating point value in round-down mode.

`__device__ __half __short2half_rn (short int i)`

Convert a signed short integer to a half in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the signed short integer value `i` to a half-precision floating point value in round-to-nearest-even mode.

`__device__ __half __short2half_ru (short int i)`

Convert a signed short integer to a half in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the signed short integer value `i` to a half-precision floating point value in round-up mode.

**\_\_device\_\_ \_\_half \_\_short2half\_rz (short int i)**

Convert a signed short integer to a half in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the signed short integer value *i* to a half-precision floating point value in round-towards-zero mode.

**\_\_device\_\_ \_\_half \_\_uint2half\_rd (unsigned int i)**

Convert an unsigned integer to a half in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned integer value *i* to a half-precision floating point value in round-down mode.

**\_\_device\_\_ \_\_half \_\_uint2half\_rn (unsigned int i)**

Convert an unsigned integer to a half in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned integer value *i* to a half-precision floating point value in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half \_\_uint2half\_ru (unsigned int i)**

Convert an unsigned integer to a half in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned integer value `i` to a half-precision floating point value in round-up mode.

`__device__ __half __uint2half_rz (unsigned int i)`

Convert an unsigned integer to a half in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned integer value `i` to a half-precision floating point value in round-towards-zero mode.

`__device__ __half __ull2half_rd (unsigned long long int i)`

Convert an unsigned 64-bit integer to a half in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned 64-bit integer value `i` to a half-precision floating point value in round-down mode.

`__device__ __half __ull2half_rn (unsigned long long int i)`

Convert an unsigned 64-bit integer to a half in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned 64-bit integer value `i` to a half-precision floating point value in round-to-nearest-even mode.

`__device__ __half __ull2half_ru (unsigned long long int i)`

Convert an unsigned 64-bit integer to a half in round-up mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned 64-bit integer value `i` to a half-precision floating point value in round-up mode.

`__device__ __half __ull2half_rz (unsigned long long int i)`

Convert an unsigned 64-bit integer to a half in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned 64-bit integer value `i` to a half-precision floating point value in round-towards-zero mode.

`__device__ __half __ushort2half_rd (unsigned short int i)`

Convert an unsigned short integer to a half in round-down mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned short integer value `i` to a half-precision floating point value in round-down mode.

`__device__ __half __ushort2half_rn (unsigned short int i)`

Convert an unsigned short integer to a half in round-to-nearest-even mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned short integer value `i` to a half-precision floating point value in round-to-nearest-even mode.

`__device__ __half __ushort2half_ru (unsigned short int i)`

Convert an unsigned short integer to a half in round-up mode.

**Returns**

Returns converted value.



**Description**

Convert the unsigned short integer value `i` to a half-precision floating point value in round-up mode.

**`__device__ __half __ushort2half_rz (unsigned short int i)`**

Convert an unsigned short integer to a half in round-towards-zero mode.

**Returns**

Returns converted value.

**Description**

Convert the unsigned short integer value `i` to a half-precision floating point value in round-towards-zero mode.

## 1.9.6. Half Math Functions

Half Precision Intrinsics

**`__device__ __half hceil (const __half h)`**

Calculate ceiling of the input argument.

**Returns**

Returns ceiling expressed as a half-precision floating point number.

**Description**

Compute the smallest integer value not less than `h`.

**`__device__ __half hcos (const __half a)`**

Calculates `half` cosine in round-to-nearest-even mode.

**Returns**

Returns `half` cosine of `a`.

**Description**

Calculates `half` cosine of input `a` in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half hexp (const \_\_half a)**

Calculates `half` natural exponential function in round-to-nearest mode.

**Returns**

Returns `half` natural exponential function of `a`.

**Description**

Calculates `half` natural exponential function of input `a` in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half hexp10 (const \_\_half a)**

Calculates `half` decimal exponential function in round-to-nearest mode.

**Returns**

Returns `half` decimal exponential function of `a`.

**Description**

Calculates `half` decimal exponential function of input `a` in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half hexp2 (const \_\_half a)**

Calculates `half` binary exponential function in round-to-nearest mode.

**Returns**

Returns `half` binary exponential function of `a`.

**Description**

Calculates `half` binary exponential function of input `a` in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half hfloor (const \_\_half h)**

Calculate the largest integer less than or equal to `h`.

**Returns**

Returns floor expressed as half-precision floating point number.

**Description**

Calculate the largest integer value which is less than or equal to `h`.

**\_\_device\_\_ \_\_half hlog (const \_\_half a)**

Calculates `half` natural logarithm in round-to-nearest-even mode.

**Returns**

Returns `half` natural logarithm of `a`.

**Description**

Calculates `half` natural logarithm of input `a` in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half hlog10 (const \_\_half a)**

Calculates `half` decimal logarithm in round-to-nearest-even mode.

**Returns**

Returns `half` decimal logarithm of `a`.

**Description**

Calculates `half` decimal logarithm of input `a` in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half hlog2 (const \_\_half a)**

Calculates `half` binary logarithm in round-to-nearest-even mode.

**Returns**

Returns `half` binary logarithm of `a`.

**Description**

Calculates `half` binary logarithm of input `a` in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half hrcp (const \_\_half a)**

Calculates `half` reciprocal in round-to-nearest-even mode.

**Returns**

Returns `half` reciprocal of `a`.

**Description**

Calculates `half` reciprocal of input `a` in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half hrint (const \_\_half h)**

Round input to nearest integer value in half-precision floating point number.

**Returns**

Returns rounded integer value expressed as half-precision floating point number.

**Description**

Round `h` to the nearest integer value in half-precision floating point format, with halfway cases rounded to the nearest even integer value.

**\_\_device\_\_ \_\_half hrsqrt (const \_\_half a)**

Calculates `half` reciprocal square root in round-to-nearest-even mode.

**Returns**

Returns `half` reciprocal square root of `a`.

**Description**

Calculates `half` reciprocal square root of input `a` in round-to-nearest mode.

**\_\_device\_\_ \_\_half hsin (const \_\_half a)**

Calculates `half` sine in round-to-nearest-even mode.

**Returns**

Returns `half` sine of `a`.

**Description**

Calculates `half` sine of input `a` in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half hsqrt (const \_\_half a)**

Calculates `half` square root in round-to-nearest-even mode.

**Returns**

Returns `half` square root of `a`.

**Description**

Calculates `half` square root of input `a` in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half htrunc (const \_\_half h)**

Truncate input argument to the integral part.

### Returns

Returns truncated integer value.

### Description

Round `h` to the nearest integer value that does not exceed `h` in magnitude.

## 1.9.7. Half2 Math Functions

Half Precision Intrinsics

**\_\_device\_\_ \_\_half2 h2ceil (const \_\_half2 h)**

Calculate `half2` vector ceiling of the input argument.

### Returns

Returns `half2` vector ceiling expressed as a pair of half-precision floating point numbers.

### Description

For each component of vector `h` compute the smallest integer value not less than `h`.

**\_\_device\_\_ \_\_half2 h2cos (const \_\_half2 a)**

Calculates `half2` vector cosine in round-to-nearest-even mode.

### Returns

Returns `half2` cosine of vector `a`.

### Description

Calculates `half2` cosine of input vector `a` in round-to-nearest-even mode.

**\_\_device\_\_ \_\_half2 h2exp (const \_\_half2 a)**

Calculates `half2` vector exponential function in round-to-nearest mode.

### Returns

Returns `half2` exponential function of vector `a`.

**Description**

Calculates `half2` exponential function of input vector `a` in round-to-nearest-even mode.

`__device__ __half2 h2exp10 (const __half2 a)`

Calculates `half2` vector decimal exponential function in round-to-nearest-even mode.

**Returns**

Returns `half2` decimal exponential function of vector `a`.

**Description**

Calculates `half2` decimal exponential function of input vector `a` in round-to-nearest-even mode.

`__device__ __half2 h2exp2 (const __half2 a)`

Calculates `half2` vector binary exponential function in round-to-nearest-even mode.

**Returns**

Returns `half2` binary exponential function of vector `a`.

**Description**

Calculates `half2` binary exponential function of input vector `a` in round-to-nearest-even mode.

`__device__ __half2 h2floor (const __half2 h)`

Calculate the largest integer less than or equal to `h`.

**Returns**

Returns `half2` vector floor expressed as a pair of half-precision floating point number.

**Description**

For each component of vector `h` calculate the largest integer value which is less than or equal to `h`.

`__device__ __half2 h2log (const __half2 a)`

Calculates `half2` vector natural logarithm in round-to-nearest-even mode.

**Returns**

Returns `half2` natural logarithm of vector `a`.

**Description**

Calculates `half2` natural logarithm of input vector `a` in round-to-nearest-even mode.

`__device__ __half2 h2log10 (const __half2 a)`

Calculates `half2` vector decimal logarithm in round-to-nearest-even mode.

**Returns**

Returns `half2` decimal logarithm of vector `a`.

**Description**

Calculates `half2` decimal logarithm of input vector `a` in round-to-nearest-even mode.

`__device__ __half2 h2log2 (const __half2 a)`

Calculates `half2` vector binary logarithm in round-to-nearest-even mode.

**Returns**

Returns `half2` binary logarithm of vector `a`.

**Description**

Calculates `half2` binary logarithm of input vector `a` in round-to-nearest mode.

`__device__ __half2 h2rcp (const __half2 a)`

Calculates `half2` vector reciprocal in round-to-nearest-even mode.

**Returns**

Returns `half2` reciprocal of vector `a`.

**Description**

Calculates `half2` reciprocal of input vector `a` in round-to-nearest-even mode.

`__device__ __half2 h2rint (const __half2 h)`

Round input to nearest integer value in half-precision floating point number.

**Returns**

Returns `half2` vector of rounded integer values expressed as half-precision floating point numbers.

**Description**

Round each component of `half2` vector `h` to the nearest integer value in half-precision floating point format, with halfway cases rounded to the nearest even integer value.

`__device__ __half2 h2rsqrt (const __half2 a)`

Calculates `half2` vector reciprocal square root in round-to-nearest mode.

**Returns**

Returns `half2` reciprocal square root of vector `a`.

**Description**

Calculates `half2` reciprocal square root of input vector `a` in round-to-nearest-even mode.

`__device__ __half2 h2sin (const __half2 a)`

Calculates `half2` vector sine in round-to-nearest-even mode.

**Returns**

Returns `half2` sine of vector `a`.

**Description**

Calculates `half2` sine of input vector `a` in round-to-nearest-even mode.

`__device__ __half2 h2sqrt (const __half2 a)`

Calculates `half2` vector square root in round-to-nearest-even mode.

**Returns**

Returns `half2` square root of vector `a`.

**Description**

Calculates `half2` square root of input vector `a` in round-to-nearest mode.

`__device__ __half2 h2trunc (const __half2 h)`

Truncate `half2` vector input argument to the integral part.

**Returns**

Returns `half2` vector truncated integer value.



**Description**

Round each component of vector  $h$  to the nearest integer value that does not exceed  $h$  in magnitude.

## **Notice**

ALL NVIDIA DESIGN SPECIFICATIONS, REFERENCE BOARDS, FILES, DRAWINGS, DIAGNOSTICS, LISTS, AND OTHER DOCUMENTS (TOGETHER AND SEPARATELY, "MATERIALS") ARE BEING PROVIDED "AS IS." NVIDIA MAKES NO WARRANTIES, EXPRESSED, IMPLIED, STATUTORY, OR OTHERWISE WITH RESPECT TO THE MATERIALS, AND EXPRESSLY DISCLAIMS ALL IMPLIED WARRANTIES OF NONINFRINGEMENT, MERCHANTABILITY, AND FITNESS FOR A PARTICULAR PURPOSE.

Information furnished is believed to be accurate and reliable. However, NVIDIA Corporation assumes no responsibility for the consequences of use of such information or for any infringement of patents or other rights of third parties that may result from its use. No license is granted by implication of otherwise under any patent rights of NVIDIA Corporation. Specifications mentioned in this publication are subject to change without notice. This publication supersedes and replaces all other information previously supplied. NVIDIA Corporation products are not authorized as critical components in life support devices or systems without express written approval of NVIDIA Corporation.

## **Trademarks**

NVIDIA and the NVIDIA logo are trademarks or registered trademarks of NVIDIA Corporation in the U.S. and other countries. Other company and product names may be trademarks of the respective companies with which they are associated.

## **Copyright**

© 2007-2016 NVIDIA Corporation. All rights reserved.